
Envy-Free Chore Allocation with Unit Subsidies

under Negative Dichotomous Valuations

Mainak Sarkar Soumyarup Sadhukhan
Indian Institute of Technology Kanpur

September 22, 2026

Outline

- ① Setting and result
- ② Model
- ③ The construction
- ④ Failure of the goods construction
- ⑤ Pareto optimality
- ⑥ Summary

The setting

M : a set of m indivisible *chores*; $N = [n]$: agents, each with $v_i : 2^M \rightarrow \mathbb{R}$, $v_i(\emptyset) = 0$.

Each agent also receives a subsidy $p_i \geq 0$ from a third party.

Envy-free solution

$(\mathcal{A}, p)$ is an *envy-free solution* if

$$v_i(A_i) + p_i \geq v_i(A_j) + p_j \quad \text{for all } i, j.$$

Money enters linearly, at the same rate for everyone, and each item is worth at most one unit — the usual normalisation before absolute subsidy bounds mean anything.

The question is not existence. It is **how much money**.

Negative dichotomous valuations

For every agent i , every $S \subseteq M$ and every $e \notin S$,

$$v_i(S \cup \{e\}) - v_i(S) \in \{-1, 0\}.$$

The sign analogue of the dichotomous goods model. **Nothing further is assumed** — not additivity, not submodularity, not even subadditivity.

Negative dichotomous valuations

For every agent i , every $S \subseteq M$ and every $e \notin S$,

$$v_i(S \cup \{e\}) - v_i(S) \in \{-1, 0\}.$$

The sign analogue of the dichotomous goods model. **Nothing further is assumed** — not additivity, not submodularity, not even subadditivity.

Why binary marginals are natural for chores. An extra task either creates a fresh unit of burden, or is absorbed by a commitment already made:

- an agent already covering an on-call window absorbs another incident in that window at no extra cost;
- a driver already routed through a district absorbs one more stop.

Whether a chore is costly depends on the rest of the bundle — hence non-additivity is the point, not an artefact.

Main result

Theorem

For every instance with negative dichotomous valuations there exist a *complete* allocation $\mathcal{A}$ and a subsidy vector $p \in \{0, 1\}^n$ with

$$\sum_{i \in N} p_i \leq n - 1$$

such that $(\mathcal{A}, p)$ is an envy-free solution. Moreover $\mathcal{A}$ is *EF1 even before subsidies are paid*, and both are computable in polynomial time in the value-oracle model.

Main result

Theorem

For every instance with negative dichotomous valuations there exist a *complete* allocation $\mathcal{A}$ and a subsidy vector $p \in \{0, 1\}^n$ with

$$\sum_{i \in N} p_i \leq n - 1$$

such that $(\mathcal{A}, p)$ is an envy-free solution. Moreover $\mathcal{A}$ is *EF1 even before subsidies are paid*, and both are computable in polynomial time in the value-oracle model.

- the per-agent bound is **0 or 1** — not merely $O(1)$;
- the total bound $n - 1$ is **tight**;
- EF1 is **extra** — it is not guaranteed by the goods-side construction.

What was already available here

Negative dichotomous valuations are a subclass of the *doubly monotone* valuations, so one bound comes for free.

Kawase, Makino, Sumita, Tamura, Yokoo (2025)

For doubly monotone valuations with marginals bounded by one in absolute value, any EF1 allocation converts into an envy-free solution with

$$n - 1 \text{ per agent,} \quad n(n - 1)/2 \text{ in total.}$$

That is already independent of m — and it was the best known here.

What was already available here

Negative dichotomous valuations are a subclass of the *doubly monotone* valuations, so one bound comes for free.

Kawase, Makino, Sumita, Tamura, Yokoo (2025)

For doubly monotone valuations with marginals bounded by one in absolute value, any EF1 allocation converts into an envy-free solution with

$$n - 1 \text{ per agent,} \quad n(n - 1)/2 \text{ in total.}$$

That is already independent of m — and it was the best known here.

We sharpen $n - 1$ per agent to 1 , and $n(n - 1)/2$ to $n - 1$ in total: a linear-factor improvement in both.

The total bound $n - 1$ is tight

Take n agents, $n - 1$ chores, and identical *binary additive* costs:

$$v_i(S) = -|S| \quad \text{for every agent } i.$$

In any complete allocation some agent h receives the empty bundle. Envy-freeness of agent i towards h reads

$$-|A_i| + p_i \geq v_i(A_h) + p_h = p_h,$$

hence

$$p_i \geq |A_i| + p_h \geq |A_i|.$$

The total bound $n - 1$ is tight

Take n agents, $n - 1$ chores, and identical *binary additive* costs:

$$v_i(S) = -|S| \quad \text{for every agent } i.$$

In any complete allocation some agent h receives the empty bundle. Envy-freeness of agent i towards h reads

$$-|A_i| + p_i \geq v_i(A_h) + p_h = p_h,$$

hence

$$p_i \geq |A_i| + p_h \geq |A_i|.$$

Summing over all agents,

$$\sum_{i \in N} p_i \geq \sum_{i \in N} |A_i| = n - 1.$$

So no construction can do better in total.

□

Model and cost form

M : m indivisible chores. $N = [n]$: agents. $v_i : 2^M \rightarrow \mathbb{R}$ with $v_i(\emptyset) = 0$.

Algorithms run in the *value-oracle* model: an oracle returns $v_i(S)$ for a queried S .

Model and cost form

M : m indivisible chores. $N = [n]$: agents. $v_i : 2^M \rightarrow \mathbb{R}$ with $v_i(\emptyset) = 0$.

Algorithms run in the *value-oracle* model: an oracle returns $v_i(S)$ for a queried S .

It is convenient to pass to non-negative *costs*

$$c_i(S) := -v_i(S).$$

Dichotomous cost function

$c(\emptyset) = 0$ and $c(S \cup \{e\}) - c(S) \in \{0, 1\}$ for all S and $e \notin S$.

Negative dichotomous valuations $\leftrightarrow$ dichotomous costs, under $c_i = -v_i$. In particular every c_i is **non-decreasing** and **integer-valued** — both used constantly below.

Write $c_i(e \mid S) := c_i(S \cup \{e\}) - c_i(S) \in \{0, 1\}$ for a marginal.

Fairness notions, in cost form

Envy-freeness

$c_i(A_i) \leq c_i(A_j)$ for all i, j .

Envy-free solution

$c_i(A_i) - p_i \leq c_i(A_j) - p_j$ for all i, j .

Fairness notions, in cost form

Envy-freeness

$c_i(A_i) \leq c_i(A_j)$ for all i, j .

Envy-free solution

$c_i(A_i) - p_i \leq c_i(A_j) - p_j$ for all i, j .

EF1 for chores

A complete $\mathcal{A}$ is *EF1* if for all i, j either

$$c_i(A_i) \leq c_i(A_j), \quad \text{or} \quad \exists e \in A_i : c_i(A_i \setminus \{e\}) \leq c_i(A_j).$$

The item removed comes from the *envious agent's own* bundle — the natural form for chores, and equivalent to the usual two-sided definition on all-chores instances.

A partial allocation X has *residue* $R(X) := M \setminus \bigcup_i X_i$.

The equality graph

Given a partial allocation X , define the *equality graph* $G = (N, E)$ on the agents:

$$(i, j) \in E \iff i \neq j \text{ and } c_i(X_i) = c_i(X_j).$$

An arc $i \rightarrow j$ records that agent i is *exactly indifferent* between her own bundle and j 's — she is on the verge of envying j .

The equality graph

Given a partial allocation X , define the *equality graph* $G = (N, E)$ on the agents:

$$(i, j) \in E \iff i \neq j \text{ and } c_i(X_i) = c_i(X_j).$$

An arc $i \rightarrow j$ records that agent i is *exactly indifferent* between her own bundle and j 's — she is on the verge of envying j .

Strongly connected component; tail component

u, v are *strongly connected* if $u = v$, or each is reachable from the other. The equivalence classes are the *strongly connected components* (SCCs).

An SCC S is a *tail component* if no arc leaves it: there is no $(i, j) \in E$ with $i \in S$ and $j \notin S$.

Tail components exist

Lemma

Every finite non-empty directed graph has at least one tail strongly connected component.

PROOF. Start at any SCC. If it is not a tail component, follow an outgoing arc to another SCC and repeat.

No SCC can be visited twice: that would make the two components mutually reachable, hence a single component.

There are finitely many SCCs, so the walk terminates — at a component with no outgoing arc, i.e. a tail component. □

Tail components exist


Lemma

Every finite non-empty directed graph has at least one tail strongly connected component.

PROOF. Start at any SCC. If it is not a tail component, follow an outgoing arc to another SCC and repeat.

No SCC can be visited twice: that would make the two components mutually reachable, hence a single component.

There are finitely many SCCs, so the walk terminates — at a component with no outgoing arc, i.e. a tail component. □



Starting point: an envy-free *partial* allocation

Theorem [Tao, Wu, Yu, Zhou 2025, Thm. 6.1]

For every n and all dichotomous costs $c_1, \dots, c_n$ there is a polynomial-time algorithm returning a partial allocation X that is envy-free,

$$c_i(X_i) \leq c_i(X_j) \quad \text{for all } i, j,$$

and leaves $|R(X)| \leq n - 1$ chores unassigned.

Starting point: an envy-free *partial* allocation

Theorem [Tao, Wu, Yu, Zhou 2025, Thm. 6.1]

For every n and all dichotomous costs $c_1, \dots, c_n$ there is a polynomial-time algorithm returning a partial allocation X that is envy-free,

$$c_i(X_i) \leq c_i(X_j) \quad \text{for all } i, j,$$

and leaves $|R(X)| \leq n - 1$ chores unassigned.

We need more than the statement — we need the **state the algorithm halts in**. So we recall its three rules.

It maintains X together with the equality graph G , and while $R(X) \neq \emptyset$ applies, *in order*:

The three rules

- (R1) If some $e \in R(X)$ and some i have $c_i(e \mid X_i) = 0$ — a *free* chore for i — assign e to i .
- (R2) Otherwise, if some $e \in R(X)$ and some arc $(i, j) \in E$ on a *directed cycle* of G have $c_i(e \mid X_j) = 0$, rotate the bundles along that cycle and assign e to i .
- (R3) Otherwise, select a tail SCC S of G . If $|R(X)| \geq |S|$, hand out distinct residual chores to the agents of S , one each. Otherwise **halt**.

The three rules

- (R1) If some $e \in R(X)$ and some i have $c_i(e \mid X_i) = 0$ — a *free* chore for i — assign e to i .
- (R2) Otherwise, if some $e \in R(X)$ and some arc $(i, j) \in E$ on a *directed cycle* of G have $c_i(e \mid X_j) = 0$, rotate the bundles along that cycle and assign e to i .
- (R3) Otherwise, select a tail SCC S of G . If $|R(X)| \geq |S|$, hand out distinct residual chores to the agents of S , one each. Otherwise **halt**.

So if the algorithm halts with residue left over:

- neither (R1) nor (R2) applies at that state, and
- for the tail SCC S selected at the halting step, $r := |R(X)| < |S|$.

That last inequality is what gives us *room*: more agents in S than chores left to place.

What the terminal state gives us

Fix X at termination; $R := R(X)$, $r := |R|$, G the equality graph there, and S the tail SCC from the halting step (so $r < |S|$).

Lemma (terminal state)

Suppose $r \geq 1$. Then

- (i) $c_i(e \mid X_i) = 1$ for every $i \in N$, $e \in R$;
- (ii) $c_i(e \mid X_j) = 1$ for every arc $(i, j) \in E$ with $i, j \in S$, and every $e \in R$;
- (iii) $c_i(X_i) < c_i(X_j)$ for every $i \in S$, $j \in N \setminus S$.

In words:

- (i) no leftover chore is free to anybody;
- (ii) inside S , no leftover chore is free even on a neighbour's bundle;
- (iii) agents in S are *strictly* better off than everyone outside S .

Proof of the terminal lemma

Since $r \geq 1$, the algorithm halted with non-empty residue. Hence neither (R1) nor (R2) was applicable.

(i) If $c_i(e \mid X_i) = 0$ for some i and $e \in R$, rule (R1) would apply. Marginals lie in $\{0, 1\}$, so $c_i(e \mid X_i) = 1$.

Proof of the terminal lemma

Since $r \geq 1$, the algorithm halted with non-empty residue. Hence neither (R1) nor (R2) was applicable.

- (i) If $c_i(e \mid X_i) = 0$ for some i and $e \in R$, rule (R1) would apply. Marginals lie in $\{0, 1\}$, so $c_i(e \mid X_i) = 1$.
- (ii) Let $(i, j) \in E$ with $i, j \in S$. As S is strongly connected there is a directed path $j \rightsquigarrow i$; take one without repeated vertices and append the arc (i, j) — a directed cycle through (i, j) .

If $c_i(e \mid X_j) = 0$ for some $e \in R$, rule (R2) would apply to that arc and that cycle. So the marginal is not 0, hence 1.

Proof of the terminal lemma

Since $r \geq 1$, the algorithm halted with non-empty residue. Hence neither (R1) nor (R2) was applicable.

(i) If $c_i(e \mid X_i) = 0$ for some i and $e \in R$, rule (R1) would apply. Marginals lie in $\{0, 1\}$, so $c_i(e \mid X_i) = 1$.

(ii) Let $(i, j) \in E$ with $i, j \in S$. As S is strongly connected there is a directed path $j \rightsquigarrow i$; take one without repeated vertices and append the arc (i, j) — a directed cycle through (i, j) .

If $c_i(e \mid X_j) = 0$ for some $e \in R$, rule (R2) would apply to that arc and that cycle. So the marginal is not 0, hence 1.

(iii) Let $i \in S, j \notin S$. As S is a *tail* component no arc leaves it, so $(i, j) \notin E$, i.e.

$$c_i(X_i) \neq c_i(X_j).$$

Envy-freeness of X gives $c_i(X_i) \leq c_i(X_j)$. The inequality is therefore strict. □

The completion

Assume $r \geq 1$ and write $R = \{e_1, \dots, e_r\}$.

Since $r < |S|$, we may pick r *distinct* agents

$$T = \{t_1, \dots, t_r\} \subseteq S,$$

arbitrarily, and define

$$A_i := \begin{cases} X_{t_k} \cup \{e_k\}, & i = t_k \quad (1 \leq k \leq r), \\ X_i, & i \in N \setminus T. \end{cases}$$

The t_k are distinct and each chore of R is used once, so the bundles stay disjoint and now cover M : **$\mathcal{A}$ is complete.**

Agents in T are the *recipients* — each absorbs exactly one leftover chore, and by (i) that costs her exactly +1.

The completion

Assume $r \geq 1$ and write $R = \{e_1, \dots, e_r\}$.

Since $r < |S|$, we may pick r *distinct* agents

$$T = \{t_1, \dots, t_r\} \subseteq S,$$

arbitrarily, and define

$$A_i := \begin{cases} X_{t_k} \cup \{e_k\}, & i = t_k \quad (1 \leq k \leq r), \\ X_i, & i \in N \setminus T. \end{cases}$$

The t_k are distinct and each chore of R is used once, so the bundles stay disjoint and now cover M : **$\mathcal{A}$ is complete.**

Agents in T are the *recipients* — each absorbs exactly one leftover chore, and by (i) that costs her exactly +1.

What remains: decide *who gets paid*, and show one unit each is enough.

Who gets paid: the subsidy set

Definition (subsidy set)

Let $P \subseteq N$ be the *smallest* set with

(P1) $T \subseteq P$; and

(P2) if $i \in N \setminus S$ and $(i, j) \in E$ for some $j \in P$, then $i \in P$.

Then set $p_i := 1$ if $i \in P$, and $p_i := 0$ otherwise.

Equivalently: start from $P := T$ and repeatedly absorb any agent *outside* S with an equality arc *into* P — a **backward closure** in G .

Who gets paid: the subsidy set

Definition (subsidy set)

Let $P \subseteq N$ be the *smallest* set with

(P1) $T \subseteq P$; and

(P2) if $i \in N \setminus S$ and $(i, j) \in E$ for some $j \in P$, then $i \in P$.

Then set $p_i := 1$ if $i \in P$, and $p_i := 0$ otherwise.

Equivalently: start from $P := T$ and repeatedly absorb any agent *outside* S with an equality arc *into* P — a **backward closure** in G .

The intuition. A recipient's cost rose by one, so she needs a unit. Paying her may make her *neighbours* envious — exactly those who were indifferent to her, i.e. those with an equality arc into her. The payment must therefore propagate backwards.

The restriction $i \notin S$ in (P2) stops the closure from swallowing S itself — and that is what keeps the budget below n .

The algorithm

Step 2 is due to Tao et al.; steps 5–10 are the contribution

Input: agents $N = [n]$, chores M , value oracles $\{v_i\}$.

Output: complete $\mathcal{A}$ that is EF1, and $p \in \{0, 1\}^n$ with $(\mathcal{A}, p)$ envy-free, $\sum_i p_i \leq n - 1$.

1. Set $c_i = -v_i$ for every i .
2. Run the partial-allocation algorithm of Tao et al. on $\{c_i\}$, keeping the tail SCC selected at the terminating application of (R3) if the residue is non-empty.
3. Let X be its terminal partial allocation, $R = M \setminus \bigcup_i X_i$.
4. **if** $R = \emptyset$ **then return** $(X, 0)$.
5. Build the terminal equality graph $G = (N, E)$.
6. Write $R = \{e_1, \dots, e_r\}$; let S be that tail SCC. $r < |S|$
7. Choose distinct $T = \{t_1, \dots, t_r\} \subseteq S$.
8. $A_{t_k} \leftarrow X_{t_k} \cup \{e_k\}$ for $k \leq r$; $A_i \leftarrow X_i$ for $i \notin T$.
9. $P \leftarrow T$; **while** $\exists i \in N \setminus S, j \in P$ with $(i, j) \in E, i \notin P$: $P \leftarrow P \cup \{i\}$.
10. $p_i \leftarrow \mathbf{1}[i \in P]$; **return** $(\mathcal{A}, p)$.

The closure does not eat into S

Lemma

$P \cap S = T$. Equivalently $P \setminus T \subseteq N \setminus S$: every subsidised agent inside S is a recipient.

PROOF. $T \subseteq P$ by (P1) and $T \subseteq S$, so $T \subseteq P \cap S$.

Suppose for contradiction some $i \in (P \cap S) \setminus T$, and put $P' := P \setminus \{i\}$.

P' still satisfies (P1), since $i \notin T$.

The closure does not eat into S

Lemma

$P \cap S = T$. Equivalently $P \setminus T \subseteq N \setminus S$: every subsidised agent inside S is a recipient.

PROOF. $T \subseteq P$ by (P1) and $T \subseteq S$, so $T \subseteq P \cap S$.

Suppose for contradiction some $i \in (P \cap S) \setminus T$, and put $P' := P \setminus \{i\}$.

P' still satisfies (P1), since $i \notin T$.

P' also satisfies (P2). Let $k \in N \setminus S$ with $(k, j) \in E$ for some $j \in P'$. Since $P' \subseteq P$ and P satisfies (P2), $k \in P$. Moreover $k \neq i$, because $k \notin S$ whereas $i \in S$. Hence $k \in P'$.

So P' satisfies both conditions and is strictly smaller than P — contradicting minimality. □

The budget

$$p \in \{0, 1\}^n \text{ and } \sum_{i \in N} p_i \leq n - 1.$$

PROOF. $p \in \{0, 1\}^n$ is immediate from the definition.

We have $|T| = r < |S|$, so $S \setminus T \neq \emptyset$.

By the previous lemma $P \cap S = T$, so every agent of $S \setminus T$ lies *outside* P and therefore receives **zero** subsidy.

At least one agent is unpaid, hence

$$\sum_{i \in N} p_i = |P| \leq n - 1.$$

□

The budget

$$p \in \{0, 1\}^n \text{ and } \sum_{i \in N} p_i \leq n - 1.$$

PROOF. $p \in \{0, 1\}^n$ is immediate from the definition.

We have $|T| = r < |S|$, so $S \setminus T \neq \emptyset$.

By the previous lemma $P \cap S = T$, so every agent of $S \setminus T$ lies *outside* P and therefore receives **zero** subsidy.

At least one agent is unpaid, hence

$$\sum_{i \in N} p_i = |P| \leq n - 1.$$

□

The slack came from $r < |S|$ — the halting condition of (R3).

The key lemma

Lemma

Let $i \in S$, $t \in T$ and $e \in R$. Then

$$c_i(X_t \cup \{e\}) \geq c_i(X_i) + 1.$$

That is: to an agent of S , *any* recipient's completed bundle looks at least one unit worse than her own.

PROOF. Envy-freeness of X gives $c_i(X_i) \leq c_i(X_t)$.

The key lemma

Lemma

Let $i \in S$, $t \in T$ and $e \in R$. Then

$$c_i(X_t \cup \{e\}) \geq c_i(X_i) + 1.$$

That is: to an agent of S , *any* recipient's completed bundle looks at least one unit worse than her own.

PROOF. Envy-freeness of X gives $c_i(X_i) \leq c_i(X_t)$.

Strict case. If $c_i(X_i) < c_i(X_t)$, integrality gives $c_i(X_t) \geq c_i(X_i) + 1$, and monotonicity gives $c_i(X_t \cup \{e\}) \geq c_i(X_t)$. Done.

The key lemma

Lemma

Let $i \in S$, $t \in T$ and $e \in R$. Then

$$c_i(X_t \cup \{e\}) \geq c_i(X_i) + 1.$$

That is: to an agent of S , *any* recipient's completed bundle looks at least one unit worse than her own.

PROOF. Envy-freeness of X gives $c_i(X_i) \leq c_i(X_t)$.

Strict case. If $c_i(X_i) < c_i(X_t)$, integrality gives $c_i(X_t) \geq c_i(X_i) + 1$, and monotonicity gives $c_i(X_t \cup \{e\}) \geq c_i(X_t)$. Done.

Equality case. Suppose $c_i(X_i) = c_i(X_t)$.

- If $i = t$: terminal lemma (i) gives $c_i(e \mid X_i) = 1$.
- If $i \neq t$: then $(i, t) \in E$, and $i, t \in S$, so terminal lemma (ii) gives $c_i(e \mid X_t) = 1$.

Either way $c_i(X_t \cup \{e\}) = c_i(X_i) + 1$. □

Envy-freeness of the completion

Theorem

With $\mathcal{A}$ and p as constructed, $(\mathcal{A}, p)$ is an envy-free solution.

Fix i . Since $X_j \subseteq A_j$ for every j , monotonicity and envy-freeness of X give the standing inequality

$$c_i(A_j) \geq c_i(X_j) \geq c_i(X_i) \quad \text{for every } j. \quad (*)$$

We must check $c_i(A_i) - p_i \leq c_i(A_j) - p_j$ for every j . Three cases, according to where i sits:

$$\begin{array}{ccc} \underbrace{i \in T}_{\text{recipient}} & \underbrace{i \in P \setminus T}_{\text{paid, not a recipient}} & \underbrace{i \notin P}_{\text{unpaid}} \end{array}$$

(These are exhaustive: N is the disjoint union of T , $P \setminus T$ and $N \setminus P$.)

Case 1: $i \in T$ is a recipient

Then $p_i = 1$, and if $e \in R$ is the chore she absorbed, terminal lemma (i) gives

$$c_i(A_i) = c_i(X_i) + c_i(e \mid X_i) = c_i(X_i) + 1, \quad \text{so} \quad c_i(A_i) - p_i = c_i(X_i).$$

So we need $c_i(A_j) - p_j \geq c_i(X_i)$ for every j .

Case 1: $i \in T$ is a recipient

Then $p_i = 1$, and if $e \in R$ is the chore she absorbed, terminal lemma (i) gives

$$c_i(A_i) = c_i(X_i) + c_i(e \mid X_i) = c_i(X_i) + 1, \quad \text{so} \quad c_i(A_i) - p_i = c_i(X_i).$$

So we need $c_i(A_j) - p_j \geq c_i(X_i)$ for every j .

$j \notin P$ Then $j \notin T$, so $A_j = X_j$ and $p_j = 0$. Envy-freeness of X :

$$c_i(A_j) - p_j = c_i(X_j) \geq c_i(X_i). \quad \checkmark$$

Case 1: $i \in T$ is a recipient

Then $p_i = 1$, and if $e \in R$ is the chore she absorbed, terminal lemma (i) gives

$$c_i(A_i) = c_i(X_i) + c_i(e \mid X_i) = c_i(X_i) + 1, \quad \text{so} \quad c_i(A_i) - p_i = c_i(X_i).$$

So we need $c_i(A_j) - p_j \geq c_i(X_i)$ for every j .

$j \notin P$ Then $j \notin T$, so $A_j = X_j$ and $p_j = 0$. Envy-freeness of X :

$$c_i(A_j) - p_j = c_i(X_j) \geq c_i(X_i). \checkmark$$

$j \in P$ Here $p_j = 1$, so we need $c_i(A_j) \geq c_i(X_i) + 1$.

- If $j \in T$: this is exactly the key lemma.
- If $j \in P \setminus T$: then $j \notin S$ by the closure lemma, and $i \in S$, so terminal lemma (iii) gives $c_i(X_i) < c_i(X_j)$; integrality then gives $c_i(A_j) = c_i(X_j) \geq c_i(X_i) + 1$. $\checkmark$

Cases 2 and 3

Case 2: $i \in P \setminus T$ Here $A_i = X_i$ and $p_i = 1$, so $c_i(A_i) - p_i = c_i(X_i) - 1$.

For every j , the standing inequality (*) and $p_j \leq 1$ give

$$c_i(A_j) - p_j \geq c_i(X_i) - 1 = c_i(A_i) - p_i. \checkmark$$

(Being paid without taking on extra work is comfortable.)

Cases 2 and 3

Case 2: $i \in P \setminus T$ Here $A_i = X_i$ and $p_i = 1$, so $c_i(A_i) - p_i = c_i(X_i) - 1$.

For every j , the standing inequality (*) and $p_j \leq 1$ give

$$c_i(A_j) - p_j \geq c_i(X_i) - 1 = c_i(A_i) - p_i. \checkmark$$

(Being paid without taking on extra work is comfortable.)

Case 3: $i \notin P$ Then $i \notin T$, so $A_i = X_i$, $p_i = 0$, and $c_i(A_i) - p_i = c_i(X_i)$.

If $j \notin P$: $A_j = X_j$, $p_j = 0$, and envy-freeness of X closes it.

If $j \in P$, we claim $c_i(A_j) \geq c_i(X_i) + 1$:

- $i \notin S$. Envy-freeness gives $c_i(X_i) \leq c_i(X_j)$. Equality would put $(i, j) \in E$; with $i \notin S$ and $j \in P$, rule (P2) would force $i \in P$ — contradiction. So the inequality is strict, and integrality with monotonicity gives the claim.
- $i \in S$. If $j \in T$ this is the key lemma. If $j \in P \setminus T$ then $j \notin S$ by the closure lemma, and terminal lemma (iii) plus integrality give the claim.

Since $p_j = 1$, we get $c_i(A_j) - p_j \geq c_i(X_i)$. $\checkmark$

□

EF1 comes for free

Corollary

$\mathcal{A}$ is EF1 — before any subsidy is paid. Precisely:

- (i) every non-recipient $i \in N \setminus T$ is *envy-free* in $\mathcal{A}$;
- (ii) every recipient t_k satisfies $c_{t_k}(A_{t_k} \setminus \{e_k\}) \leq c_{t_k}(A_j)$ for every j .

PROOF. Recall X is envy-free and $X_j \subseteq A_j$ for every j .

(i) For $i \notin T$ we have $A_i = X_i$, so for every j

$$c_i(A_i) = c_i(X_i) \leq c_i(X_j) \leq c_i(A_j),$$

using envy-freeness of X and then monotonicity.

EF1 comes for free

Corollary

$\mathcal{A}$ is EF1 — before any subsidy is paid. Precisely:

- (i) every non-recipient $i \in N \setminus T$ is *envy-free* in $\mathcal{A}$;
- (ii) every recipient t_k satisfies $c_{t_k}(A_{t_k} \setminus \{e_k\}) \leq c_{t_k}(A_j)$ for every j .

PROOF. Recall X is envy-free and $X_j \subseteq A_j$ for every j .

(i) For $i \notin T$ we have $A_i = X_i$, so for every j

$$c_i(A_i) = c_i(X_i) \leq c_i(X_j) \leq c_i(A_j),$$

using envy-freeness of X and then monotonicity.

(ii) For $i = t_k$ we have $A_i = X_i \cup \{e_k\}$, so $A_i \setminus \{e_k\} = X_i$, and the same two steps give

$$c_i(A_i \setminus \{e_k\}) = c_i(X_i) \leq c_i(X_j) \leq c_i(A_j).$$

Removing the single chore she absorbed erases all her envy.

□

Putting it together

PROOF OF THE MAIN THEOREM. Pass to costs $c_i = -v_i$, dichotomous. Run the Tao et al. algorithm; let X be its terminal partial allocation, R the residue, $r = |R|$.

$r = 0$ X is already complete and envy-free, so $(X, \mathbf{0})$ is an envy-free solution with zero subsidy — and every envy-free allocation is EF1.

Putting it together

PROOF OF THE MAIN THEOREM. Pass to costs $c_i = -v_i$, dichotomous. Run the Tao et al. algorithm; let X be its terminal partial allocation, R the residue, $r = |R|$.

$r = 0$ X is already complete and envy-free, so $(X, \mathbf{0})$ is an envy-free solution with zero subsidy — and every envy-free allocation is EF1.

$r \geq 1$ Build G , take the tail SCC S , and construct $T, \mathcal{A}, p$. Then

- $\mathcal{A}$ is complete;
- $(\mathcal{A}, p)$ is envy-free (envy-freeness theorem)
- $p \in \{0, 1\}^n, \sum_i p_i \leq n - 1$ (budget)
- $\mathcal{A}$ is EF1 (corollary)

Putting it together

PROOF OF THE MAIN THEOREM. Pass to costs $c_i = -v_i$, dichotomous. Run the Tao et al. algorithm; let X be its terminal partial allocation, R the residue, $r = |R|$.

$r = 0$ X is already complete and envy-free, so $(X, \mathbf{0})$ is an envy-free solution with zero subsidy — and every envy-free allocation is EF1.

$r \geq 1$ Build G , take the tail SCC S , and construct $T, \mathcal{A}, p$. Then

- $\mathcal{A}$ is complete;
- $(\mathcal{A}, p)$ is envy-free (envy-freeness theorem)
- $p \in \{0, 1\}^n, \sum_i p_i \leq n - 1$ (budget)
- $\mathcal{A}$ is EF1 (corollary)

Running time Tao et al. is polynomial. G has n vertices and is read off the n^2 values $c_i(X_j)$ — one oracle query each. SCCs, hence a tail component, in polynomial time. Completing $\mathcal{A}$ costs $r \leq n - 1$ assignments, and P is a backward reachability search in $O(n^2)$. $\square$

Why not just mirror the goods construction?

The goods-side algorithm of Barman et al. is *incremental*: allocate items one at a time, maintaining the invariant that the partial allocation is envy-freeable with subsidies in $\{0, 1\}^n$.

The obvious chore analogue: give each new chore to an agent who currently receives *no* subsidy.

Why not just mirror the goods construction?

The goods-side algorithm of Barman et al. is *incremental*: allocate items one at a time, maintaining the invariant that the partial allocation is envy-freeable with subsidies in $\{0, 1\}^n$.

The obvious chore analogue: give each new chore to an agent who currently receives *no* subsidy.

This does not work.

There are negative dichotomous instances where a partial allocation admits an envy-free solution with subsidies in $\{0, 1\}^n$, yet after one more chore arrives *no* assignment of it preserves the unit-subsidy guarantee — even if the already-allocated bundles may first be permuted among the agents.

So the one-item extension principle fails for chores, and a different argument is needed. Here is the witness.

The instance

Three agents $N = \{1, 2, 3\}$, three chores $M = \{x, y, z\}$. Costs depend only on bundle size:

$$c_1(S) = \max \{0, |S| - 1\}, \quad c_2(S) = c_3(S) = \min \{|S|, 2\}.$$

$ S $	0	1	2	3
$c_1(S)$	0	0	1	2
$c_2(S) = c_3(S)$	0	1	2	2

Every marginal is 0 or 1: these are dichotomous costs. Agent 1 gets her first chore free; agents 2 and 3 saturate at two.

The instance

Three agents $N = \{1, 2, 3\}$, three chores $M = \{x, y, z\}$. Costs depend only on bundle size:

$$c_1(S) = \max \{0, |S| - 1\}, \quad c_2(S) = c_3(S) = \min \{|S|, 2\}.$$

$ S $	0	1	2	3
$c_1(S)$	0	0	1	2
$c_2(S) = c_3(S)$	0	1	2	2

Every marginal is 0 or 1: these are dichotomous costs. Agent 1 gets her first chore free; agents 2 and 3 saturate at two.

Consider the partial allocation

$$X_1 = \{x, y\}, \quad X_2 = X_3 = \emptyset,$$

with z still unallocated.

The state is legitimately reachable

$$(\emptyset, \emptyset, \emptyset) \longrightarrow (\{x\}, \emptyset, \emptyset) \longrightarrow (\{x, y\}, \emptyset, \emptyset)$$

with pointwise-minimal subsidies

$$(0, 0, 0), \quad (0, 0, 0), \quad (1, 0, 0)$$

respectively — all within $\{0, 1\}^3$. So an incremental algorithm can legitimately arrive here.

At X : agent 1 costs her bundle at 1 and either empty bundle at 0, so she needs one unit; agents 2, 3 cost their own bundles at 0 and X_1 at 2, so they have no envy.

The state is legitimately reachable

$$(\emptyset, \emptyset, \emptyset) \longrightarrow (\{x\}, \emptyset, \emptyset) \longrightarrow (\{x, y\}, \emptyset, \emptyset)$$

with pointwise-minimal subsidies

$$(0, 0, 0), \quad (0, 0, 0), \quad (1, 0, 0)$$

respectively — all within $\{0, 1\}^3$. So an incremental algorithm can legitimately arrive here.

At X : agent 1 costs her bundle at 1 and either empty bundle at 0, so she needs one unit; agents 2, 3 cost their own bundles at 0 and X_1 at 2, so they have no envy.

Claim

z cannot be inserted so that the result admits an envy-free subsidy vector in $\{0, 1\}^3$ — even allowing $\{x, y\}, \emptyset, \emptyset$ to be permuted among the agents first.

The three ways to place z

(a) z joins the two-chore bundle Bundle sizes $(3, 0, 0)$. Every agent values a three-chore bundle at 2 and an empty bundle at 0, so whoever holds it needs a subsidy at least **two** units above an agent holding nothing. Impossible in $\{0, 1\}^3$.

The three ways to place z

(a) z joins the two-chore bundle Bundle sizes $(3, 0, 0)$. Every agent values a three-chore bundle at 2 and an empty bundle at 0, so whoever holds it needs a subsidy at least **two** units above an agent holding nothing. Impossible in $\{0, 1\}^3$.

(b) z joins an empty bundle, and an agent of $\{2, 3\}$ holds the pair Bundle sizes $(2, 1, 0)$. That agent values her pair at 2 and the empty bundle at 0 — again a gap of **two**. Impossible.

The three ways to place z

(a) z joins the two-chore bundle Bundle sizes $(3, 0, 0)$. Every agent values a three-chore bundle at 2 and an empty bundle at 0, so whoever holds it needs a subsidy at least **two** units above an agent holding nothing. Impossible in $\{0, 1\}^3$.

(b) z joins an empty bundle, and an agent of $\{2, 3\}$ holds the pair Bundle sizes $(2, 1, 0)$. That agent values her pair at 2 and the empty bundle at 0 — again a gap of **two**. Impossible.

(c) z joins an empty bundle, and agent 1 holds the pair The only remaining option. Let $i \in \{2, 3\}$ take the singleton and j the empty bundle. Then

$$\underbrace{1 - p_1 \leq 0 - p_i}_{\text{agent 1 vs } i} \implies p_1 \geq p_i + 1, \quad \underbrace{1 - p_i \leq -p_j}_{\text{agent } i \text{ vs } j} \implies p_i \geq p_j + 1.$$

Chaining: $p_1 \geq p_j + 2$ — impossible for $p \in \{0, 1\}^3$.

□

The instance itself is fine

The obstruction is to the *method*, not to the instance.

The complete allocation $(\{x\}, \{y\}, \{z\})$ is envy-free **with no subsidy at all**.

Indeed agent 1 values every singleton at cost 0, and agents 2 and 3 value every singleton at cost 1 — nobody envies anybody.

The instance itself is fine

The obstruction is to the *method*, not to the instance.

The complete allocation $(\{x\}, \{y\}, \{z\})$ is envy-free **with no subsidy at all**.

Indeed agent 1 values every singleton at cost 0, and agents 2 and 3 value every singleton at cost 1 — nobody envies anybody.

What went wrong. Reaching that allocation from $(\{x, y\}, \emptyset, \emptyset)$ requires *taking back* chores already placed on agent 1. An incremental construction that only ever adds items cannot do it.

This is why our proof starts from a *global* object — the terminal state of the Tao et al. algorithm — rather than building up one chore at a time.

A further result on the domain

A separate question about negative dichotomous instances, independent of which algorithm produces the allocation: can one ask for the allocation to be *Pareto optimal* as well?

Pareto optimality is judged on the chore allocation alone, ignoring subsidies: $\mathcal{B}$ *dominates* $\mathcal{A}$ if $c_i(\mathcal{B}_i) \leq c_i(\mathcal{A}_i)$ for every i , strictly for some.

A further result on the domain

A separate question about negative dichotomous instances, independent of which algorithm produces the allocation: can one ask for the allocation to be *Pareto optimal* as well?

Pareto optimality is judged on the chore allocation alone, ignoring subsidies: $\mathcal{B}$ *dominates* $\mathcal{A}$ if $c_i(\mathcal{B}_i) \leq c_i(\mathcal{A}_i)$ for every i , strictly for some.

Proposition

There are negative dichotomous instances — already with two agents and identical binary *submodular* costs — in which *no* Pareto-optimal allocation admits an envy-free subsidy vector in $\{0, 1\}^n$.

So on this domain, unit subsidies and Pareto optimality are simply incompatible — for *any* construction, not just ours.

The instance

Two agents, three chores $\{a, b, c\}$, both with

$$c_1(S) = c_2(S) = \min \{|S|, 2\}.$$

Marginals in $\{0, 1\}$, so dichotomous. Submodular, since $k \mapsto \min \{k, 2\}$ is non-decreasing and concave. Not additive: $c_i(\{a, b, c\}) = 2 < 3 = \sum_g c_i(\{g\})$.

The instance

Two agents, three chores $\{a, b, c\}$, both with

$$c_1(S) = c_2(S) = \min \{|S|, 2\}.$$

Marginals in $\{0, 1\}$, so dichotomous. Submodular, since $k \mapsto \min \{k, 2\}$ is non-decreasing and concave. Not additive: $c_i(\{a, b, c\}) = 2 < 3 = \sum_g c_i(\{g\})$.

Up to symmetry every complete allocation has bundle sizes $(3, 0)$ or $(2, 1)$, giving cost profiles

$$(2, 0), \quad (2, 1), \quad (1, 2), \quad (0, 2).$$

The clash

The concentrated allocations are Pareto optimal None of $(2, 1)$, $(1, 2)$, $(0, 2)$ dominates $(2, 0)$ — each either raises some agent's cost, or matches it without improving the other. Symmetrically for $(0, 2)$.

The clash

The concentrated allocations are Pareto optimal None of $(2, 1)$, $(1, 2)$, $(0, 2)$ dominates $(2, 0)$ — each either raises some agent's cost, or matches it without improving the other. Symmetrically for $(0, 2)$.

But they cannot be paid for If agent 1 holds all three chores, envy-freeness needs

$$2 - p_1 \leq -p_2 \implies p_1 - p_2 \geq 2,$$

impossible in $\{0, 1\}^2$. Symmetrically for agent 2.

The clash

The concentrated allocations are Pareto optimal None of $(2, 1)$, $(1, 2)$, $(0, 2)$ dominates $(2, 0)$ — each either raises some agent's cost, or matches it without improving the other. Symmetrically for $(0, 2)$.

But they cannot be paid for If agent 1 holds all three chores, envy-freeness needs

$$2 - p_1 \leq -p_2 \implies p_1 - p_2 \geq 2,$$

impossible in $\{0, 1\}^2$. Symmetrically for agent 2.

And the affordable allocations are dominated A split $(2, 1)$ has profile $(2, 1)$ and is envy-free with $(1, 0)$ — but handing the third chore to the agent already holding two moves the profile to $(2, 0)$: her cost is unchanged, the other's strictly falls. $\square$

Result

Negative dichotomous chores admit a complete allocation $\mathcal{A}$ and subsidies $p \in \{0, 1\}^n$ with $\sum_i p_i \leq n - 1$ such that $(\mathcal{A}, p)$ is envy-free; $\mathcal{A}$ is EF1 before subsidies; all computable in polynomial time.

The shape of the argument.

1. Terminal state of the Tao et al. algorithm: envy-free, $\leq n - 1$ chores left.
2. Its halting condition gives a tail SCC S with $r < |S|$ — *more agents than leftover chores*.
3. Park the leftovers on distinct $T \subseteq S$; each pays +1.
4. Pay T , propagating backwards along equality arcs, never into S .
5. $S \setminus T \neq \emptyset$ stays unpaid: exactly the $n - 1$ budget.

Summary

Result

Negative dichotomous chores admit a complete allocation $\mathcal{A}$ and subsidies $p \in \{0, 1\}^n$ with $\sum_i p_i \leq n - 1$ such that $(\mathcal{A}, p)$ is envy-free; $\mathcal{A}$ is EF1 before subsidies; all computable in polynomial time.

The shape of the argument.

1. Terminal state of the Tao et al. algorithm: envy-free, $\leq n - 1$ chores left.
2. Its halting condition gives a tail SCC S with $r < |S|$ — *more agents than leftover chores*.
3. Park the leftovers on distinct $T \subseteq S$; each pays +1.
4. Pay T , propagating backwards along equality arcs, never into S .
5. $S \setminus T \neq \emptyset$ stays unpaid: exactly the $n - 1$ budget.

Also. The goods-side incremental construction does not transfer; and on this domain, Pareto optimality and unit subsidies are incompatible.

Directions

- **General binary marginals.** Marginals in $\{-1, 0, 1\}$ — goods and chores together, without additivity. Does the unit-subsidy guarantee survive? Open even for $n = 3$.
- **Beyond binary.** Marginals in $\{-k, \dots, 0\}$: what is the right per-agent bound?
- **Efficiency.** What is the least total cost achievable among allocations that admit a subsidy vector in $\{0, 1\}^n$?
- **Truthfulness.** Goko et al. obtain truthfulness for matroidal goods under free disposal. Is anything comparable possible for dichotomous chores?

Thank you

Questions welcome.

References I

- S. Barman, A. Krishna, Y. Narahari, S. Sadhukhan. *Achieving Envy-Freeness with Limited Subsidies under Dichotomous Valuations*. IJCAI 2022, 60–66.
- U. Bhaskar, A. R. Sricharan, R. Vaish. *On Approximate Envy-Freeness for Indivisible Chores and Mixed Resources*. APPROX/RANDOM 2021, LIPIcs 207, 1:1–1:23.
- J. Brustle, J. Dippel, V. V. Narayan, M. Suzuki, A. Vetta. *One Dollar Each Eliminates Envy*. ACM EC 2020, 23–39.
- E. Budish. *The Combinatorial Assignment Problem*. Journal of Political Economy 119(6), 2011.
- H. Goko, A. Igarashi, Y. Kawase, K. Makino, H. Sumita, A. Tamura, Y. Yokoi, M. Yokoo. *A Fair and Truthful Mechanism with Limited Subsidy*. Games and Economic Behavior 144, 2024, 49–70.
- D. Halpern, N. Shah. *Fair Division with Subsidy*. SAGT 2019, LNCS 11801, 374–389.
- Y. Kawase, K. Makino, H. Sumita, A. Tamura, M. Yokoo. *Towards Optimal Subsidy Bounds for Envy-Freeable Allocations*. Artificial Intelligence 348, 2025, 104406.

References II

- Z. Lin, S. Liu, B. Tao, S. Zhou. *On the Existence of EFX Allocations for Binary XOS and Supermodular Chores*. 2026.
- B. Tao, X. Wu, Z. Yu, S. Zhou. *On the Existence of EFX (and Pareto-Optimal) Allocations for Binary Chores*. Theoretical Computer Science 1042, 2025, 115248.