

# Envy-Freeness with Limited Subsidies under Negative Dichotomous Valuations

*Working draft — not for circulation*

Mainak

**What this document is.** The complete state of the work, including results that are proved and machine-checked but not yet worth putting in the report. Sections 1 to 3 are shared verbatim with `main.tex`; everything after them is parked material. When a parked section earns its place, move its file from `working/` back to `sections/` and uncomment the matching line in `main.tex`. Nothing here is stale: every claim was re-verified against the scripts on the date of its section.

## 1 Introduction

Discrete fair division studies the allocation of resources that cannot be fractionally assigned among agents with individual preferences, and sits at the interface of mathematical economics and computer science [BCE<sup>+</sup>16, End17]. The classical fairness criterion is *envy-freeness* [Fol67, Var74]: no agent should prefer the bundle assigned to another agent over her own. For indivisible items an envy-free allocation need not exist, as the standard one-item two-agent instance shows, and a large body of work therefore studies relaxations such as envy-freeness up to one good (EF1) [Bud11, LMMS04].

A second and complementary line of work retains *exact* envy-freeness and buys off the residual envy with money. Each agent  $i$  receives, in addition to her bundle, a subsidy  $p_i \geq 0$  supplied by a third party, and one asks for an allocation together with a subsidy vector under which every agent weakly prefers her own bundle-plus-subsidy to that of every other agent. The question is how much money this requires. Maskin [Mas87] answered it for unit-demand agents with  $m = n$ : under the scaling in which no agent values any single item above 1, a total subsidy of  $n - 1$  suffices. Halpern and Shah [HS19] moved to the multi-demand setting with arbitrary  $m$ , gave the characterisation of *envy-freeable* allocations that the whole line now runs on, showed that a subsidy of  $(n - 1)m$  is necessary and sufficient in the worst case when the allocation is handed to us, and conjectured that  $n - 1$  suffices when we may choose the allocation. Brustle et al. [BDN<sup>+</sup>20] proved that conjecture in the stronger per-agent form — one dollar to each agent suffices for additive valuations, with the allocation simultaneously EF1 and balanced — and showed that for general monotone valuations  $2(n - 1)$  per agent suffices, so that the total subsidy is  $O(n^2)$  and, remarkably, *independent of the number of items*. Barman et al. [BKNS22] then improved the monotone bound by a factor of  $n$  on the dichotomous subclass: when every marginal

$v_i(S \cup \{g\}) - v_i(S)$  lies in  $\{0, 1\}$ , there is an allocation whose accompanying subsidy is exactly 0 or 1 for each agent, hence at most  $n - 1$  in total, computable in polynomial time in the value-oracle model. Their bound assumes neither additivity nor submodularity nor even subadditivity, and it is tight: with  $n$  agents and a single unit-valued good,  $n - 1$  agents must each be paid 1.

**Chores.** Every result above is about *goods*. This paper asks the mirrored question, in which every item is a *chore*: an object that an agent would rather not hold, so that acquiring it weakly decreases her utility. Chores are not an exotic variant. Shift rotations, on-call duty, teaching and refereeing load, committee service, maintenance tasks and the liabilities attached to an estate are all allocation problems in which the items are burdens rather than prizes, and in all of them money is already the standard instrument of repair — as overtime pay, hardship allowance, or a course-release budget. If anything the subsidy model is more natural for chores than for goods, since compensating somebody for taking on unwanted work needs less justification than paying somebody for the privilege of not receiving a prize.

Concretely, we study *negative dichotomous* valuations: for every agent  $i$ , every  $S \subseteq M$  and every  $g \in M \setminus S$ ,

$$v_i(S) - v_i(S \cup \{g\}) \in \{0, 1\}.$$

Equivalently, each marginal  $v_i(S \cup \{g\}) - v_i(S)$  is 0 or  $-1$ : every item is a chore for every agent, and taking on one more chore costs an agent either nothing or exactly one unit. This is the exact mirror of the model of [BKNS22], and we impose no further structure — in particular the valuations need not be additive, submodular, or subadditive.

Binary marginals are the natural model whenever a task either lands on an agent as a fresh unit of burden or is absorbed by a commitment she has already made. An engineer already staffing an on-call window absorbs a second incident in that window at no extra cost; a referee already reading for a venue absorbs one more paper from the same batch; a driver already routed through a district absorbs an extra stop along the way. Whether a given chore is free depends in an arbitrary way on the rest of the agent’s assignment, which is exactly why the model must tolerate non-additive and even non-subadditive valuations, and it is the counterpart of the scheduling example used by [BKNS22] to motivate dichotomous valuations for goods.

**Why this is not a change of sign.** It is tempting to expect that a chores result follows from the corresponding goods result by negating the valuations. It does not, for three reasons that we record here because they shape everything that follows.

First, money is one-signed. The subsidy vector is constrained to  $p \in \mathbb{R}_+^n$ , and negating an instance negates the valuations but leaves the subsidies where they are. The polyhedron of feasible subsidy vectors is therefore a genuinely different object in the two models, and no formal duality carries a bound across.

Second, the obvious transformation is available only for two agents. Writing  $c_i := -v_i$  for the cost form of a negative dichotomous valuation, the complement  $\tilde{v}_i(S) := c_i(M) - c_i(M \setminus S)$  is again dichotomous in the sense of [BKNS22], so the class is closed under complementation. But the complement transformation respects the allocation only when  $M \setminus A_j = A_i$ , that is,

only when  $n = 2$ ; see Remark 9. For  $n \geq 3$  there is no such reduction, and that is where the content of the problem lies.

Third, and most tellingly, envy flows the other way. Under goods, envy concentrates *into* the agent holding the valuable items: many agents envy one, and the tight instance of [BKNS22] pays  $n - 1$  agents one dollar each so that they stop envying the single agent who received the good. Under chores, envy emanates *out of* the agent carrying the load: one agent envies many. Since the minimum subsidy to an agent is the weight of the heaviest directed path leaving her in the envy graph [HS19, Theorem 2], the two situations are not interchangeable, and one should not assume in advance that the same bound is the right target.

**What is already known, and what is open.** Two facts delimit the problem before any work is done, and we state them explicitly so that the contribution can be measured against them.

On the one hand, existence of a bounded subsidy is not in question. Negative dichotomous valuations are doubly monotone — every item is a chore for every agent, which is a special case of every item being a good or a chore for each agent — and they satisfy the normalisation  $|v_i(S \cup \{e\}) - v_i(S)| \leq 1$  that [KMS<sup>+</sup>24, §2] imposes. The reduction of Kawase et al., which converts any EF1 allocation into an envy-free one, therefore applies off the shelf and yields a subsidy of at most  $n - 1$  per agent and  $n(n - 1)/2$  in total [KMS<sup>+</sup>24, Theorem 1]. Two details make the containment work and are easy to get wrong. Their normalisation is two-sided, so it bounds chore disutility as well as value; and the EF1 they require is the two-sided form of Theorem 5, not the goods-only form of [HS19] — the two differ precisely when chores are present. An EF1 allocation is guaranteed to exist and to be computable in polynomial time for doubly monotone valuations, so the reduction has an input; for that step [KMS<sup>+</sup>24, §2] cites [LMMS04] and [BSV21], and the citation matters, because the envy-cycle argument of [LMMS04] does *not* port to chores unaltered and it is [BSV21] that supplies the repair.

On the other hand, the lower bound of the goods model survives intact.

**Example 1** (The  $n - 1$  lower bound transfers). Take  $n$  agents and  $m = n - 1$  chores, with  $c_i(S) = |S|$  for every agent  $i$  — identical, binary, additive costs, which are in particular negative dichotomous. Every complete allocation leaves at least one agent empty-handed. Under the balanced allocation, in which  $n - 1$  agents hold one chore each and one agent holds nothing, each of the  $n - 1$  loaded agents envies the empty agent by exactly 1 and no other envy is present, so the minimum subsidy assigns 1 to each loaded agent and 0 to the empty one, for a total of  $n - 1$ . No allocation does better: any complete allocation gives some agent a bundle of size  $k \geq 1$  while some other agent is empty, and that agent’s heaviest outgoing path already has weight  $k$ .

So the target statement is precisely the analogue of [BKNS22, Theorem 4].

**Conjecture 2.** *For every instance with negative dichotomous valuations there is an envy-free solution  $(\mathcal{A}, p)$  with  $p \in \{0, 1\}^n$ , hence with total subsidy at most  $n - 1$ ; and it can be computed in polynomial time given value-oracle access.*

Example 1 shows Conjecture 2 would be tight. Between the  $n - 1$  per agent that [KMS<sup>+</sup>24] already gives and the  $\{0, 1\}$  per agent being asked for lies a factor of  $n$  in the total, exactly

the gap that [BKNS22] closed on the goods side. Our objective is to close it on the chores side or to exhibit an instance showing that it cannot be closed.

**Scope of this report.** Conjecture 2 is open. This report sets up the problem and stops there: Section 2 fixes the model and notation and records what transfers from the goods theory unchanged, and Section 3 isolates three structural facts about the chore problem — an arc-update lemma (Lemma 14) reducing the entire goods/chore asymmetry to two lines, a two-tier characterisation of  $\{0, 1\}$  solutions (Proposition 15), and a size-shift transform (Theorem 16) exhibiting the chore problem as exactly the goods problem of [BKNS22] together with a cardinality-balancing requirement. The last of these is the precise sense in which Conjecture 2 does not follow from [BKNS22] as a black box.

Work towards the conjecture is in progress and is deliberately not reported here. Several special cases are settled and two natural proof strategies have been ruled out by explicit machine-checked instances; none of it is stated below until it is worth stating. The remainder of this working draft, from Section 4 onwards, carries that material.

**Related Work.** The subsidy line for goods is described above. Within it, the dichotomous subclass has attracted separate attention: Halpern and Shah [HS19] settle binary additive valuations with a total subsidy of  $n - 1$ , Goko et al. [GIK<sup>+</sup>21] obtain the analogous bound for binary submodular valuations together with a strategy-proof mechanism, and Barman et al. [BKNS22] subsume both by dropping every structural assumption beyond binary marginals. Dichotomous preferences are themselves a well-studied restriction in social choice [BMS05, BEF21]. On the computational side, Caragiannis and Ioannidis [CI21] show that minimising the subsidy is NP-hard and give a fully polynomial-time approximation scheme for a constant number of agents, and Narayan et al. [NSV21] study the welfare cost of achieving envy-freeness with transfers.

The paper closest to the present setting is that of Kawase et al. [KMS<sup>+</sup>24], the only work in the envy-freeness-with-subsidy line that admits chores at all. Working with doubly monotone valuations under the two-sided normalisation, it decouples the two jobs that [BDN<sup>+</sup>20] performed together: rather than constructing a specific allocation and showing that it is cheap to subsidise, it takes *any* EF1 allocation as input and bounds the subsidy needed to convert it, obtaining  $n - 1$  per agent and  $n(n - 1)/2$  in total. Since EF1 allocations are known to exist for doubly monotone valuations, the wider class comes along at no cost. That reduction is what makes the present question well-posed rather than open-ended, and it is the baseline we are trying to beat.

The complementary repair — no money, weaker fairness notion — is pursued for binary valuations by Bu et al. [BSY23], who show that EFX allocations always exist and can be computed in polynomial time for general binary valuations, thus extending the matroid-rank result of Babaioff et al. [BEF21]. Under binary valuations one therefore has a clean dichotomy on the goods side: either pay 0 or 1 per agent and obtain exact envy-freeness [BKNS22], or pay nothing and obtain EFX [BSY23]. Whether either branch survives the passage to chores is open; this paper addresses the first.

**The chores side.** Work on indivisible chores has largely developed alongside rather than inside the subsidy line, and the two have met less often than one might expect. Where money has been brought to chores, it has been in the service of *proportionality* rather than envy-freeness: the question studied is the total subsidy needed to bring every agent down to her proportional share of the burden, for which the best bounds under additive disutilities normalised to at most 1 per chore are  $n/4$  and then  $n/3 - 1/6$ , recently obtained together with Pareto optimality by Garg et al. [GSW25]. That is a different fairness notion and a different accounting; it does not bound the subsidy needed for envy-freeness.

The no-money branch, by contrast, has a direct chore analogue of [BSY23]. Tao et al. [TWYZ23] study chores with binary marginals and show that for binary *additive* costs an allocation that is simultaneously EFX and Pareto optimal always exists and is computable in polynomial time. So on the chores side the second half of the binary dichotomy — pay nothing, obtain EFX — is established at least for additive costs, while the first half, pay 0 or 1 per agent and obtain exact envy-freeness, is what Conjecture 2 asks for. That money is genuinely needed here is not obvious a priori but is known: Bhaskar et al. [BSV21] show that deciding whether an exactly envy-free allocation of chores exists is NP-complete already for binary additive costs, and give a polynomial-time EF1 algorithm for chores and for doubly monotone instances.

Where the two lines have met is on the *additive* side, and recently. Lu et al. [LMS26] prove the goods-and-chores analogue of [BDN<sup>+</sup>20]: for additive utilities in which each item is a good for some agents and a chore for others, with every  $|v_i(g)| \leq 1$ , a subsidy of one dollar per agent suffices, the bound is tight, and the allocation is computable in polynomial time — hence  $n - 1$  in total. This settles the additive chore problem completely.

It does not settle ours, and the reason is exactly the reason [BKNS22] was not a corollary of [BDN<sup>+</sup>20]. Dichotomous valuations are not additive and additive valuations are not dichotomous; the two classes are incomparable. The four corners are

|             | goods                 | chores              |
|-------------|-----------------------|---------------------|
| additive    | [BDN <sup>+</sup> 20] | [LMS26]             |
| dichotomous | [BKNS22]              | <b>Conjecture 2</b> |

each of the three settled corners giving one dollar per agent and  $n - 1$  in total. The present question is the fourth. Note also that [LMS26] does not displace [KMS<sup>+</sup>24] as the baseline of Section 1: for valuations that are dichotomous but not additive,  $n - 1$  per agent from [KMS<sup>+</sup>24] remains the best published bound.

We are not aware of any treatment of envy-freeness with subsidy for chores under dichotomous costs. This is the outcome of a literature search rather than a proof of novelty — [LMS26] appeared in July 2026, which is a reminder of how quickly the surrounding corners are being filled in — and it should be repeated before the paper is circulated.

## 2 Notation and Preliminaries

We study the allocation of  $m$  indivisible items among  $n$  agents, with subsidies. Write  $N = [n]$  for the set of agents and  $M$  for the set of items,  $|M| = m$ . Each agent  $i \in N$  has a valuation

$v_i : 2^M \rightarrow \mathbb{R}$  with  $v_i(\emptyset) = 0$ , and we represent an instance by the tuple  $\langle N, M, \{v_i\}_{i \in N} \rangle$ . Algorithms operate in the standard *value-oracle* model: for each  $v_i$  we assume only an oracle that returns  $v_i(S)$  for a queried set  $S \subseteq M$ , and running time is measured in  $n, m$  and the number of oracle calls.

An allocation  $\mathcal{A} = (A_1, \dots, A_n)$  is an ordered tuple of pairwise disjoint *bundles*, with  $A_i$  assigned to agent  $i$ . The allocation is *complete* if  $\bigcup_{i \in N} A_i = M$  and *partial* otherwise. The ordering matters throughout: much of the theory below concerns permuting a fixed family of bundles among the agents, and for a permutation  $\sigma \in \mathbb{S}_n$  we write  $\mathcal{A}_\sigma := (A_{\sigma(1)}, \dots, A_{\sigma(n)})$  for the allocation in which agent  $i$  receives  $A_{\sigma(i)}$ . The utilitarian welfare of  $\mathcal{A}$  is  $\text{SW}(\mathcal{A}) = \sum_{i \in N} v_i(A_i)$ .

**Definition 3** (Envy-Freeness). *An allocation  $\mathcal{A} = (A_1, \dots, A_n)$  is envy-free (EF) iff  $v_i(A_i) \geq v_i(A_j)$  for all agents  $i, j \in N$ .*

Envy-free allocations of indivisible items need not exist, so we allow a third party to supplement the allocation with money. Agents have quasi-linear utilities: agent  $i$  receiving bundle  $A_i$  and subsidy  $p_i$  enjoys utility  $v_i(A_i) + p_i$ , with money entering linearly and at the same exchange rate for every agent.

**Definition 4** (Envy-Free Solution). *An allocation  $\mathcal{A} = (A_1, \dots, A_n)$  and a subsidy vector  $p = (p_1, \dots, p_n) \in \mathbb{R}_+^n$  constitute an envy-free solution  $(\mathcal{A}, p)$  iff  $v_i(A_i) + p_i \geq v_i(A_j) + p_j$  for all  $i, j \in N$ .*

An allocation  $\mathcal{A}$  is *envy-freeable* if some  $p \in \mathbb{R}_+^n$  makes  $(\mathcal{A}, p)$  an envy-free solution; this is a property of the allocation alone. Following [BKNS22] we write  $M(p) := \{i \in N \mid p_i \geq p_j \text{ for all } j \in N\}$  for the set of maximally subsidised agents.

We will occasionally need the standard relaxation of envy-freeness. Note that two inequivalent forms of it appear in this literature: the goods-only form of [HS19], which removes one item from the envied bundle, and the two-sided form used by [KMS<sup>+</sup>24], which removes one item from *either* bundle. They coincide when all items are goods and differ once chores are present, and only the two-sided form supports the doubly monotone results. We adopt the two-sided form.

**Definition 5** (EF1, two-sided form). *An allocation  $\mathcal{A}$  is envy-free up to one item (EF1) iff for all  $i, j \in N$  there exists  $X \subseteq A_i \cup A_j$  with  $|X| \leq 1$  such that  $v_i(A_i \setminus X) \geq v_i(A_j \setminus X)$ .*

## 2.1 Negative Dichotomous Valuations

**Definition 6** (Negative Dichotomous Valuation). *A valuation  $v_i : 2^M \rightarrow \mathbb{R}$  with  $v_i(\emptyset) = 0$  is negative dichotomous iff*

$$v_i(S) - v_i(S \cup \{g\}) \in \{0, 1\} \quad \text{for every } S \subseteq M \text{ and every } g \in M \setminus S,$$

*equivalently iff every marginal  $\Delta_i^+(S, g) := v_i(S \cup \{g\}) - v_i(S)$  lies in  $\{-1, 0\}$ .*

Every item is thus a chore for every agent:  $v_i$  is non-increasing, so  $v_i(S) \geq v_i(T)$  whenever  $S \subseteq T$ . We stress that nothing further is assumed — the valuations need not be additive,



submodular, or subadditive — which is the generality at which [BKNS22] works on the goods side.

It is often more convenient to work with the non-negative *cost form*

$$c_i := -v_i, \quad \text{so} \quad c_i(\emptyset) = 0, \quad c_i(S) \leq c_i(T) \text{ for } S \subseteq T, \quad c_i(S \cup \{g\}) - c_i(S) \in \{0, 1\}.$$

**Definition 7** (Dichotomous cost function). *A set function  $c : 2^M \rightarrow \mathbb{R}$  is dichotomous if  $c(\emptyset) = 0$  and every marginal  $c(S \cup \{g\}) - c(S)$  lies in  $\{0, 1\}$ .*

Monotonicity is not a separate hypothesis: marginals in  $\{0, 1\}$  are non-negative, so a dichotomous  $c$  is automatically non-decreasing. We nonetheless list it above because it is the property most often used directly.

**Remark 8** (Three adjectives that are routinely confused). *Monotone* means only  $c(S) \leq c(T)$  for  $S \subseteq T$ , with no bound whatever on how large an increment can be; it is far weaker than Theorem 7 and, on its own, supports almost none of the arguments below. *Dichotomous* adds the binary-marginal condition and is the class of this paper. *Binary additive* is the strictly smaller class  $c_i(S) = |S \cap D_i|$  for a set  $D_i$  of disliked chores, i.e. additive with per-item costs in  $\{0, 1\}$ ; it is the case settled by [LMS26].

A warning about the literature: several papers say “binary valuations” for what is called dichotomous here — [BSY23] and [TWYZ23] both do — and reserve “binary additive” for the additive subclass. When a bound is quoted, check which is meant, since the classes are not the same and the gap between them is where this paper lives.

The correspondence  $v_i \leftrightarrow c_i$  is a bijection between negative dichotomous valuations and dichotomous cost functions in the sense of [BKNS22]: the class we study is exactly the pointwise negation of the class they study. In cost form an envy-free solution is the requirement  $c_i(A_i) - p_i \leq c_i(A_j) - p_j$  for all  $i, j$ , read as “my own burden net of my compensation is no worse than what I perceive of yours”. We use whichever of  $v_i$  and  $c_i$  is clearer locally and always state which.

The negation, however, does not carry results across, for the reasons given in Section 1. The following remark isolates the one case in which a transformation does exist, and shows why it is uninformative.

**Remark 9** (Complementation, and why it stops at  $n = 2$ ). For a negative dichotomous  $v_i$  with cost form  $c_i$ , define the *complement*  $\tilde{v}_i(S) := c_i(M) - c_i(M \setminus S)$ . Then  $\tilde{v}_i(\emptyset) = 0$  and, for  $g \notin S$ , writing  $T := M \setminus (S \cup \{g\})$ ,

$$\tilde{v}_i(S \cup \{g\}) - \tilde{v}_i(S) = c_i(M \setminus S) - c_i(T) = c_i(T \cup \{g\}) - c_i(T) \in \{0, 1\},$$

so  $\tilde{v}_i$  is dichotomous in the sense of [BKNS22]. The class is therefore closed under complementation, without any submodularity assumption. Now let  $n = 2$  and let  $\mathcal{A} = (A_1, A_2)$  be complete, so that  $M \setminus A_2 = A_1$  and  $M \setminus A_1 = A_2$ . Then  $-c_i(A_i) = \tilde{v}_i(A_j) - c_i(M)$  and  $-c_i(A_j) = \tilde{v}_i(A_i) - c_i(M)$ , so  $(\mathcal{A}, p)$  is an envy-free solution for the chores instance  $\{v_i\}$  if and only if  $((A_2, A_1), p)$  is an envy-free solution for the goods instance  $\{\tilde{v}_i\}$ : swap the bundles and leave each subsidy attached to its agent. Two-agent chore division is thus exactly two-agent goods division, and [BKNS22] settles it. The step  $M \setminus A_j = A_i$  has no analogue for  $n \geq 3$ , where the complement of one bundle is a union of several others, and no reduction of this kind is available.

## 2.2 The Envy Graph and the Halpern–Shah Characterisation

For an allocation  $\mathcal{A}$ , the *envy graph*  $G_{\mathcal{A}}$  is the complete weighted directed graph on vertex set  $N$  in which the arc  $(i, j)$  carries weight

$$w_{\mathcal{A}}(i, j) := v_i(A_j) - v_i(A_i) = c_i(A_i) - c_i(A_j),$$

the envy that agent  $i$  has towards agent  $j$ . The weight of a directed path  $P$  is  $w_{\mathcal{A}}(P) = \sum_{(i,j) \in P} w_{\mathcal{A}}(i, j)$ , and  $\ell_{\mathcal{A}}(i)$  denotes the maximum weight of any path in  $G_{\mathcal{A}}$  starting at  $i$ , the empty path of weight 0 included.

The two results below are the engine of the entire subsidy line. Both are stated by Halpern and Shah [HS19] for additive valuations, but their proofs use only the arc weights and permutations of the bundles — neither additivity nor monotonicity nor any sign condition on  $v_i$  enters — and [BKNS22] invokes them in exactly this generality for (non-additive) dichotomous valuations. They therefore apply verbatim to negative dichotomous valuations, and we use them without further comment.

**Theorem 10** ([HS19]). *For any allocation  $\mathcal{A} = (A_1, \dots, A_n)$ , the following are equivalent.*

- (i)  $\mathcal{A}$  is *envy-freeable*.
- (ii)  $\mathcal{A}$  *maximises utilitarian welfare across all reassignments of its own bundles among the agents: for every permutation  $\sigma$  over  $N$ ,  $\sum_{i \in N} v_i(A_i) \geq \sum_{i \in N} v_i(A_{\sigma(i)})$ .*
- (iii) *The envy graph  $G_{\mathcal{A}}$  has no positive-weight directed cycle.*

**Theorem 11** ([HS19]). *For any envy-freeable allocation  $\mathcal{A}$ , the subsidy vector  $p^*$  given by  $p_i^* := \ell_{\mathcal{A}}(i)$  realises an envy-free solution  $(\mathcal{A}, p^*)$ , and  $p_i^* \leq p_i$  for every  $i \in N$  and every  $p$  such that  $(\mathcal{A}, p)$  is an envy-free solution. It can be computed in strongly polynomial time by running an all-pairs longest-path computation on  $G_{\mathcal{A}}$ .*

Condition (ii) of Theorem 10 has the standing consequence that an envy-freeable allocation can always be manufactured from an arbitrary one: fix the bundles and reassign them according to a maximum-weight perfect matching between agents and bundles. What is at issue is never the existence of an envy-freeable allocation, only the size of the subsidy that Theorem 11 then charges for it.

Specialising to our model gives the following, which we use throughout.

**Observation 12** (Integrality). *Let  $\{v_i\}_{i \in N}$  be negative dichotomous. Then  $v_i(S) \in \mathbb{Z}_{\leq 0}$  for every  $i$  and every  $S \subseteq M$ ; every arc weight  $w_{\mathcal{A}}(i, j)$  is an integer; and for every envy-freeable allocation  $\mathcal{A}$  the minimum subsidy vector satisfies  $p^* \in \mathbb{Z}_{\geq 0}^n$ .*

**Proof** Fix  $i$  and  $S = \{g_1, \dots, g_k\}$ . Telescoping from  $v_i(\emptyset) = 0$  along any ordering of  $S$  writes  $v_i(S)$  as a sum of  $k$  marginals, each of which lies in  $\{-1, 0\}$  by Definition 6; hence  $v_i(S) \in \mathbb{Z}$  and  $-k \leq v_i(S) \leq 0$ . Consequently each arc weight  $w_{\mathcal{A}}(i, j) = v_i(A_j) - v_i(A_i)$  is a difference of integers, and the weight of any directed path, being a finite sum of arc weights, is an integer. By Theorem 11,  $p_i^* = \ell_{\mathcal{A}}(i)$  is the maximum such weight over paths starting at  $i$ , and it is non-negative because the empty path has weight 0.  $\square$



Observation 12 is what makes the target statement — a subsidy vector in  $\{0, 1\}^n$  — a statement about a *bound* rather than about rounding: once  $p^*$  is known to be integral, showing  $p_i^* \leq 1$  for every  $i$  is the same as showing  $p^* \in \{0, 1\}^n$ . The same observation underlies the corresponding step in [BKNS22], where the arc weights of a dichotomous instance are likewise integers.

A second normalisation costs nothing and removes a clause from the target.

**Proposition 13** (Some agent is always unpaid). *For every envy-freeable allocation  $\mathcal{A}$  there is an agent  $h$  with  $p_h^* = 0$ . Consequently, if  $p^* \in \{0, 1\}^n$  then  $\sum_{i \in N} p_i^* \leq n - 1$ .*

**Proof** By Theorem 10(iii) the envy graph  $G_{\mathcal{A}}$  has no positive-weight cycle, so a walk of maximum weight may be taken simple and  $\ell_{\mathcal{A}}$  is finite; it is non-negative because the empty path qualifies. Choose a path  $P$  attaining  $\max_i \ell_{\mathcal{A}}(i)$  and let  $h$  be its final vertex. If  $\ell_{\mathcal{A}}(h) > 0$  there is a positive-weight walk  $Q$  leaving  $h$ , and  $P$  followed by  $Q$  is a walk from  $P$ 's start of weight  $w_{\mathcal{A}}(P) + w_{\mathcal{A}}(Q) > w_{\mathcal{A}}(P)$ , contradicting the maximality of  $P$  among walks. Hence  $\ell_{\mathcal{A}}(h) = 0$ . The second claim is immediate: at most  $n - 1$  coordinates of  $p^*$  can equal 1.  $\square$

Proposition 13 justifies the word “hence” in Conjecture 2: the bound on the *total* is not an independent requirement but a consequence of the per-agent one. The conjecture is therefore a single-clause statement,

$$\exists \mathcal{A} : \max_{i \in N} \ell_{\mathcal{A}}(i) \leq 1,$$

and it is this form that the rest of the report works with. Two cautions are worth recording alongside it. First, the bound cannot be improved: the family of Example 1 attains  $n - 1$  for every  $n$ , verified exhaustively for  $n \leq 7$ . Second, the conjecture is *not* the assertion that the utilitarian-optimal subsidy is at most  $n - 1$ . Minimising the total over all allocations can do strictly better than any allocation with  $p^* \in \{0, 1\}^n$ , by concentrating the subsidy on a single agent: there are instances at  $n = 4$  whose cheapest envy-free solution costs 2, borne entirely by one agent, while every allocation with  $p^* \in \{0, 1\}^4$  costs 3. The  $\{0, 1\}$  shape is a genuine restriction, not a normalisation of the total.

### 3 Structure of the Chore Problem

Throughout this section we work in cost form, so  $w_{\mathcal{A}}(i, j) = c_i(A_i) - c_i(A_j)$ .

#### 3.1 The arc-update table, and the whole of the asymmetry

Every incremental algorithm in this literature adds one item to one bundle and then asks what happened to the envy graph. The answer is two lines long, and comparing those two lines against their goods counterparts accounts for the entire difference between the present problem and that of [BKNS22].

For an allocation  $\mathcal{A}$ , an agent  $x$  and an unallocated chore  $g \notin \bigcup_i A_i$ , write

$$Z^x := (A_1, \dots, A_x \cup \{g\}, \dots, A_n)$$

for the allocation obtained by giving  $g$  to  $x$ , and put

$$\delta_x := c_x(A_x \cup \{g\}) - c_x(A_x), \quad \delta_{i,x} := c_i(A_x \cup \{g\}) - c_i(A_x),$$

both of which lie in  $\{0, 1\}$ . So  $\delta_x$  is what the chore costs  $x$  herself, and  $\delta_{i,x}$  is what it costs in  $i$ 's opinion.

**Lemma 14** (Arc update). *For all  $i, j \in N$ ,*

$$w_{Z^x}(i, j) = w_{\mathcal{A}}(i, j) \quad (i, j \neq x), \quad w_{Z^x}(x, j) = w_{\mathcal{A}}(x, j) + \delta_x, \quad w_{Z^x}(i, x) = w_{\mathcal{A}}(i, x) - \delta_{i,x}.$$

**Proof** Only  $A_x$  changes, so an arc between two agents other than  $x$  has both endpoints untouched. For an arc leaving  $x$ ,  $w_{Z^x}(x, j) = c_x(A_x \cup \{g\}) - c_x(A_j) = w_{\mathcal{A}}(x, j) + \delta_x$ . For an arc entering  $x$ ,  $w_{Z^x}(i, x) = c_i(A_i) - c_i(A_x \cup \{g\}) = w_{\mathcal{A}}(i, x) - \delta_{i,x}$ .  $\square$

The mirror statement for goods is obtained by replacing costs with values: there the arcs *out of* the receiving agent fall and the arcs *into* her rise. The consequences are opposite in a way that matters.

|                          | goods [BKNS22]              | chores (here)               |
|--------------------------|-----------------------------|-----------------------------|
| arcs out of the receiver | fall                        | <b>rise by</b> $\delta_x$   |
| arcs into the receiver   | <b>rise</b>                 | fall by $\delta_{i,x}$      |
| so the receiver must be  | <i>maximally</i> subsidised | <i>minimally</i> subsidised |
| which is guaranteed by   | $M(p) \neq \emptyset$       | $p_x = 0$ for some $x$      |

Both guarantees hold trivially — some agent is always maximally subsidised, and some agent always has  $p_x = 0$  under the minimum subsidy vector. The asymmetry is subtler than the table suggests, and it is what Theorem 31 exploits. Under goods, the harm done by an insertion is confined to paths *ending* at the receiver, and requiring  $x \in M(p)$  caps every such path at  $p_u - p_x \leq 0$  in one stroke. Under chores, the harm is done to paths *leaving* the receiver, and  $p_x = 0$  caps only the paths that *start* at  $x$ . A path that merely passes through  $x$  — entering along an arc from some  $i$  with  $\delta_{i,x} = 0$  and leaving along an arc that has just risen by  $\delta_x = 1$  — is not controlled by any condition on  $p_x$ . That gap is not a defect of the analysis; Section 6 shows it is realised.

### 3.2 The shape of a $\{0, 1\}$ solution

**Proposition 15** (Two-tier characterisation). *Let  $p \in \{0, 1\}^n$  and put  $S := \{i \in N \mid p_i = 1\}$ . Then  $(\mathcal{A}, p)$  is an envy-free solution if and only if*

- (a)  $c_i(A_i) \leq c_i(A_j)$  whenever  $i, j$  lie both in  $S$  or both in  $N \setminus S$ ;
- (b)  $c_i(A_i) \leq c_i(A_j) + 1$  for  $i \in S$  and  $j \in N \setminus S$ ;
- (c)  $c_i(A_i) + 1 \leq c_i(A_j)$  for  $i \in N \setminus S$  and  $j \in S$ .

**Proof** In cost form Definition 4 reads  $c_i(A_i) \leq c_i(A_j) + p_i - p_j$  for all  $i, j$ . Substituting the three possible values 0, 1,  $-1$  of  $p_i - p_j$  gives (a), (b) and (c) respectively.  $\square$

The reading is worth keeping in mind when designing an algorithm. The unpaid agents are exactly envy-free among themselves and *strictly* prefer their own bundle to that of every paid agent, by a full unit; the paid agents are exactly envy-free among themselves and envy-free up to one unit towards the unpaid. This is a certificate format, and it also suggests choosing the partition  $S \mid N \setminus S$  *first* and allocating afterwards, rather than building the allocation up one chore at a time. That is the one place where the incremental template killed in Section 6 is not obviously an obstacle. In the proof of Theorem 19 the set  $S$  turns out to be exactly the agents holding a  $\lceil q/n \rceil$ -share of the universally disliked chores.

### 3.3 The size-shift transform

Remark 9 showed that complementation relates the chore and goods models only for  $n = 2$ . There is a second transform, available for every  $n$ , which does not solve the problem but says precisely how far it is from being a corollary of [BKNS22].

**Theorem 16** (Size shift). *For a negative dichotomous instance with cost functions  $c_i$ , define  $\tilde{v}_i(S) := |S| - c_i(S)$ . Then each  $\tilde{v}_i$  is a dichotomous goods valuation in the sense of [BKNS22], and for every allocation  $\mathcal{A}$  and all  $i, j$ ,*

$$w_{\mathcal{A}}(i, j) = \tilde{w}_{\mathcal{A}}(i, j) + |A_i| - |A_j|,$$

where  $\tilde{w}_{\mathcal{A}}$  denotes the envy graph of the goods instance  $\{\tilde{v}_i\}$ . Consequently every cycle has the same weight in both graphs — so  $\mathcal{A}$  is envy-freeable for the chores  $\{c_i\}$  if and only if it is envy-freeable for the goods  $\{\tilde{v}_i\}$  — while for a path  $P$  from  $u$  to  $v$ ,

$$w_{\mathcal{A}}(P) = \tilde{w}_{\mathcal{A}}(P) + |A_u| - |A_v|, \quad \text{hence} \quad \ell_{\mathcal{A}}(u) = \max_{v \in N} \left[ \tilde{d}_{\mathcal{A}}(u, v) + |A_u| - |A_v| \right],$$

with  $\tilde{d}_{\mathcal{A}}(u, v)$  the maximum weight of a path from  $u$  to  $v$  in the goods graph and  $\tilde{d}_{\mathcal{A}}(u, u) = 0$ .

**Proof**  $\tilde{v}_i(\emptyset) = 0$ , and for  $g \notin S$ ,  $\tilde{v}_i(S \cup \{g\}) - \tilde{v}_i(S) = 1 - (c_i(S \cup \{g\}) - c_i(S)) \in \{0, 1\}$ , so  $\tilde{v}_i$  is dichotomous. For the identity,

$$\tilde{w}_{\mathcal{A}}(i, j) = \tilde{v}_i(A_j) - \tilde{v}_i(A_i) = (|A_j| - c_i(A_j)) - (|A_i| - c_i(A_i)) = w_{\mathcal{A}}(i, j) - |A_i| + |A_j|.$$

Summing along a path  $(i_1, \dots, i_r)$ , the terms  $|A_{i_t}| - |A_{i_{t+1}}|$  telescope to  $|A_{i_1}| - |A_{i_r}|$ ; along a cycle they telescope to 0. The characterisation of envy-freeability by the absence of positive cycles (Theorem 10) then transfers, and the formula for  $\ell_{\mathcal{A}}(u)$  follows from Theorem 11.  $\square$

**Remark 17** (Why Conjecture 2 is not a corollary of [BKNS22]). Theorem 16 says the chore problem is the goods problem of [BKNS22] *plus a cardinality-balancing requirement*. A sufficient condition falls straight out: since  $\tilde{d}_{\mathcal{A}}(u, v) \leq \tilde{p}_u - \tilde{p}_v$  for any envy-eliminating  $\tilde{p}$ , an allocation solving the goods instance in which the quantity  $|A_i| + \tilde{p}_i$  has spread at most 1 across agents satisfies  $\ell_{\mathcal{A}} \leq 1$ . But [BKNS22] offers no control whatever on bundle cardinalities — its output need not even be EF1, by its own Appendix C — so the requirement is not met by invoking it as a black box. Nor can its algorithm simply be re-run with a balancing tie-break, because Theorem 31 rules out the entire item-by-item template on which it is built. The transform also shows the chore problem is at least as expressive as the goods problem, so no easier.

## 4 Our Results

The target is Conjecture 2, stated in Section 1: an envy-free solution with  $p \in \{0, 1\}^n$  and hence total subsidy at most  $n - 1$ . Example 1 supplies the matching lower bound, so it would be tight. By Observation 12 the conjecture is equivalent to the purely graph-theoretic statement that some allocation  $\mathcal{A}$  has no positive-weight cycle in  $G_{\mathcal{A}}$  and satisfies  $\ell_{\mathcal{A}}(i) \leq 1$  for every agent  $i$ .

We do not settle Conjecture 2. What we contribute is a map of the problem: two solved special cases, three structural tools, and — the substance of the paper — two theorems showing that the two most natural proof strategies both fail, one of them the strategy that succeeds on the goods side.

**Solved cases.** Conjecture 2 holds, constructively and in polynomial time, for **binary additive** costs (Theorem 19) and for **identical** costs (Theorem 21). The case  $n = 2$  follows from [BKNS22] through the complementation of Remark 9. The binary additive proof is the one to imitate: dump every chore on somebody who finds it free, then balance whatever nobody wants. Its bound telescopes along paths in the envy graph, which is the mechanism all three tools below try to reproduce.

**Structural tools** (Section 3). The arc-update table of Lemma 14 isolates the entire difference between the goods and chore problems in two lines. Proposition 15 characterises  $\{0, 1\}$  solutions as a two-tier structure, which is the natural certificate format and suggests choosing the tiers before allocating. Theorem 16 exhibits the chore problem as *exactly* the goods problem of [BKNS22] plus a cardinality-balancing requirement, which is the precise sense in which Conjecture 2 is not a corollary of [BKNS22].

**First negative result** (Section 6). Item-by-item insertion — allocate one chore at a time, never move an already-allocated chore, and repair only by permuting whole bundles — is the template of [BKNS22]. Theorem 31 exhibits a three-agent instance on which that template provably cannot be completed: a legal partial state with  $p \in \{0, 1\}^3$  and an unallocated chore such that *every* agent who receives it forces a subsidy of 2, while the full instance admits a zero-subsidy allocation. Any proof of Conjecture 2 must therefore be allowed to re-open bundles that have already been allocated. We state this bluntly because the template is seductive and patching the insertion invariant is a dead end, not an unsolved exercise.

**Second negative result** (Section 7). The natural repair is to stop building the allocation incrementally and instead select a good allocation globally, the obvious candidate being one that minimises total cost. This suggests the following strengthening.

**Conjecture 18.** *The allocation in Conjecture 2 may in addition be taken utilitarian-optimal, i.e. minimising  $\sum_{i \in N} c_i(A_i)$  over all allocations.*

Proposition 34 refutes Conjecture 18: there is a three-agent, four-chore instance whose *unique* utilitarian-optimal allocation requires a subsidy of 2, although another allocation — of strictly higher total cost — requires no subsidy at all. Utilitarian optimality is therefore not

merely an unhelpful restriction but an actively wrong one, and the programme of searching for a tie-breaking rule inside the utilitarian-optimal set is unsound as posed.

**A third approach, still open** (Section 8). Both failures above are failures of *timing*: insertion commits a chore to an owner and cannot take it back, and utilitarian optimality commits to the wrong allocation outright. Section 8 changes coordinates so that ownership is decided only by the last move that touches a chore. Theorem 39 re-expresses the chore instance as a *goods* instance on an item set carrying  $n - 1$  copies of each chore, with identical envy graphs, so the dichotomous class is closed under the duality and [BKNS22] applies to the replica verbatim. Corollary 41 then reduces Conjecture 2 to a single residual constraint — *coverage*, that no agent receive two copies of one type. Read dynamically this is the *peel* process, in which the orientation of [BKNS22] is restored because the quantity handed out is relief rather than burden, and the witness of Theorem 31 is handled with the invariant intact. The approach is not known to work: Proposition 48 exhibits legal states from which no legal continuation exists, so what remains is to find a peel rule that never enters one.

**A third negative result** (Section 9). Given Corollary 41, the natural move is to change the numbers so that a coverage violation becomes unattractive, and thereby reuse [BKNS22] as a black box. Section 9 closes that at both levels. No dichotomous reweighting of the replica instance separates coverage from non-coverage (Theorems 59 and 60, the latter exhaustive over all  $990^3$  candidate triples), and no inflation of item weights can steer the algorithm of [BKNS22], because its repair subroutine provably runs where every candidate’s marginal is 0 and a duplicate moves no arc at all (Theorem 62, Corollary 63). Both failures have one cause, named in Section 9.6: a rule fixed before the allocation is chosen cannot act on a quantity the allocation determines. Coverage must be enforced by the procedure that builds the allocation, not by the objective it is scored against.

**A working algorithm** (Section 11). The size-shift transform of Theorem 16 is an involution (Lemma 76), so the chore problem restates as a pure *goods* question, Target G, with no chores or scheduling in it (Proposition 78). Running [BKNS22]’s own algorithm on that goods instance fails for a structural reason — item-by-item insertion can lock two items into the same bundle permanently, verified by an exhaustive backtracking search that finds a 0-spread allocation *unreachable* by any legal execution of the algorithm (Section 11.3). Abandoning insertion for a construct-and-repair procedure — one pass of iterated maximum-weight perfect matching, repaired by cardinality-preserving single-item swaps, restarted on failure — has **zero failures across 389,215 test instances**, including every hard instance that defeated Approaches 1 through 5.

Four properties of that algorithm are now *theorems*: its outputs are certified (Theorem 82), it terminates in polynomial time  $O(n^3 m^3 \tau + n^4 m^3)$  (Theorem 83), its search space is genuinely the unordered balanced partitions (Lemma 81), and — the main theoretical result of this line — **it is complete for two agents** (Theorem 84): every two-agent dichotomous instance admits a balanced split of spread at most 1, found constructively in  $O(m\tau)$  time with no restarts, by a discrete intermediate-value argument on the aggregate  $u = \tilde{v}_1 + \tilde{v}_2$  along swap paths. This strengthens even the known  $n = 2$  chore result, adding balancedness that the

complementation route does not provide. For  $n \geq 3$ , completeness remains the one open statement (Section 11.8).

**Evidence** (Section 14). Conjecture 2 survives exhaustive verification for  $n = m = 3$  over all 9,880 instances up to agent symmetry, and randomised search over some  $3.5 \times 10^4$  further instances. Section 14 also explains why that evidence is weaker than the counts suggest, and what would strengthen it.

## 5 Solved Cases

Conjecture 2 holds in three restrictions of the model, by explicit polynomial-time constructions, and on one further class of instances cut out by a certificate rather than by a restriction on the valuations (Section 5.4). The first three proofs run the same way: bound  $w_{\mathcal{A}}(i, j)$  by a difference  $\phi_i - \phi_j$  of a per-agent potential, so that the bound telescopes to  $\phi_{i_1} - \phi_{i_r}$  along a path and to 0 around a cycle. Envy-freeability and the  $\{0, 1\}$  bound then both follow from controlling the spread of  $\phi$ . Reproducing that mechanism in general is the open problem.

**What is and is not new here.** This section should be read with its overlap against the literature stated plainly, because two of the three cases are largely or wholly not ours.

- **Binary additive costs** (Theorem 19) are *subsumed* by [LMS26], which settles all additive goods-and-chores instances. See Remark 20. The statement is not new; only the potential-and-telescoping mechanism is retained, because that is what the general case has to reproduce.
- **Identical costs** (Theorem 21) are *not* subsumed, for the reason given in Remark 22 — but the result is elementary and is best read as a warm-up.
- **Two agents** is [BKNS22, Theorem 4]. Only the reduction carrying it to chores (Remark 9) is ours.
- **The uniformly balanced class** (Section 5.4) is new, and is the only case here not implied by a cited paper. It is a certificate rather than an algorithm, and it is incomplete.

Nothing in this section, in other words, is load-bearing. The substance of the paper is in Sections 6, 7, 9 and 10, which are negative, and in Sections 3 and 8, which are structural.

### 5.1 Binary additive costs

Let  $c_i(S) = |S \cap D_i|$  for a set  $D_i \subseteq M$  of chores that agent  $i$  dislikes; every other chore is free for her. Call a chore *universally bad* if it lies in  $D_i$  for every  $i$ , write  $U$  for the set of universally bad chores and  $q := |U|$ .

**Theorem 19.** *For binary additive costs, the following allocation is envy-freeable with minimum subsidy  $p^* \in \{0, 1\}^n$ , hence total subsidy at most  $n - 1$ , and it is computable in time  $O(nm)$ .*



Give each chore  $g \notin U$  to some agent  $i$  with  $g \notin D_i$ ; at least one exists, by definition of  $U$ .  
Distribute  $U$  as evenly as possible, so that  $u_i := |A_i \cap U| \in \{\lfloor q/n \rfloor, \lceil q/n \rceil\}$  for every  $i$ .

The bound is tight, by Example 1.

**Proof** By construction every chore in  $A_i \setminus U$  is free for  $i$ , so  $c_i(A_i) = |A_i \cap D_i| = u_i$ . For any  $j$  we have  $A_j \cap U \subseteq A_j \cap D_i$ , since  $U \subseteq D_i$ , and therefore  $c_i(A_j) = |A_j \cap D_i| \geq u_j$ . Hence

$$w_{\mathcal{A}}(i, j) = c_i(A_i) - c_i(A_j) \leq u_i - u_j.$$

Summing along a path  $P = (i_1, \dots, i_r)$  the right-hand side telescopes:

$$w_{\mathcal{A}}(P) \leq u_{i_1} - u_{i_r} \leq \lceil q/n \rceil - \lfloor q/n \rfloor \leq 1.$$

Around a cycle the same computation gives  $w_{\mathcal{A}}(C) \leq 0$ , so  $G_{\mathcal{A}}$  has no positive-weight cycle and  $\mathcal{A}$  is envy-freeable by Theorem 10. By Theorem 11,  $p_i^* = \ell_{\mathcal{A}}(i) \leq 1$ , and Observation 12 forces  $p^* \in \{0, 1\}^n$ . Since some agent receives  $\lfloor q/n \rfloor$  universally bad chores and hence requires no subsidy, the total is at most  $n - 1$ . Example 1 is binary additive with  $D_i = M$  for every  $i$ , so the bound cannot be improved.  $\square$

**Remark 20** (Superseded, and by how much). Theorem 19 is subsumed by [LMS26], which appeared in July 2026 and gives one dollar per agent for *all* additive goods-and-chores instances with  $|v_i(g)| \leq 1$ ; binary additive costs are the special case  $v_i(g) \in \{-1, 0\}$ . We keep the proof because the mechanism, not the statement, is what the general case needs: the potential  $\phi_i = u_i$  and its telescoping along paths is the only device in this paper that produces a  $\{0, 1\}$  bound from scratch, and the general conjecture is precisely the problem of finding the right  $\phi$ . Theorem 21 and the  $n = 2$  case are *not* subsumed, since dichotomous costs need not be additive.

The construction is a two-phase version of the shape the general case must reproduce: *dump every chore on somebody who finds it free, then balance whatever nobody wants*. The obstruction to running the same argument for general negative dichotomous costs is that “free for  $i$ ” becomes context dependent. A chore may be free on top of  $A_j$  and cost a full unit on top of  $A_i$ , so there is no fixed set  $D_i$  to compute with and no fixed potential  $u_i$  to telescope against. Theorem 31 shows that this is a genuine obstacle rather than a bookkeeping nuisance.

## 5.2 Identical costs, and two agents

**Theorem 21.** Suppose  $c_1 = \dots = c_n = c$ . Then every allocation is envy-freeable, and greedily adding each chore to a currently cheapest bundle yields  $p^* \in \{0, 1\}^n$  in time  $O(m(n + \tau))$ , where  $\tau$  is the cost of one oracle call.

**Proof** For any permutation  $\sigma$  we have  $\sum_i c(A_{\sigma(i)}) = \sum_i c(A_i)$ , since both sides sum the same function over the same multiset of bundles. So condition (ii) of Theorem 10 holds with equality for every allocation, and every allocation is envy-freeable. Moreover  $w_{\mathcal{A}}(i, j) = c(A_i) - c(A_j)$  telescopes *exactly* along a path  $P = (i_1, \dots, i_r)$ , giving  $w_{\mathcal{A}}(P) = c(A_{i_1}) - c(A_{i_r})$  and hence

$$\ell_{\mathcal{A}}(u) = c(A_u) - \min_{j \in N} c(A_j).$$

It therefore suffices to keep  $\max_i c(A_i) - \min_i c(A_i) \leq 1$ . This is an invariant of the greedy rule. It holds vacuously at the empty allocation. Suppose it holds before a step, so that  $c(A_i) \in \{\mu, \mu + 1\}$  for every  $i$ , where  $\mu$  is the current minimum, and let the next chore go to a bundle attaining  $\mu$ . Its cost becomes  $\mu$  or  $\mu + 1$ , because marginals lie in  $\{0, 1\}$ ; every other bundle is unchanged. So all costs still lie in  $\{\mu, \mu + 1\}$  and the invariant survives. The bound  $p^* \in \{0, 1\}^n$  follows with Observation 12.  $\square$

**Remark 22** (Why this one survives the literature). Theorem 21 is not covered by [LMS26], because identical dichotomous costs need not be additive. The instance  $c_i(S) = \max(0, |S| - 1)$  for every  $i$  is identical and dichotomous, with marginals  $0, 1, 1, \dots$ , and is supermodular rather than additive — it is the very cost function that drives the obstruction of Theorem 31. Nor does [BKNS22] cover it, that being a goods result, and [KMS<sup>+</sup>24] gives only  $n - 1$  per agent on this class. So the statement is new, in the narrow sense that no cited paper implies it.

It is also easy, and should not be oversold. With identical costs every agent agrees on the cost of every bundle, the problem degenerates to load balancing, and the potential  $\phi_i = c(A_i)$  telescopes *exactly* rather than merely bounding the arcs. That exactness is what makes the proof three lines long, and it is precisely what fails as soon as two agents disagree.

For  $n = 2$  the conjecture holds as well, with total subsidy at most 1, but here the theorem is somebody else's: by Remark 9 two-agent chore division *is* two-agent goods division with the bundles swapped, so the content is [BKNS22, Theorem 4] and ours is only the reduction. As noted there, that reduction consumes the identity  $M \setminus A_j = A_i$  and so says nothing for  $n \geq 3$ .

### 5.3 Small bundles, and fewer chores than agents

Every other case in this section, and every approach in the paper, eventually needs an *exchange* step: given a bad allocation, produce a better one by moving items. That is exactly what dichotomous costs refuse to supply, since a bundle of positive cost need not contain any single item whose removal lowers it — for instance  $c(S) = \min(\max(0, |S| - 1), 1)$  on a three-element set has  $c = 1$  while every single removal leaves it at 1. The following theorem needs no exchange at all; it reads the conclusion straight off the cycle-closing bound.

**Theorem 23** (Small-bundle theorem). *Let  $\mathcal{A}$  be envy-freeable and suppose every bundle costs at most one unit to every agent:  $c_v(A_i) \leq 1$  for all  $v, i \in N$ . Then  $p^* \in \{0, 1\}^n$  and the total subsidy is at most  $n - 1$ .*

**Proof** By Theorem 29,  $\ell_{\mathcal{A}}(u) \leq \max(0, \max_{v \neq u} [c_v(A_u) - c_v(A_v)]) \leq \max(0, 1 - 0) = 1$ , using  $c_v(A_u) \leq 1$  and  $c_v(A_v) \geq 0$ . Observation 12 then gives  $p^* \in \{0, 1\}^n$ , and the endpoint of a heaviest path has  $\ell = 0$ , so the total is at most  $n - 1$ .  $\square$

**Corollary 24** (Fewer chores than agents). *If  $m \leq n$  then Conjecture 2 holds, and a witness is computable by a single minimum-cost bipartite matching.*

**Proof** Among all allocations giving each agent at most one chore, take one minimising  $\sum_i c_i(A_i)$ ; such an allocation exists because  $m \leq n$ , and it is a minimum-cost matching

between agents and chores. Every permutation of its own bundles is again an allocation of this form, so the minimisation was over a set containing all of them, and condition (ii) of Theorem 10 holds:  $\mathcal{A}$  is envy-freeable. Each bundle has  $|A_i| \leq 1$ , so  $c_v(A_i) \leq |A_i| \leq 1$  for every  $v$  — because marginals lie in  $\{0, 1\}$  — and Theorem 23 applies. The bound is tight by Example 1, which has  $m = n - 1$ .  $\square$

The hypothesis of Theorem 23 is genuinely weaker than  $m \leq n$ : it also covers instances with many chores in which every agent finds all but a few of them free, so that no bundle is expensive to anybody. Both statements were checked semantically against the true longest-path computation — exhaustively for  $m \leq 3$ ,  $n \leq 3$ , on 3,000 random instances for the theorem, and on 111,232 instances for the corollary — with no violations (`structural.py`).

**Remark 25** (Why this is worth stating despite being easy). The proof is three lines and the class is small, so the content is not the theorem but the *mechanism*: it is the only argument in this paper that reaches  $\ell \leq 1$  without an exchange step, a balance hypothesis, a schedule, or a search. It does so by making the cycle-closing bound tight from the far side — bounding  $c_v(A_u)$  from above rather than  $c_v(A_v)$  from below. Whether that trick has a wider reach is open; the obvious next question is what happens when bundles cost at most 2, where the same computation gives only  $\ell \leq 2$  and the argument stops.

## 5.4 Instances admitting a uniformly balanced family

The fourth case is not a restriction on the valuation class but on the instance, and it is the only one of the four that is genuinely new here.

**Theorem 26.** *Conjecture 2 holds for every instance admitting a uniformly balanced family, i.e. a partition of  $M$  into  $n$  bundles that every agent values within one unit of each other (Theorem 65). Given such a family, an envy-free solution with  $p^* \in \{0, 1\}^n$  and total subsidy at most  $n - 1$  is obtained by assigning the bundles by a single maximum-weight matching.*

**Proof** Proposition 66.  $\square$

Three things should be said about the scope of this, none of which the proof makes evident.

- **It is a certificate, not an algorithm.** Verifying a proposed family and assigning it are polynomial; *finding* one is a search over partitions, and no polynomial procedure for that is known. So Theorem 26 decides an instance only once somebody supplies the family.
- **The class is large but not everything.** It contains every instance at  $n = m = 3$  — all 9,880, checked exhaustively — and every hard instance the project had accumulated before it was tested. It does *not* contain the discrepancy instance of Proposition 67, at  $n = 3$  and  $m = 4$ , nor the certified-hard set-splitting instance.
- **Membership is not the same as difficulty.** The instances outside the class are not the instances that need large subsidies; on the counterexample of Proposition 67, 19 allocations are good and several need no subsidy at all. Uniform balance is a property of the certificate, not of the instance’s hardness.

So this is a solved case in the same sense as the others in this section: real, checkable, and not load-bearing. What it contributes is the counterexample and Remark 69, which say that any complete method must reason about paths rather than arcs.

## 6 Approach 1: Item-by-Item Insertion

The algorithm of [BKNS22] maintains an envy-free solution  $(\mathcal{A}^t, p^t)$  over a partial allocation with  $p^t \in \{0, 1\}^n$ , and at each step allocates one further good, repairing the invariant by permuting whole bundles among the agents when necessary. Nothing already allocated is ever moved between bundles. This section follows that template on the chore side. Two of its steps come out *easier* than in [BKNS22]; the template nevertheless cannot be completed, and Section 6.3 proves it.

Throughout,  $\mathcal{A}$  is a partial allocation,  $g \notin \bigcup_i A_i$  is an unallocated chore, and  $Z^x, \delta_x, \delta_{i,x}$  are as in Section 3.1.

### 6.1 The easy case is easier than for goods

**Theorem 27** (Free insertion). *Let  $\mathcal{A}$  be envy-freeable and suppose  $\delta_x = 0$  for some agent  $x$ , that is,  $c_x(A_x \cup \{g\}) = c_x(A_x)$ . Then  $Z^x$  is envy-freeable and  $\ell_{Z^x}(i) \leq \ell_{\mathcal{A}}(i)$  for every  $i$ .*

**Proof** By Lemma 14, arcs between agents other than  $x$  are unchanged; arcs out of  $x$  satisfy  $w_{Z^x}(x, j) = w_{\mathcal{A}}(x, j) + \delta_x = w_{\mathcal{A}}(x, j)$ ; and arcs into  $x$  satisfy  $w_{Z^x}(i, x) = w_{\mathcal{A}}(i, x) - \delta_{i,x} \leq w_{\mathcal{A}}(i, x)$ . So every arc weight weakly decreases. No positive-weight cycle can therefore appear, and no path weight can rise; the claims follow from Theorems 10 and 11.  $\square$

It is worth recording what is *absent* from that proof. There is no hypothesis that  $x$  be maximally subsidised, no permutation of bundles, and no extendability machinery: the corresponding step of [BKNS22] needs all three, because for goods the insertion raises the arcs *into* the receiver and those have to be controlled by placing her in  $M(p)$ . Here the insertion lowers them for free. Consequently:

**Observation 28.** *The only case requiring work is a chore whose marginal cost is 1 for every agent on her own current bundle:  $\delta_x = 1$  for all  $x \in N$ . We call this the hard case.*

### 6.2 A cycle-closing bound

The following is the cheapest general handle on  $\ell_{\mathcal{A}}$  available, and it is the chore form of the device [BDN<sup>+</sup>20] uses in its Section 4, where a lower bound of  $-1$  on every arc weight is shown to force a subsidy of at most one per agent.

**Theorem 29** (Cycle-closing bound). *For any envy-freeable allocation  $\mathcal{A}$  and any agent  $u$ ,*

$$\ell_{\mathcal{A}}(u) \leq \max \left( 0, \max_{v \neq u} \left[ c_v(A_u) - c_v(A_v) \right] \right).$$

**Proof** The empty path contributes 0. Let  $P$  be a nonempty path from  $u$  to some  $v \neq u$ . Appending the arc  $(v, u)$  closes  $P$  into a directed cycle, whose weight is at most 0 because  $\mathcal{A}$  is envy-freeable (Theorem 10). Hence  $w_{\mathcal{A}}(P) \leq -w_{\mathcal{A}}(v, u) = c_v(A_u) - c_v(A_v)$ .  $\square$

**Corollary 30.** *If  $c_v(A_u) \leq c_v(A_v) + 1$  for all  $u, v \in N$  — no agent would suffer more than one extra unit by taking anybody else's bundle — then  $p^* \in \{0, 1\}^n$ .*

The hypothesis is sufficient and far from necessary: an agent with  $c_i \equiv 0$  can absorb everything, which breaks the hypothesis while the subsidy stays 0. It is nevertheless what does the work in Theorem 19, where the telescoping bound  $w_{\mathcal{A}}(i, j) \leq u_i - u_j$  supplies exactly this.

### 6.3 The obstruction

In the hard case, Lemma 14 says every arc leaving the receiving agent  $x$  rises by exactly 1. Since a path starting at  $x$  then rises by 1, the condition  $p_x = 0$  is *necessary*. The following theorem shows that no condition of this kind can be sufficient.

Say that  $(\mathcal{A}, p)$  is *extendable with  $g$*  if there is a permutation  $\sigma$  such that  $\mathcal{A}_\sigma$  is envy-freeable and some agent  $\kappa$  has  $c_\kappa(A_{\sigma(\kappa)} \cup \{g\}) = c_\kappa(A_{\sigma(\kappa)})$  — that is, if after some welfare-preserving reassignment of the existing bundles, somebody can take  $g$  for free and Theorem 27 applies.

**Theorem 31** (The insertion template cannot be completed). *There is a three-agent instance with negative dichotomous valuations, an envy-freeable partial allocation  $\mathcal{A}$  with  $p^* \in \{0, 1\}^3$ , and an unallocated chore  $g$ , such that for every agent  $x \in N$  the allocation  $Z^x$  requires a subsidy of 2 — even though the full instance admits an allocation with subsidy 0.*

**Proof** Take  $M = \{a_1, a_2, g\}$  and

$$c_1(S) = \max(0, |S| - 1), \quad c_2(S) = c_3(S) = |S|.$$

Each is dichotomous:  $c_1$  has marginals 0, 1, 1 as  $|S|$  runs 0, 1, 2, and  $c_2, c_3$  are binary additive. Note that  $c_1$  is *supermodular*, which is legal — neither [BKNS22] nor the present paper restricts to submodular valuations — and the instance genuinely needs it.

Consider the partial allocation  $\mathcal{A} = (\{a_1, a_2\}, \emptyset, \emptyset)$ . Its envy graph has weight matrix

$$(w_{\mathcal{A}}(i, j))_{i,j} = \begin{pmatrix} 0 & 1 & 1 \\ -2 & 0 & 0 \\ -2 & 0 & 0 \end{pmatrix},$$

which has no positive-weight cycle, and  $p^* = (1, 0, 0)$ . So  $(\mathcal{A}, p^*)$  is a legal state of the invariant.

The marginal cost of  $g$  is 1 for every agent on her own bundle:  $c_1(\{a_1, a_2, g\}) - c_1(\{a_1, a_2\}) = 2 - 1 = 1$  and  $c_i(\{g\}) - c_i(\emptyset) = 1$  for  $i = 2, 3$ . So we are in the hard case. Moreover  $(\mathcal{A}, p^*)$  is not extendable with  $g$ : agent 1 must keep  $\{a_1, a_2\}$  under any envy-freeable reassignment, since that bundle costs her 1 but costs agents 2 and 3 two apiece, and no reassignment gives anybody a bundle on which  $g$  is free.

Inserting  $g$  into each bundle in turn gives the minimum subsidy vectors

$$Z^1 \mapsto (2, 0, 0), \quad Z^2 \mapsto (2, 1, 0), \quad Z^3 \mapsto (2, 0, 1),$$

each of which has an entry equal to 2. Yet the complete allocation  $(\{a_1\}, \{a_2\}, \{g\})$  has an identically zero envy matrix and subsidy  $(0, 0, 0)$ .  $\square$

The failure is not that envy-freeability breaks — it does not.

**Lemma 32.** *In the hard case, if  $(\mathcal{A}, p^*)$  is not extendable with  $g$ , then  $Z^x$  is envy-freeable for every  $x \in N$ .*

**Proof** Cycles avoiding  $x$  are unchanged by Lemma 14. Let  $C$  be a cycle through  $x$  and let  $i$  be its predecessor of  $x$ . The only arcs of  $C$  that change are the one leaving  $x$ , which rises by  $\delta_x = 1$ , and the one entering  $x$  from  $i$ , which falls by  $\delta_{i,x}$ ; hence  $w_{Z^x}(C) = w_{\mathcal{A}}(C) + 1 - \delta_{i,x}$ .

If  $w_{\mathcal{A}}(C) \leq -1$  then  $w_{Z^x}(C) \leq 0$  immediately. If  $w_{\mathcal{A}}(C) = 0$ , rotating the bundles along  $C$  is a reassignment that preserves total cost, so by Theorem 10(ii) the rotated allocation is again envy-freeable, and in it agent  $i$  receives  $A_x$ . Non-extendability therefore forces  $c_i(A_x \cup \{g\}) \neq c_i(A_x)$ , i.e.  $\delta_{i,x} = 1$ , and again  $w_{Z^x}(C) \leq 0$ . Since  $\mathcal{A}$  is envy-freeable no cycle has  $w_{\mathcal{A}}(C) \geq 1$ , so these cases are exhaustive.  $\square$

So envy-freeability survives every insertion; it is only the  $\{0, 1\}$  bound on the subsidy that breaks. The reason is visible in Section 3.1. A path  $u \rightsquigarrow i \rightarrow x \rightarrow y \rightsquigarrow$  of weight 1 that merely *passes through*  $x$  has its outgoing arc raised by  $\delta_x = 1$  and its incoming arc lowered by  $\delta_{i,x}$ , so it becomes a path of weight 2 whenever  $\delta_{i,x} = 0$ . No condition on  $p_x$  prevents this. In the instance above both  $x = 2$  and  $x = 3$  are hit this way by  $i = 1$ .

Nor can the repair subroutine of [BKNS22] be mirrored. Its navigation rule walks to an agent that has just required a subsidy of at least 2; in the goods world such an agent is automatically maximally subsidised, which is the property the argument needs. In the chore world the same agent automatically has  $p_u = 1$ , and is therefore precisely the kind of agent to whom the chore must *not* be given.

**Remark 33** (What this rules out). Any proof of Conjecture 2 must be allowed to **re-open already-allocated bundles**. The template of [BKNS22] — allocate one item at a time, never move an item that has been allocated, and repair only by permuting whole bundles — is provably insufficient, by Theorem 31. This is a genuine asymmetry between the goods and chore problems and not a gap in the present analysis, so effort spent strengthening the insertion invariant is effort wasted.

## 7 Approach 2: Selecting a Utilitarian-Optimal Allocation

Theorem 31 rules out building the allocation one chore at a time. The natural response is to stop building it at all and instead *select* a good allocation globally, from a set small enough to reason about and rich enough to contain a witness. Utilitarian optimality is the obvious candidate for two reasons. First, by Theorem 10 an allocation minimising  $\sum_i c_i(A_i)$  over *all* allocations, and not merely over reassignments of its own bundles, is automatically envy-freeable, so half of what we need is free. Second, it carries a strong exchange property: if  $g \in A_i$  is *active* for  $i$ , meaning  $c_i(A_i) - c_i(A_i \setminus \{g\}) = 1$ , then  $c_j(A_j \cup \{g\}) = c_j(A_j) + 1$  for every  $j$  — nobody can absorb it for free, or the allocation would not have been optimal. Combined with Theorem 29, a path of weight at least 2 from  $u$  to  $v$  forces  $c_v(A_u) \geq c_v(A_v) + 2$ , so  $A_u$  carries at least two units that are live for  $v$ , and one would like to transfer one of them along the path and show a potential drops.

This is what Conjecture 18 proposes. It is false.



## 7.1 Utilitarian optimality is the wrong restriction

**Proposition 34.** *Conjecture 18 is false. There is an instance with  $n = 3$  agents and  $m = 4$  chores whose unique utilitarian-optimal allocation requires a subsidy of 2, although another allocation — of strictly greater total cost — admits a subsidy of 0.*

**Proof** Take  $M = \{g_1, g_2, g_3, g_4\}$  and the costs tabulated below.

| $S$         | $c_1$ | $c_2$ | $c_3$ | $S$            | $c_1$ | $c_2$ | $c_3$ |
|-------------|-------|-------|-------|----------------|-------|-------|-------|
| $\emptyset$ | 0     | 0     | 0     | $g_3g_4$       | 2     | 1     | 1     |
| $g_1$       | 1     | 1     | 1     | $g_1g_2g_3$    | 2     | 3     | 2     |
| $g_2$       | 1     | 1     | 1     | $g_1g_2g_4$    | 3     | 2     | 2     |
| $g_3$       | 1     | 1     | 0     | $g_1g_3g_4$    | 2     | 2     | 2     |
| $g_4$       | 1     | 1     | 1     | $g_2g_3g_4$    | 2     | 2     | 2     |
| $g_1g_2$    | 2     | 2     | 2     | $g_1g_2g_3g_4$ | 3     | 3     | 3     |
| $g_1g_3$    | 1     | 2     | 1     |                |       |       |       |
| $g_1g_4$    | 2     | 2     | 2     |                |       |       |       |
| $g_2g_3$    | 1     | 2     | 1     |                |       |       |       |
| $g_2g_4$    | 2     | 2     | 2     |                |       |       |       |

Each column is 0 at  $\emptyset$  and has every marginal in  $\{0, 1\}$ , so all three are negative dichotomous in cost form.

*The utilitarian optimum.* Agent 1 pays 1 for every singleton, so  $c_1(A_1) = 0$  only if  $A_1 = \emptyset$ ; agent 3 pays 0 for  $\{g_3\}$  but 1 for every superset, so  $c_3(A_3) = 0$  only if  $A_3 \subseteq \{g_3\}$ . Any allocation of total cost below 2 would need at least two agents at cost 0, forcing  $A_1 = \emptyset$  and  $A_3 \subseteq \{g_3\}$  and hence  $A_2 \supseteq \{g_1, g_2, g_4\}$ , which costs agent 2 at least 2. So the minimum total cost is 2, attained only by

$$\mathcal{A}^* = (\emptyset, \{g_1, g_2, g_4\}, \{g_3\}), \quad 0 + 2 + 0 = 2.$$

Its envy graph has  $w_{\mathcal{A}^*}(2, 1) = 2 - 0 = 2$ , so  $\ell_{\mathcal{A}^*}(2) \geq 2$ ; in fact the minimum subsidy is  $p^* = (0, 2, 0)$ .

*A zero-subsidy allocation.* Take  $\mathcal{B} = (\{g_1, g_3\}, \{g_2\}, \{g_4\})$ , of total cost  $1 + 1 + 1 = 3$ . Its envy graph has weights

$$(w_{\mathcal{B}}(i, j))_{i,j} = \begin{pmatrix} 0 & 0 & 0 \\ -1 & 0 & 0 \\ 0 & 0 & 0 \end{pmatrix},$$

every entry of which is at most 0. Hence there is no positive-weight cycle, every directed path has weight at most 0, and  $p^* = (0, 0, 0)$ .  $\square$

Two features make this sharper than a bare counterexample. The utilitarian-optimal allocation is *unique*, so there is no tie to break and no selection rule inside the optimal set can rescue the conjecture. And the gap is the largest possible in one step: subsidy 2 against subsidy 0, at a cost of one extra unit of total burden. Restricting attention to utilitarian-optimal allocations is therefore not a neutral simplification but an actively wrong one.

**Remark 35** (What this rules out). Conjecture 2, if true, is *not* witnessed by welfare-maximising allocations, and the programme of finding a tie-breaking rule inside the utilitarian-optimal set is unsound as posed — on this instance the set is a singleton and the witness lies outside it. The exchange property quoted at the head of this section remains true and remains usable; what fails is the assumption that a witness may be sought among cost-minimising allocations at all. Any search procedure must be free to give up total cost.

The consequence is worth stating in the strongest available form, because refinement schemes are easy to propose and all of them die here at once.

**Corollary 36** (No refinement of the utilitarian-optimal set can work). *Let  $\mathcal{R}$  be any rule selecting a nonempty subset of the utilitarian-optimal allocations. Then  $\mathcal{R}$  fails on the instance of Proposition 34: every allocation it can return requires a subsidy of 2, while some allocation outside the utilitarian-optimal set requires none.*

**Proof** On that instance the utilitarian-optimal set is the singleton  $\{\mathcal{A}^*\}$ , so  $\mathcal{R}$  returns  $\{\mathcal{A}^*\}$ , and  $\ell_{\mathcal{A}^*}(2) = 2$ .  $\square$

**Remark 37** (A tempting refinement, and why it dies). Three independent signals — Theorem 16, the balancedness guarantee of [BDN<sup>+</sup>20], and Remark 49 — point at *cardinality balance* as the missing ingredient, which suggests refining the utilitarian-optimal set lexicographically by the descending-sorted cardinality vector and then by the descending-sorted own-cost vector. Randomised search over roughly 1050 instances at  $n \in \{3, 4\}$  and  $m \leq 5$  finds no failure, and neither refinement alone survives that same test, so the pair looks strong. It is nevertheless false, by Corollary 36: on the instance of Proposition 34 both refinements are no-ops on a singleton, and the selected allocation needs subsidy 2 against 0 achievable elsewhere. Verified directly by `checkD.py`.

The episode is worth recording for two reasons. First, it is the sampler bias of Section 14.1 firing again: the failure occurs in roughly one instance in  $2.6 \times 10^4$ , so a sweep of  $10^3$  was two orders of magnitude short of being able to see it. Second, the slip underneath is a subtle misreading of Theorem 10. That theorem says an envy-freeable allocation maximises welfare across reassignments of *its own bundles*; it does *not* say that an envy-freeable allocation is globally welfare-maximising. Restricting to utilitarian-optimal allocations is therefore not free, and the zero-subsidy allocation of Proposition 34 — of total cost 3 against the optimum 2, and envy-freeable — is a direct witness that the two notions differ.

## 7.2 Selection rules that fail

Before Proposition 34 was found, three selection rules inside the utilitarian-optimal set were tested, and all three fail on instances where some optimal allocation *does* work. They are recorded because each is a natural first guess.

- **Minimising total perceived load**  $\Psi(\mathcal{A}) := \sum_i \sum_j c_j(A_i)$ , which prefers bundles that are cheap in everybody’s opinion rather than only in their holder’s.
- **Leximin on the perceived-load vector**  $(\max_j c_j(A_i))_{i \in N}$ , which tries to prevent any single bundle from looking heavy to anyone.

- **Local search by single chore transfers** on the objective  $(\max_i \ell_{\mathcal{A}}(i), \sum_i \ell_{\mathcal{A}}(i))$  restricted to the utilitarian-optimal set. This has bad local minima: there are optimal allocations with subsidy 2 from which no single transfer that stays optimal improves the objective.

The local-search failure is the informative one. Single transfers are too coarse a move set, because the utilitarian-optimal allocations of a dichotomous instance are the bases of a structure whose natural moves are exchange *paths and cycles*, not single steps — the same phenomenon that makes Yankee Swap and matroid union the right tools in the matroid-rank literature [BEF21]. A move set matched to that structure is the obvious thing to try next, and remains untried: exchange-path local search over the full allocation lattice — not restricted to the utilitarian-optimal set, per Remark 35 — with the counterexample hunt re-run under it.

## 8 Approach 3: The Replica Transform and the Peel Reformulation

Both preceding approaches fail because of *when* decisions are made. Item-by-item insertion (Section 6) commits a chore to an owner and cannot take it back; selecting a utilitarian optimum (Section 7) commits to the wrong allocation outright. This section changes coordinates so that the commitment is deferred. The chore problem is re-expressed as a *goods* problem on a replicated item set, and then read dynamically as a process that relieves agents of chores one at a time. Ownership of a chore is decided only by the last move that touches it.

The payoff is that the proof orientation of [BKNS22] becomes correct again and the obstruction of Theorem 31 ceases to bind. The price is a new constraint — coverage — which is then shown to be non-vacuous.

### 8.1 The replica instance

Fix a negative dichotomous instance with cost functions  $c_i$  on chores  $M = \{a_1, \dots, a_m\}$ . Build a *replica instance*  $\hat{I}$  whose item set carries  $n - 1$  copies of a good  $b_j$  for each chore  $a_j$ :

$$\hat{M} := \left\{ b_j^{(1)}, \dots, b_j^{(n-1)} : j \in M \right\}.$$

For  $S \subseteq \hat{M}$  let  $\tau(S) \subseteq M$  be the set of *types* occurring in  $S$ , and define

$$\hat{v}_i(S) := c_i(M) - c_i(M \setminus \tau(S)).$$

The reading is that *handing agent  $i$  a copy of  $b_j$  relieves  $i$  of chore  $a_j$* . Each chore ends with exactly one owner, so exactly  $n - 1$  agents must be relieved of it — hence  $n - 1$  copies — and no agent may be relieved of the same chore twice, so no agent may hold two copies of one type. That last clause is load-bearing rather than cosmetic: it is what makes  $\hat{v}_i$  well defined as a dual, since a second copy of a type has marginal 0 and never +1 again.

**Remark 38** (The naive reading is false). If one instead declares a copy of  $b_j$  to be worth +1 to whoever receives it, the construction collapses. An agent indifferent to  $a_j$  would gain from  $b_j$ , every  $\hat{v}_i$  would reduce to  $|\tau(S)|$ , and the induced envy graph would be  $\hat{w}(i, k) = |B_i| - |B_k|$

— the pure cardinality term already isolated in Theorem 16, carrying none of the cost information. The correct reading is *marginal*: a copy of  $b_j$  is worth to agent  $i$  exactly  $i$ 's marginal cost of  $a_j$  on what remains of  $i$ 's workload.

This also completes Remark 9. There the dual  $\tilde{v}_i(S) = c_i(M) - c_i(M \setminus S)$  settled  $n = 2$  and broke at  $n \geq 3$  because  $M \setminus A_i$  stops being a bundle of the partition. Replication is precisely the repair: with  $n - 1$  copies of every type, the  $n$  complements  $M \setminus A_1, \dots, M \setminus A_n$  become *simultaneously* realisable as disjoint bundles.

## 8.2 The transform

Call an allocation  $B = (B_1, \dots, B_n)$  of all of  $\hat{M}$  a *coverage allocation* if no  $B_i$  contains two copies of the same type.

**Theorem 39** (Replica transform). *Let  $c_1, \dots, c_n$  be dichotomous in the sense of Theorem 7 — that is, every marginal  $c_i(S \cup \{g\}) - c_i(S)$  lies in  $\{0, 1\}$ , which already forces  $c_i$  to be monotone. Then*

- (i) *each  $\hat{v}_i$  is a dichotomous valuation on  $\hat{M}$  in the sense of [BKNS22];*
- (ii)  *$\Phi : \mathcal{A} \mapsto B$ , where  $B_i$  receives one copy of each type in  $M \setminus A_i$ , is a bijection between partitions of  $M$  and coverage allocations of  $\hat{M}$ , up to relabelling copies of a type;*
- (iii) *for corresponding  $\mathcal{A}$  and  $B$  the envy graphs agree arc for arc:  $\hat{w}_B(i, k) = w_{\mathcal{A}}(i, k)$  for all  $i, k$ .*

**Proof** (i) Normalisation:  $\hat{v}_i(\emptyset) = c_i(M) - c_i(M) = 0$ . Monotonicity:  $S \subseteq S'$  gives  $\tau(S) \subseteq \tau(S')$ , hence  $M \setminus \tau(S) \supseteq M \setminus \tau(S')$ , and  $c_i$  is monotone, so  $\hat{v}_i(S) \leq \hat{v}_i(S')$ . Marginals: take  $b \notin S$  of type  $j$ . If  $j \in \tau(S)$  then  $\tau(S \cup \{b\}) = \tau(S)$  and the marginal is 0. Otherwise, writing  $T := M \setminus \tau(S)$ , we have  $M \setminus \tau(S \cup \{b\}) = T \setminus \{j\}$  with  $j \in T$ , so

$$\hat{v}_i(S \cup \{b\}) - \hat{v}_i(S) = c_i(T) - c_i(T \setminus \{j\}) \in \{0, 1\},$$

because  $c_i$  has marginals in  $\{0, 1\}$ .

(ii) Forward: type  $j$  is absent from exactly one bundle of  $\mathcal{A}$ , its owner's, so exactly  $n - 1$  agents require a copy of  $b_j$  and the  $n - 1$  available copies are routed injectively. Which physical copy goes where is immaterial, since  $\hat{v}_i$  depends on  $S$  only through  $\tau(S)$ . By construction no  $B_i$  holds a duplicate and  $\bigcup_i B_i = \hat{M}$ . Backward: in a coverage allocation the  $n - 1$  copies of type  $j$  lie in  $n - 1$  *distinct* bundles, so exactly one agent's bundle contains no copy of  $j$ ; setting  $\Psi(B)_i := M \setminus \tau(B_i)$  therefore puts  $j$  in exactly one  $\Psi(B)_i$ , making  $\Psi(B)$  a partition. Finally  $\Psi(\Phi(\mathcal{A}))_i = M \setminus (M \setminus A_i) = A_i$ , and  $\Phi(\Psi(B))$  agrees with  $B$  up to which copy of a type sits where.

(iii) For corresponding  $\mathcal{A}, B$  we have  $\tau(B_k) = M \setminus A_k$ , so  $\hat{v}_i(B_k) = c_i(M) - c_i(A_k)$  and

$$\hat{w}_B(i, k) = \hat{v}_i(B_k) - \hat{v}_i(B_i) = (c_i(M) - c_i(A_k)) - (c_i(M) - c_i(A_i)) = c_i(A_i) - c_i(A_k) = w_{\mathcal{A}}(i, k). \quad \square$$

**Remark 40** (Monotonicity alone is not enough, and which clause does what). The two halves of the hypothesis do different jobs in the proof of (i), and it is worth separating them because the statement is easy to misread. Monotonicity of  $c_i$  is what makes  $\hat{v}_i$  monotone, and *nothing*

else. The binary-marginal condition is what makes  $\hat{v}_i$ 's marginals binary, because those marginals are *downward* marginals of  $c_i$ :

$$\hat{v}_i(S \cup \{b\}) - \hat{v}_i(S) = c_i(T) - c_i(T \setminus \{j\}), \quad T = M \setminus \tau(S),$$

so they inherit exactly the range of  $c_i$ 's own marginals. Monotonicity alone bounds this only by  $\geq 0$ .

The failure is real rather than an artefact of the proof. Take  $m = 2$  with  $c(\emptyset) = c(\{a\}) = c(\{b\}) = 0$  and  $c(\{a, b\}) = 2$ . This  $c$  is monotone and normalised but has a marginal of 2, and the resulting dual has  $\hat{v}(\{b_a\}) - \hat{v}(\emptyset) = 2 - 0 = 2$ , so  $\hat{v}$  is not dichotomous and [BKNS22] does not apply to  $\hat{I}$  at all. Conversely, the normalisation  $c_i(\emptyset) = 0$  is *not* used in Theorem 39:  $\hat{v}_i(\emptyset) = c_i(M) - c_i(M) = 0$  regardless, and the constant  $c_i(M)$  cancels in part (iii). It is retained only because it is part of the model. Verified by `monotone_check.py`, which also confirms by exhaustive enumeration that binary marginals imply monotonicity, so the operative hypothesis is a single clause.

Part (i) is worth stating on its own: *the dual of a dichotomous cost function is a dichotomous value function*, so the class is closed under this duality and the hypotheses of [BKNS22] apply to  $\hat{I}$  verbatim, with no additivity, submodularity or subadditivity assumed anywhere. Because the graphs of part (iii) are equal, envy-freeability, every  $\ell(i)$ , the whole set of feasible subsidy vectors and the  $\{0, 1\}$  target all transfer in both directions without loss.

**Corollary 41** (Coverage is the entire gap). *Conjecture 2 holds if and only if every replica instance  $\hat{I}$  admits a coverage allocation  $B$  with  $\ell_B(i) \leq 1$  for all  $i$ .*

Applying [BKNS22] to  $\hat{I}$  as a black box is legal by Theorem 39(i) and returns some  $B$  with  $\hat{p} \in \{0, 1\}^n$  and total at most  $n - 1$ . But it is free to place two copies of one type in one bundle. If it does, that type is absent from at least two bundles, the chore has at least two owners, and  $\Psi(B)$  is not a partition. Coverage is the sole residual difficulty — and, by Section 8.5, a real one.

**Lemma 42** (Duplicate extraction is envy-neutral). *Removing from  $B_x$  a copy whose type occurs at least twice in  $B_x$  leaves  $\tau(B_x)$  unchanged, hence leaves every  $\hat{v}_k(B_x)$  unchanged, hence leaves the entire envy graph unchanged.*

**Proof** Immediate, since  $\hat{v}_k(S)$  depends on  $S$  only through  $\tau(S)$ . □

So any output of [BKNS22] may be stripped to a duplicate-free partial allocation for free. The difficulty is never in removing duplicates; it is in *re-inserting* the stripped copies into agents that do not already hold that type.

### 8.3 The peel process

Read a coverage allocation as being built one copy at a time and translate back into chores.

**Definition 43** (Peel process). *A state is a workload profile  $W = (W_1, \dots, W_n)$  with  $W_i \subseteq M$ , such that the owner-candidate set  $S_j := \{i : j \in W_i\}$  is nonempty for every  $j \in M$ . Read  $W_i$  as “agent  $i$  is still on the hook for  $W_i$ ”.*

- **Root:**  $W_i = M$  for every  $i$ .
- **Peel**  $\text{peel}(x, j)$ : legal iff  $j \in W_x$  and  $|S_j| \geq 2$ ; it sets  $W_x \leftarrow W_x \setminus \{j\}$ .
- **Permutation:** replace  $W$  by  $(W_{\sigma(1)}, \dots, W_{\sigma(n)})$ .
- **Terminal:**  $|S_j| = 1$  for every  $j$ ; the unique element of  $S_j$  is the survivor, i.e. the owner of  $j$ , and the induced allocation is  $A_i = \{j : S_j = \{i\}\}$ .
- **Graph:**  $w_W(i, k) = c_i(W_i) - c_i(W_k)$ . The root graph is identically zero.

A state is legal if its graph has no positive-weight cycle and  $\ell_W(i) \leq 1$  for every  $i$ .

By Theorem 39 a peel state is exactly a duplicate-free partial allocation of  $\hat{M}$ , a peel is exactly the insertion of one good into one bundle, a permutation is exactly a bundle reassignment, and terminal states are exactly coverage allocations. Nothing is added or lost in the translation.

**Lemma 44** (Peel dynamics). Write  $\mu_k := c_k(W_x) - c_k(W_x \setminus \{j\})$  for  $k$ 's marginal cost of  $j$  on  $x$ 's residual workload. Then  $\text{peel}(x, j)$  changes only the arcs incident to  $x$ , by

$$\Delta w(k, x) = +\mu_k \in \{0, 1\} \quad \text{and} \quad \Delta w(x, k) = -\mu_x \in \{-1, 0\} \quad (k \neq x).$$

**Proof** Only  $W_x$  changes. For  $k \neq x$ ,  $w(k, x) = c_k(W_k) - c_k(W_x)$  and  $c_k(W_x)$  drops by  $\mu_k$ , so the arc rises by  $\mu_k$ . For arcs out of  $x$ ,  $w(x, k) = c_x(W_x) - c_x(W_k)$  and  $c_x(W_x)$  drops by  $\mu_x$ . Arcs between two agents other than  $x$  involve neither  $W_x$  nor  $c_x$ .  $\square$

**The orientation reversal.** This is the point of the whole construction. Compare Lemma 44 against Lemma 14:

| frame                  | move               | in-arcs at $x$ | $x$ should be           |
|------------------------|--------------------|----------------|-------------------------|
| goods [BKNS22]         | insert a good      | rise           | maximally subsidised    |
| chores, insertion (§6) | insert a chore     | fall           | unsubsidised, $p_x = 0$ |
| chores, <b>peel</b>    | relieve of a chore | <b>rise</b>    | maximally subsidised    |

The obstruction of Theorem 31 was exactly this sign clash. The repair subroutine of [BKNS22] navigates to an agent whose tentative subsidy reached 2; in the goods world that agent is automatically maximally subsidised, which is the set the argument needs, whereas in the chore-insertion world she automatically has  $p_u = 1$  and is precisely the agent who must *not* receive the chore. In the peel frame the quantity handed out is *relief*, relief correctly flows to the most envious agents, and the orientation points the right way again.

**Lemma 45** (Spectator-free peels are safe). If  $\mu_k = 0$  for every  $k \neq x$ , then  $\text{peel}(x, j)$  weakly decreases every arc weight, hence  $\ell$  pointwise, and preserves envy-freeability — whatever  $\mu_x$  may be.

**Proof** By Lemma 44 the arcs into  $x$  are unchanged, the arcs out of  $x$  change by  $-\mu_x \leq 0$ , and all other arcs are untouched. Path and cycle weights are sums of arc weights.  $\square$



Lemma 45 is the peel-frame replacement for Theorem 27, and note where the safety condition has moved. In the insertion frame a chore that was free *for the receiver* was unconditionally safe; here what matters is that the chore be free *for everyone else*. That is the goods-side safety structure of [BKNS22], as the identification predicts. The hard case is now a pair  $(x, j)$  for which some other agent  $k$  has  $\mu_k = 1$ .

What does *not* transfer is the coverage constraint, which has no analogue in [BKNS22]: any good may go in any bundle there, whereas a peel demands  $x \in S_j$  and a type is peelable only while  $|S_j| \geq 2$ .

#### 8.4 The insertion obstruction dissolves

Theorem 31 remains true — it is a theorem about the insertion template, and the peel frame is a different template, not a repair of that one. What follows is that its witness no longer obstructs.

Recall the witness:  $M = \{a_1, a_2, g\}$  with  $c_1(S) = \max(0, |S| - 1)$  and  $c_2 = c_3 = |S|$ .

**The trap state is refused before it can form.** The peel-frame counterpart of the trap is  $W = (\{a_1, a_2, g\}, \{g\}, \{g\})$ , and

$$w_W(1, 2) = c_1(\{a_1, a_2, g\}) - c_1(\{g\}) = 2 - 0 = 2,$$

so  $\ell_W(1) = 2$  and the state is already illegal. No legal peel schedule passes through it.

**Remark 46** (The conditioned-remainder principle). The two frames disagree because they compare different things. At a stage where the finished types form a partial assignment  $\mathcal{A}^{(t)}$  and the untouched types form  $R$ , every agent is still on the hook for  $R$ , so the peel-frame arc is

$$c_i(\mathcal{A}_i^{(t)} \cup R) - c_i(\mathcal{A}_k^{(t)} \cup R), \quad \text{not} \quad c_i(\mathcal{A}_i^{(t)}) - c_i(\mathcal{A}_k^{(t)}).$$

Every comparison is conditioned on the common unfinished remainder. For additive costs the  $R$  terms cancel and the two frames coincide — consistent with the additive case already being closed by Theorem 19. For non-additive costs they differ, and on the witness it is exactly the supermodularity of  $c_1$  that the conditioning catches: with  $g$  still pending, loading agent 1 with both of  $a_1, a_2$  already costs 2, and the peel invariant sees it. *The insertion invariant was blind to pending load; the peel invariant has hindsight built in.* This is the substantive content of the change of coordinates.

**A legal schedule reaches the zero-subsidy allocation.** Six peels, every one of them the hard case ( $\mu_k = 1$  for all  $k$  — this instance never offers a spectator-free move), with the invariant holding throughout. Verified by peel.py.

| step | move                                    | resulting $W$                 | $\ell$      | $\sum_i \ell(i)$ |
|------|-----------------------------------------|-------------------------------|-------------|------------------|
| 0    | root                                    | $(a_1a_2g, a_1a_2g, a_1a_2g)$ | $(0, 0, 0)$ | 0                |
| 1    | relieve 1 of $g$                        | $(a_1a_2, a_1a_2g, a_1a_2g)$  | $(0, 1, 1)$ | 2                |
| 2    | relieve 2 of $g$ [owner( $g$ ) = 3]     | $(a_1a_2, a_1a_2, a_1a_2g)$   | $(0, 0, 1)$ | 1                |
| 3    | relieve 3 of $a_1$                      | $(a_1a_2, a_1a_2, a_2g)$      | $(0, 0, 0)$ | 0                |
| 4    | relieve 1 of $a_2$                      | $(a_1, a_1a_2, a_2g)$         | $(0, 1, 1)$ | 2                |
| 5    | relieve 2 of $a_1$ [owner( $a_1$ ) = 1] | $(a_1, a_2, a_2g)$            | $(0, 0, 1)$ | 1                |
| 6    | relieve 3 of $a_2$ [owner( $a_2$ ) = 2] | $(a_1, a_2, g)$               | $(0, 0, 0)$ | 0                |

The terminal allocation is  $(\{a_1\}, \{a_2\}, \{g\})$ , the known zero-subsidy solution. This establishes that the peel template is not killed by the existing obstruction; it does not establish that the template always works, which is Section 8.5. Note also how it meets the requirement of Remark 33 that a correct proof be free to re-open allocated bundles: *peeling never commits a chore's owner until the last peel of that type, so ownership is a deferred decision and there is nothing to re-open.*

## 8.5 The new obstruction: peel dead ends

The natural conjecture is local.

**Conjecture 47** (Local extendability — **refuted**). *From every legal non-terminal state, some legal move exists.*

**Proposition 48.** *Conjecture 47 is false. There are legal, non-terminal peel states from which no sequence of peels and permutations reaches a terminal state without violating the invariant.*

**Proof** A complete backward fixed-point computation over the entire state space of the witness, with both move types and no heuristics, is carried out by `peel3.py`; it reports the following.

| quantity                                                             | value    |
|----------------------------------------------------------------------|----------|
| states in total ( $7^3$ )                                            | 343      |
| legal states                                                         | 175      |
| terminal (partition) states that are legal                           | 6 of 27  |
| legal states from which a terminal is reachable within the invariant | 166      |
| <b>legal dead ends</b>                                               | <b>9</b> |
| root is good                                                         | yes      |

The smallest dead end is  $W = (\{a_1, a_2\}, \{g\}, \{g\})$ , which is legal with  $\ell = (1, 0, 0)$ . The only peelable type is  $g$ , and both continuations fail: relieving agent 2 or agent 3 of  $g$  creates a weight-2 path out of agent 1. Permutations do not help — the profile  $(\{g\}, \{a_1, a_2\}, \{g\})$  is legal but its two continuations fail identically.  $\square$

**What the dead ends are, in the language of [BKNS22].** This is the sharpest statement available. At  $W = (\{a_1, a_2\}, \{g\}, \{g\})$  the remaining unallocated copy is a copy of  $b_g$ , and it has three possible recipients. Giving it to agent 2 or 3 is a genuine peel and violates the invariant. Giving it to agent 1 — who already holds a copy of  $b_g$  — creates a *duplicate*, which is perfectly legal in  $\hat{I}$  and, by Lemma 42, changes nothing whatever in the graph. So the algorithm of [BKNS22] never gets stuck on  $\hat{I}$ : it gets stuck only in the sense that its one legal move is to spend a copy on a duplicate, and spending a copy on a duplicate is exactly the loss of coverage. If agent 1 takes the second  $b_g$  then  $g$  ends with owner-candidate set  $\{2, 3\}$  — two owners, not a partition. Corollary 41 is therefore not a formal restatement; the gap it names is realised.

**Remark 49** (The balance signal — **observation only**). All nine dead ends of the witness have the same shape: *one agent is already committed to a two-element terminal bundle*. They are exactly the three choices of which pair times the three choices of which agent. Correspondingly all six reachable terminals are the six perfect matchings, so on this instance the only legal partitions are the matchings. This is the same signal as Theorem 16 — chores are goods plus cardinality balancing — and as the balancedness guarantee of [BDN<sup>+</sup>20], arriving from a third direction, and it suggests the missing ingredient is a *balance* discipline rather than a subsidy-based one. Consistent with that, restricting recipients to  $\arg \max_i \ell(i)$ , the direct mirror of  $M(p)$ , shrinks the reachable state space but does *not* avoid the dead ends: a subsidy-based rule alone is not enough. This is a single instance and should not be generalised without more data.

## 8.6 The reduced open problem

**Conjecture 50** (Reachability — **open**). *From the root, some sequence of peels and permutations reaches a terminal state with the invariant holding at every intermediate state.*

Conjecture 50 implies Conjecture 2. The root graph is identically zero, so the invariant holds there; if a legal schedule reaches a terminal state  $W^{\text{end}}$  then the induced allocation  $\mathcal{A}$  has  $G_{\mathcal{A}} = G_{W^{\text{end}}}$ , with no positive cycle and  $\ell_{\mathcal{A}}(i) \leq 1$  for all  $i$ ; Observation 12 then forces  $p^* \in \{0, 1\}^n$ , and the endpoint of a maximum-weight path has  $\ell = 0$ , giving a total of at most  $n - 1$ .

**Remark 51** (The converse is open, and it matters for what the experiments test). Conjecture 50 implies Conjecture 2, but the implication has not been established in the other direction and should not be assumed. Conjecture 2 asserts that a good terminal state *exists*; Conjecture 50 asserts that one is *reachable along a path every state of which is legal*. These can come apart: a good terminal state may exist while every schedule leading to it passes through an illegal intermediate state, and Proposition 48 shows that illegal-looking intermediate states are not hypothetical. The practical consequence is that an instance with a *bad root* would refute Conjecture 50, and hence kill this approach, but would say nothing directly about Conjecture 2, which would have to be attacked by exhaustive search over allocations as in Section 14. Establishing the converse — that a good terminal state, when one exists, is always reachable legally — would make the peel frame a complete reformulation rather than a sufficient one, and is worth attempting on its own.

It is not automatic. Conjecture 50 is a reachability statement about a state space with genuine dead ends, so a proof must supply a *rule* for choosing the next peel that provably never enters the bad region, or else avoid the state space altogether by a potential or exchange argument. Two resources are available that [BKNS22] never had:

1. **Scheduling freedom.** Many types are unfinished at once, and we need only *some* type to admit a legal peel, not a prescribed one.
2.  $|S_j| \geq 2$ . A peelable type always has at least two candidate recipients, so the forbidden recipient set is never all of  $N$ .

Call a legal non-terminal state a *deadlock* if for every unfinished type  $j$  and every permutation, no recipient in  $S_j$  preserves the invariant. Proposition 48 shows deadlocks exist; Conjecture 50 asks only that the *root* avoid them. So the target theorem has the shape *there is a peel rule under which no deadlock is ever entered*, and Remark 49 is the first candidate for what that rule must enforce.

## 8.7 The type-ordered restriction

One restriction of the peel process survives every test so far and shrinks the search substantially. Present the chores in *blocks*: decide the owner of  $a_1$  completely, then  $a_2$ , and so on. Coverage is automatic in the peel frame — a peel deletes  $j$  from  $W_x$ , so the same pair  $(x, j)$  cannot be peeled twice — and after  $r$  complete blocks the state is  $W_i = \mathcal{A}_i^{(r)} \cup R$  with  $R$  the undecided chores, which is literally the conditioned-remainder principle of Remark 46. A state is then just a triple (partial allocation, current chore, set of agents already relieved of it), which is a far smaller space than the general peel process.

**Conjecture 52** (Type-ordered reachability — **open**). *From the root, some chore order  $\pi$ , together with some choice of owner and within-block sequence for each chore, reaches a terminal state with the invariant holding at every intermediate state.*

Conjecture 52 implies Conjecture 50 implies Conjecture 2, and it is a strictly smaller search problem than either. The evidence, from `typeorder.py` and `stress.py`, is that the restriction costs nothing: across the witness of Theorem 31 and 890 random dichotomous instances at  $n \in \{3, 4\}$ ,  $m \leq 4$ , *every* instance that admits a legal schedule at all admits a type-ordered one. The chore *order* is not free — in 8/120 instances at  $n = m = 3$ , 16/60 at  $n = 3, m = 4$  and 8/60 at  $n = 4, m = 3$  only some orders work — but in every instance tested some order does. Note that the schedule exhibited in Section 8.4 is *not* type-ordered, since it interleaves  $a_1$  and  $a_2$ ; the table says a type-ordered one exists for that instance anyway, which was not obvious.

What does not survive is the natural recipient rule to pair with it.

**Proposition 53** (arg max recipient selection is insufficient — **refuted**). *Giving each copy to an agent of maximum current subsidy does not suffice, even under the type-ordering.*

**Proof** Take  $n = 3, m = 2$  with the costs

$$c_1(\emptyset) = 0, \quad c_1(\{a_1\}) = c_1(\{a_2\}) = 0, \quad c_1(M) = 1; \quad c_2(S) = \min(|S|, 1); \quad c_3 = c_1.$$

All three are dichotomous (Theorem 7). Exhaustive search over both chore orders, both roles and all within-block sequences finds no legal type-ordered schedule in which every recipient lies in  $\arg \max_i \ell_W(i)$ . Dropping only the  $\arg \max$  restriction, the schedule that relieves agent 1 of  $a_1$ , then agent 3 of  $a_1$ , then agent 1 of  $a_2$ , then agent 3 of  $a_2$  is legal and terminates at  $\mathcal{A} = (\emptyset, \{a_1, a_2\}, \emptyset)$  with  $\ell = (0, 1, 0)$ .  $\square$

This is the signal of Remark 49 again — a subsidy-based recipient rule alone is not enough — now confirmed *inside* the type-ordered restriction, which is the sharper statement. Combined with Theorem 62, which says the recipient in the hard case has zero marginal, the rule provably cannot be read off the subsidy vector: it must see the residual workload, which is exactly what the conditioned-remainder arcs already encode.

## 8.8 A subclass that may fall first — conjectural

On type space the dual of a *supermodular* dichotomous cost is submodular with binary marginals, that is, a matroid rank function; lifting through  $\tau$  to the copies is the parallel extension of that matroid, still a matroid rank function. Additive costs dualise to partition-matroid ranks. If that is right — and the parallel-extension step in particular deserves a full proof before being relied on — then supermodular dichotomous chores are exactly matroid-rank goods coverage problems, and the witness of Theorem 31 lies entirely inside this class, since  $c_1$  dualises to the rank of  $U_{2,3}$ , parallel extended.

Matroid rank is where the machinery of the surrounding literature is strongest, through the exchange properties used in [BSY23, BEF21]. That makes Conjecture 2 restricted to supermodular dichotomous costs, attacked by basis exchange in the parallel-copy matroid, the natural intermediate target — sitting between Theorem 19, which is closed, and the full conjecture.

## 9 Approach 4: Encoding Coverage into the Numbers

Corollary 41 reduces Conjecture 2 to a single constraint: the replica instance  $\hat{I}$  must admit an allocation in which no agent holds two copies of one type. Applying [BKNS22] to  $\hat{I}$  is legal but returns an allocation free to violate that. The obvious response is to change the numbers so that a violation becomes unattractive — either at the level of the valuations, by charging for duplicates, or at the level of the algorithm, by inflating weights so that the wrong recipient is never selected.

This section closes both. They fail for the same reason, isolated in Section 9.6.

### 9.1 The literal proposal is a no-op

**Theorem 54.** *In  $\hat{I}$  the marginal value of a second copy of a type is already exactly 0, for every agent and every set: if  $b$  has type  $j$  and  $j \in \tau(S)$  then  $\hat{v}_i(S \cup \{b\}) - \hat{v}_i(S) = 0$ .*

**Proof**  $\tau(S \cup \{b\}) = \tau(S)$ , and  $\hat{v}_i$  factors through  $\tau$ .  $\square$

So “let a first copy carry weight and a second carry none” is not a modification of  $\hat{I}$ ; it is  $\hat{I}$ . Remark 38 already records that the other choice, a flat +1 per copy, destroys the transform outright.

The reason this does not help is that zero marginal means *indifferent*, not *averse*. Lemma 42 is the exact statement of the damage: a duplicate is envy-neutral, changing no arc of the envy graph at all. An objective built out of envy comparisons cannot see duplicates, and the dead ends of Proposition 48 exist because of this rather than in spite of it — at the dead end the one legal move is to spend a copy on a duplicate, which is free. To repel a duplicate rather than merely fail to attract it, the duplicate must be made *visible*.

## 9.2 The separation programme, correctly stated

Write  $d(S) := |S| - |\tau(S)|$  for the number of duplicated copies in  $S$ . Call  $u = (u_1, \dots, u_n)$  on  $\hat{M}$  a *faithful reweighting* if

- (D) each  $u_i$  is dichotomous —  $u_i(\emptyset) = 0$ , monotone, all marginals in  $\{0, 1\}$  — so that [BKNS22] applies to  $(\hat{M}, u)$ ; and
- (F)  $u_i(S) = \hat{v}_i(S)$  for every duplicate-free  $S$ , so that the envy graph of a coverage allocation is unchanged.

Say  $u$  *separates* if every non-coverage allocation  $B$  has  $\max_i \ell_B^u(i) \geq 2$ .

**Observation 55** (The programme is logically valid). *If some faithful  $u$  separates, then Conjecture 2 holds for that instance.*

**Proof** Applying [BKNS22] to  $(\hat{M}, u)$  returns  $B$  with  $\ell_B^u \in \{0, 1\}^n$ ; separation forces  $B$  to be a coverage allocation; on a coverage allocation (F) gives  $u_i(B_k) = \hat{v}_i(B_k)$  for all  $i, k$ , so  $\ell_B^{\hat{v}} = \ell_B^u \leq 1$ ; and Theorem 39 hands back a chore partition with per-agent subsidy in  $\{0, 1\}$  and total at most  $n - 1$ .  $\square$

The idea is therefore not confused — it is a genuine reduction template. The rest of this subsection shows it cannot be filled in.

**Theorem 56** (Duplicate budget). *Every faithful reweighting satisfies  $\hat{v}_i(S) \leq u_i(S) \leq \hat{v}_i(S) + d(S)$  for all  $i$  and all  $S \subseteq \hat{M}$ . So the penalty  $f_i := u_i - \hat{v}_i$  is non-negative, vanishes on duplicate-free sets, and is bounded by one unit per duplicated copy.*

**Proof** Let  $T \subseteq S$  be duplicate-free with  $\tau(T) = \tau(S)$ , so  $|S \setminus T| = d(S)$  and  $\hat{v}_i(T) = \hat{v}_i(S)$ . Monotonicity gives  $u_i(S) \geq u_i(T) = \hat{v}_i(T) = \hat{v}_i(S)$  by (F). Adding the  $d(S)$  elements of  $S \setminus T$  one at a time raises  $u_i$  by at most 1 each by (D), so  $u_i(S) \leq u_i(T) + d(S)$ .  $\square$

The extreme  $f_i = d$  is attained, since  $u_i = \hat{v}_i + d$  is monotone, normalised and has marginals in  $\{0, 1\}$ . So “a duplicate is worth +1 to everybody” is the strongest legal penalty there is, and any scaling by  $\varepsilon \notin \{0, 1\}$  leaves the dichotomous class and forfeits [BKNS22] altogether.



### 9.3 The strongest uniform penalty is a potential

**Theorem 57.** *Let  $u_i = \hat{v}_i + \varepsilon d$  for every  $i$ , with the same  $\varepsilon$  and the same  $d$  for all agents. Then for every allocation  $B$ ,*

$$w_B^u(i, k) = w_B^{\hat{v}}(i, k) + \varepsilon(d(B_k) - d(B_i)),$$

*and consequently*

- (1) *every directed cycle has the same weight in both graphs, so  $B$  is envy-freeable under  $u$  if and only if it is under  $\hat{v}$ : a uniform penalty can never disqualify an allocation;*
- (2) *the agent holding the most duplicates has weakly smaller subsidy and everyone else weakly larger — the penalty lands on the wrong agents;*
- (3) *if  $d(B_1) = \dots = d(B_n)$  the two graphs are identical, so such allocations are invisible to the penalty however large  $\varepsilon$  is.*

**Proof** The identity is immediate from the definitions. The correction telescopes to 0 around any cycle, giving (1), and along a path to endpoint minus start, giving (2). For (3) all corrections vanish.  $\square$

**Example 58** (Maximally non-coverage, and invisible). Take  $n = 3, m = 3$  and  $c_i(S) = |S|$  for every  $i$ , and let  $B_i$  be both copies of  $b_i$ . Then  $d(B_i) = 1$  for every  $i$ , so Theorem 57(3) applies;  $\hat{v}_i(B_k) = |\tau(B_k)| = 1$  for all  $i, k$ , every arc is 0 and  $\ell_B^u = (0, 0, 0)$ . Yet  $\Psi(B)_i = M \setminus \{a_i\}$  gives every chore two owners — as far from coverage as it is possible to be. Applying [BKNS22] to this  $u$  may return exactly this allocation.

### 9.4 No reweighting whatsoever separates

Theorem 57 leaves open the asymmetric case, in which different agents charge for duplicates of different types. That does break the potential structure and can create positive cycles, so it has to be excluded separately.

**Theorem 59** (No-go). *There is a negative dichotomous instance on which no faithful reweighting separates:  $n = 3, m = 2, c_1 = c_2 = c_3 = |\cdot|$ .*

**Proof** The replica has types  $a_1, a_2$  with two copies each and  $\hat{v}_i(S) = |\tau(S)|$ . Write  $D$  for the pair of copies of  $b_1$ , and  $s_1, s_2$  for the two copies of  $b_2$ . For  $x \in \{1, 2, 3\}$  let  $B^{(x)}$  give  $D$  to agent  $x$  and one of  $s_1, s_2$  to each of the others; each is non-coverage.

Let  $u$  be faithful. The singletons are duplicate-free, so  $u_i(s_1) = u_i(s_2) = 1$  for every  $i$  by (F). For  $D$  we have  $\hat{v}_i(D) = 1$  and  $d(D) = 1$ , so Theorem 56 gives  $u_i(D) = 1 + \delta_i$  with  $\delta_i \in \{0, 1\}$  — and  $\delta = (\delta_1, \delta_2, \delta_3)$  is the *entire* freedom  $u$  has on this family. Fixing  $x$  and writing  $\{y, z\} = N \setminus \{x\}$ , the arcs of  $B^{(x)}$  are

$$w(y, x) = \delta_y, \quad w(z, x) = \delta_z, \quad w(x, y) = w(x, z) = -\delta_x, \quad w(y, z) = w(z, y) = 0.$$

If  $\delta_x = 1$ : every arc out of  $x$  is  $-1$ , every arc into  $x$  is at most 1, and the  $y \leftrightarrow z$  arcs are 0. Any cycle uses one arc into  $x$  and one out, so has weight at most 0; any path out of  $y$  has weight at most  $\max(0, \delta_y, \delta_z) \leq 1$ , likewise out of  $z$ , and  $\ell(x) = 0$ . So  $B^{(x)}$  survives. If  $\delta_x = 0$

and  $\delta_y = \delta_z = 0$ : every arc is 0 and  $B^{(x)}$  survives. If  $\delta_x = 0$  and  $\delta_k = 1$  for some  $k \neq x$ : the two-cycle  $k \rightarrow x \rightarrow k$  has weight  $\delta_k - \delta_x = 1 > 0$ , and  $B^{(x)}$  is excluded.

So  $B^{(x)}$  is excluded exactly when  $\delta_x = 0$  and  $\delta_k = 1$  for some  $k \neq x$ . Separation demands all three of  $B^{(1)}, B^{(2)}, B^{(3)}$  be excluded, forcing  $\delta_1 = \delta_2 = \delta_3 = 0$  and hence no  $\delta_k = 1$  — a contradiction.  $\square$

**Theorem 60** (Machine-verified strengthening). *On the same instance separation fails even if faithfulness is dropped entirely. Among all 990 dichotomous valuations on the four replica copies, none of the  $990^3$  triples  $(u_1, u_2, u_3)$  excludes all twelve non-coverage allocations of the shape “one duplicated pair plus two singletons”.*

**Proof** Exhaustive, by `dupsep.py`: collapsing the 990 valuations by their behaviour on the critical family leaves 30 classes, and 0 of the resulting class-triples survive the filter. No heuristics are used. The same script also sweeps the separable  $P$ -family  $u_i(S) = \hat{v}_i(S) + \sum_{j \in P_i} (|S_j| - 1)^+$ , which for  $n = 3$  is the whole separable penalty family, and finds no separating member on  $n = 3, m = 2$  (all 64 triples), on  $n = 3, m = 3$  with unit costs (all 512), or on the witness of Theorem 31 (all 512).  $\square$

The instance of Theorem 59 is *not* a counterexample to Conjecture 2: the coverage allocation  $B = (\{b_1^{(1)}, b_2^{(1)}\}, \{b_1^{(2)}\}, \{b_2^{(2)}\})$ , i.e. the chore partition  $(\emptyset, \{a_2\}, \{a_1\})$ , has  $\ell = (0, 1, 1)$  and total  $2 = n - 1$ . What is refuted is the *method*, on the easiest instance available.

## 9.5 The same failure at the algorithm level

The second form of the idea leaves the valuations alone and instead inflates the weight of the block of copies currently being allocated, intending that receiving a copy drive the recipient’s subsidy to 0 so that an agent already holding that type is never chosen again. The diagnosis behind it is right.

**Theorem 61.** *In any run of the algorithm of [BKNS22] on  $\hat{I}$ , the EXTEND branch never places a duplicate. Every coverage violation is created by FINDSINK.*

**Proof** EXTEND returns a pair  $(\rho, k)$  drawn from its loop over  $k \in N$  and  $\ell \in M(p)$  satisfying  $v_k(A_\ell \cup \{g\}) - v_k(A_\ell) = 1$ , and places  $g$  into the bundle  $B_k = A_\ell$ . So the receiving bundle has strictly positive marginal for  $g$ . By Theorem 54 a copy of  $b_j$  has marginal 0 whenever  $j \in \tau(S)$ , so the receiving bundle contains no copy of  $j$ .  $\square$

But the lever cannot reach the subroutine that does the damage.

**Theorem 62** (FINDSINK runs only in the zero-marginal regime). *Suppose the algorithm reaches its FINDSINK branch, i.e.  $(\mathcal{A}, p)$  is not extendable with  $g$ . Then every  $\ell \in M(p)$  satisfies  $v_\ell(A_\ell \cup \{g\}) - v_\ell(A_\ell) = 0$ .*

**Proof** Suppose some  $\ell \in M(p)$  had marginal 1, and run EXTEND’s loop at  $k = \ell$ , which is then a feasible choice. It computes a maximum-weight matching  $\rho$  on the complete bipartite graph between  $N \setminus \{\ell\}$  and itself with weights  $v_i(A_j)$ ; maximality against the identity matching gives  $\sum_{i \neq \ell} v_i(A_{\rho(i)}) \geq \sum_{i \neq \ell} v_i(A_i)$ . Setting  $\rho(\ell) = \ell$  adds  $v_\ell(A_\ell)$  to both sides, so the welfare test passes and EXTEND returns  $(\rho, \ell)$  — contradicting non-extendability.  $\square$

**Corollary 63** (The recipient’s subsidy is exactly preserved). *Let  $s$  be the agent returned by `FINDSINK` and  $\mathcal{A}' = (A_1, \dots, A_s \cup \{g\}, \dots, A_n)$ . Then  $\ell_{\mathcal{A}'}(s) = \ell_{\mathcal{A}}(s)$ .*

**Proof** [BKNS22] places  $s \in M(p)$ , so Theorem 62 gives  $v_s(A_s \cup \{g\}) = v_s(A_s)$  and every arc out of  $s$  is unchanged. Arcs not incident to  $s$  are unchanged, and only arcs into  $s$  can rise. Since  $\mathcal{A}'$  is envy-freeable it has no positive cycle, so  $\ell_{\mathcal{A}'}$  is the maximum weight over *simple* paths; a simple path leaving  $s$  never re-enters  $s$  and therefore uses only unchanged arcs.  $\square$

Three consequences, each fatal on its own.

1. **The lever multiplies a zero.** The premise — that receiving a copy forces the recipient’s subsidy to 0 — holds only where the recipient’s marginal is positive. Theorem 62 says that in the `FINDSINK` branch it is provably 0, and scaling by any weight leaves 0 where it was. Corollary 63 makes it exact: the recipient’s subsidy does not move at all. Inflation bites only in the `EXTEND` branch, which by Theorem 61 was never broken.
2. **A duplicate is `FINDSINK`’s safest move.** By Lemma 42 a second copy changes no arc whatsoever, so the subsidy vector after the tentative assignment is literally identical, the subroutine’s guard fails on the first test, and it returns immediately. A duplicate is never rejected because there is nothing to detect.
3. **The scaling forfeits the guarantee anyway.** With block weights the marginals lie in  $\{0, W_j\}$  rather than  $\{0, 1\}$ , and the step of [BKNS22] that concludes  $p \in \{0, 1\}^n$  from integrality, together with the unit path bounds throughout its analysis, is gone. Using the lever is not *applying* [BKNS22]; it is writing a new algorithm and owing a new proof.

It is worth recording where the lever *does* bite, because that is why it looks like it works. Within a block, once some agent has been relieved of  $a_j$  while others still hold it, the arcs into the relieved agent from any holder with marginal 1 carry the inflated weight, so  $M(p)$  collapses onto exactly the holders who care — and those are precisely the agents for whom `EXTEND` applies. The lever keeps `FINDSINK` from firing at all. It runs out exactly when every remaining holder of  $a_j$  has marginal 0: then nothing is inflated,  $M(p)$  reverts to stale values from earlier blocks, and an already-relieved agent can sit at the top of it. The minimal shape is  $n = 3$  with a chore of marginal 1 for one agent and 0 for the other two.

## 9.6 Why both fail: holder-blindness

The two failures have one cause, and naming it is the useful output of this section.

A duplicate penalty bites only through the arcs. The arc effect of agent  $k$  charging +1 for the duplicates in bundle  $B_x$  is +1 on  $w(k, x)$  — but the same charge levied by  $x$  on her own bundle is +1 on  $u_x(B_x)$ , which is  $-1$  on every arc *out of*  $x$ . The two cancel around every cycle. So a penalty can disqualify an allocation only if the agent who *holds* the duplicate does not feel it while somebody else does. But  $u$  is fixed before the allocation is chosen, and it is the allocation that decides who holds the duplicate. Requiring the holder to be blind for every possible holder forces the penalty to zero, which is exactly the  $\delta$  argument of Theorem 59.

Compare the orientation reversal of Section 8.3. There the difficulty was that the navigation rule of [BKNS22] sends the item to the agent who must not receive it. Here the difficulty is that the penalty must be levied on an agent who cannot be identified in advance. Both are the same species of failure: *a rule fixed ex ante against a quantity determined ex post*.

**Remark 64** (What this rules out, and what it does not). **Ruled out:** encoding coverage into the valuations of the replica instance so that [BKNS22] can be used as a black box. Any such encoding is a dichotomous reweighting, and Theorem 60 kills it. Also ruled out: steering the algorithm of [BKNS22] by inflating weights, by Theorem 62. Together these close the whole family “*encode coverage into the numbers and reuse [BKNS22] as a black box*” from two directions.

**Not ruled out:** coverage-aware *algorithms* — Conjecture 50 and the search for a peel rule that never enters a deadlock. Coverage is a constraint on the allocation, and the right place to enforce a constraint on the allocation is the procedure that builds it, not the objective it is scored against.

## 10 Approach 5: Agent Induction and the Balance Invariant

Every approach so far inducts on *items*. Inducting on *agents* instead is attractive for chores because a newcomer starts with an empty bundle, which is the cheapest bundle there is, so the newcomer is the natural sink and the orientation problem of Section 6 does not arise. This section sets up what such an induction would have to carry, and shows that the natural candidate fails.

### 10.1 What the induction would have to carry

Corollary 30 says that if  $\mathcal{A}$  is envy-freeable and every arc satisfies  $w_{\mathcal{A}}(i, j) \geq -1$  — equivalently  $c_i(A_j) \leq c_i(A_i) + 1$  — then  $p^* \in \{0, 1\}^n$ . The useful feature of that condition for an induction on agents is that it is *assignment-invariant*: it constrains only the family of bundles, not which agent holds which. That suggests carrying the family rather than the allocation.

**Definition 65** (Uniformly balanced family). *A family of  $n$  disjoint bundles covering  $M$  is uniformly balanced if every agent values all  $n$  bundles within one unit of each other:  $\max_t c_i(B_t) - \min_t c_i(B_t) \leq 1$  for every  $i \in N$ .*

**Proposition 66.** *If an instance admits a uniformly balanced family, then Conjecture 2 holds for it.*

**Proof** Assign the bundles to agents by a maximum-weight matching. By Theorem 10 the resulting allocation is envy-freeable, and by uniform balance every arc satisfies  $w_{\mathcal{A}}(i, j) = c_i(A_i) - c_i(A_j) \geq -1$  — the bound holds for *any* assignment, so in particular for this one. Corollary 30 then gives  $p^* \in \{0, 1\}^n$ , and Observation 12 plus the vanishing of  $\ell$  at the endpoint of a heaviest path give a total of at most  $n - 1$ .  $\square$

This is exactly the shape an agent induction wants: a purely combinatorial, schedule-free, coverage-free statement, with the step from  $n$  to  $n + 1$  agents being a refinement of the family. It also holds on every hard instance the project has accumulated, and on the full exhaustive family at  $n = m = 3$  — all 9,880 instances, with the implication chain of Proposition 66 verified end to end on each (routeA.py). It is nevertheless false.

## 10.2 The invariant is unreachable

**Proposition 67.** *There is a negative dichotomous instance with  $n = 3$  and  $m = 4$ , with binary additive costs, admitting no uniformly balanced family.*

**Proof** Take  $c_i(S) = |S \cap D_i|$  with

$$D_1 = \{g_1, g_2\}, \quad D_2 = \{g_1, g_3, g_4\}, \quad D_3 = \{g_2, g_3, g_4\}.$$

Uniform balance for agent  $i$  means the elements of  $D_i$  are spread across the three bundles with counts of range at most 1. Three elements over three bundles have range at most 1 only in the pattern  $(1, 1, 1)$ , so  $D_2$  forces  $g_1, g_3, g_4$  into three *different* bundles, one each, and  $D_3$  likewise forces  $g_2, g_3, g_4$  into three different bundles. Hence  $g_2$  must lie in the one bundle containing neither  $g_3$  nor  $g_4$ , which is the bundle containing  $g_1$ . But two elements over three bundles have range at most 1 only in the pattern  $(1, 1, 0)$ , so  $D_1$  forces  $g_1$  and  $g_2$  into *different* bundles. Contradiction.  $\square$

The instance is binary additive and therefore already settled by [LMS26]; what it refutes is the method, on a class the literature closes. Indeed 19 of its 81 allocations are good, several with  $\ell = (0, 0, 0)$  — it is an easy instance for the conjecture and an impossible one for the invariant.

The failure is sharper than mere unreachability.

**Proposition 68.** *On the instance of Proposition 67, every good allocation has an arc of weight  $-2$  or less. Hence no sufficient condition of the form “all arcs  $\geq -1$ ” can certify any of them.*

**Proof** Exhaustive over all  $3^4 = 81$  allocations (routeA\_cex.py): 19 are good, and the minimum arc weight over them is  $-3$ , attained at  $(\{g_1, g_3, g_4\}, \{g_2\}, \emptyset)$ , with every other good allocation reaching  $-2$ .  $\square$

**Remark 69** (What this rules out). Corollary 30 is a bound on *arcs*; the target is a bound on *paths*. Propositions 67 and 68 show the gap between the two is not slack that a cleverer construction can close — it is where all the surviving witnesses live. **Any correct invariant must be path-based.** That is precisely what makes agent induction hard, because path weights depend on which agent holds which bundle, so the invariant cannot be assignment-invariant and the newcomer’s arrival re-opens the whole graph. Agent induction is not refuted here; the only invariant that would have made it easy is.

**Remark 70** (The obstruction is hypergraph discrepancy). For binary additive costs, uniform balance says every  $D_i$  is split evenly across the  $n$  bundles, which is a discrepancy condition on the set system  $\{D_1, \dots, D_n\}$ . Proposition 67 is a small discrepancy obstruction, and it is the same combinatorial phenomenon that drives the reduction by which [BSV21] proves deciding exact envy-freeness for binary additive chores NP-complete. The certified-hard set-splitting instance recorded from that reduction also admits no uniformly balanced family, checked by routeA\_adv.py. So the balance signal of Remark 49, which three independent routes had pointed at, cannot be imposed exactly — only approximately, and the approximation is exactly the path slack.

### 10.3 The $n = 3$ specialisation

Since the attack was aimed at  $n = 3$  first, it is worth recording exactly what the target reduces to there. The point is that with three agents the path structure is finite and small, so “ $\ell \leq 1$ ” becomes a fixed list of inequalities.

**Proposition 71** (Characterisation at  $n = 3$ ). *Let  $n = 3$ . An allocation  $\mathcal{A}$  is good — envy-freeable with  $\ell_{\mathcal{A}} \leq 1$  — if and only if*

- (i) *every directed cycle of  $G_{\mathcal{A}}$  has weight at most 0;*
- (ii) *every arc has weight at most 1;*
- (iii) *every directed path of length two has weight at most 1.*

**Proof** Condition (i) is envy-freeability, by Theorem 10. Given it, no positive cycle exists, so  $\ell_{\mathcal{A}}(i)$  is the maximum weight of a *simple* path leaving  $i$ ; on three vertices such a path has length 0, 1 or 2, contributing weight 0, an arc, or a two-path respectively. So  $\ell_{\mathcal{A}} \leq 1$  is exactly (ii) together with (iii).  $\square$

**Corollary 72** (Only six inequalities, and most are free). *Let  $n = 3$  and let  $\mathcal{A}$  satisfy (i) and (ii). For distinct  $i, j, k$  the three-cycle bound gives  $w_{\mathcal{A}}(i, j) + w_{\mathcal{A}}(j, k) \leq -w_{\mathcal{A}}(k, i)$ , so the two-path  $i \rightarrow j \rightarrow k$  has weight at most 1 whenever  $w_{\mathcal{A}}(k, i) \geq -1$ . As  $(i, j, k)$  runs over the six orderings of  $N$ , the closing arc  $(k, i)$  runs over all six ordered pairs exactly once. Hence exactly those two-paths whose closing arc lies below  $-1$  require separate checking, one per offending arc. In particular, if every arc lies in  $[-1, 1]$  then (iii) is automatic and  $\mathcal{A}$  is good.*

Both statements are machine-checked against the true longest-path computation over  $5.8 \times 10^5$  (instance, allocation) pairs — exhaustively at  $n = m = 3$  and randomised to  $m = 6$  — by `n3check.py`.

Corollary 72 is the sharpest handle available at  $n = 3$ , and it also explains precisely why Section 10 failed. The clause “if every arc lies in  $[-1, 1]$ ” is the uniform-balance hypothesis; the residual work is the two-paths whose closing arc is very negative; and Proposition 68 says that on the counterexample *every* good allocation has such an arc, so the residual work is not a corner case but the whole problem.

### 10.4 What survives: the two-tier restatement

One reformulation does survive, and it is worth recording because it is equivalent rather than merely sufficient.

**Proposition 73.** *Let  $\mathcal{A}$  be envy-freeable with every arc at most 1. Then  $\mathcal{A}$  is good if and only if the digraph of weight-1 arcs contains no directed path of length two.*

**Proof** If no such path exists then every directed path carries at most one positive arc, so has weight at most 1, and  $\ell_{\mathcal{A}} \leq 1$ . Conversely if  $\ell_{\mathcal{A}} \leq 1$  then  $p^* \in \{0, 1\}^n$  by Observation 12; writing  $S = \{i : p_i^* = 1\}$ , Proposition 15 puts every arc within  $S$ , within  $N \setminus S$ , or from  $N \setminus S$  into  $S$  at weight at most 0, so weight-1 arcs run only from  $S$  to  $N \setminus S$  and cannot be composed.  $\square$



So the target is exactly: *find an envy-freeable allocation together with a bipartition of the agents such that all strict envy runs from one side to the other*. This is Proposition 15 again, now as a criterion rather than a certificate format. It is not assignment-invariant, consistent with Remark 69, but it is a condition on a bipartition rather than on a schedule, so it inherits none of the dead ends of Section 8 and none of the ex-ante failures of Section 9.

Two facts about the paid set are worth recording before that route is attempted. The first is a theorem and constrains what the set can be; the second is an observation from Section 14 and suggests where to look.

**Proposition 74** (Paid agents are those somebody finds expensive). *Let  $\mathcal{A}$  be envy-freeable. If  $p_i^* \geq 1$  then there is an agent  $v \neq i$  with  $c_v(A_i) \geq c_v(A_v) + 1$ .*

**Proof** Immediate from Theorem 29:  $\ell_{\mathcal{A}}(i) \leq \max(0, \max_{v \neq i} [c_v(A_i) - c_v(A_v)])$ , so  $\ell_{\mathcal{A}}(i) \geq 1$  forces the inner maximum to be at least 1.  $\square$

So the paid set is contained in the set of agents whose bundle looks strictly more expensive to somebody than that somebody's own — which is exactly condition (c) of Proposition 15 read backwards. An agent cannot be paid merely for holding a bundle she herself dislikes; somebody else has to corroborate.

**Remark 75** (Where the paid set sits, empirically). Across the 164 adversarial instances of Remark 133 that genuinely require a subsidy, the minimum-size paid set of the witness found is in every case contained in  $\arg \max_i c_i(A_i)$  — the paid agents are among those carrying the largest own cost.

**That lead is now withdrawn.** The 164 instances were not filtered for additivity, and the containment turns out to be an artefact of additive instances. Re-tested on 553 instances that are both envy-free-free *and* non-additive, it fails on 141 of them (update\_15/nonadditive\_residual.py). Own costs of different agents are measured by different functions, so there was never a reason to expect the comparison to mean anything; the earlier 164/164 simply reflected a sample that was effectively additive. What survives is Proposition 74, which is proved: a paid agent's bundle must look at least one unit more expensive to *somebody* than that somebody's own.

## 11 Approach 6: A Pure Goods Reformulation, and a Working Algorithm

Every approach so far has worked inside the chore model or the replica model. This one leaves both, and it ends in a concrete algorithm with no observed failures: a construct-and-repair procedure that has solved every one of 389,215 test instances tried against it, including every adversarial instance that defeated Approaches 1 through 5.

### 11.1 The transform is an involution

The size-shift transform of Theorem 16 is not merely a comparison device: it is an involution of the dichotomous class onto itself, so the chore problem can be restated as a question about *goods* with no chores, no replica, no coverage constraint and no schedule anywhere in it.

**Lemma 76.** *The map  $c \mapsto \tilde{v}$ ,  $\tilde{v}(S) := |S| - c(S)$ , is an involution of the set of dichotomous set functions (Theorem 7) onto itself.*

**Proof**  $\tilde{v}(\emptyset) = 0$ , and for  $g \notin S$  the marginal is  $\tilde{v}(S \cup \{g\}) - \tilde{v}(S) = 1 - (c(S \cup \{g\}) - c(S)) \in \{0, 1\}$ , so  $\tilde{v}$  is dichotomous. Applying the map twice returns  $|S| - (|S| - c(S)) = c(S)$ .  $\square$

## 11.2 The target

Fix an allocation  $\mathcal{A}$  and write  $\tilde{p}$  for the minimum subsidy vector of the goods instance  $\{\tilde{v}_i\}$  at  $\mathcal{A}$ , defined exactly when the chore instance is envy-freeable at  $\mathcal{A}$ , since by Theorem 16 the two envy graphs have the same cycle weights. Put  $q_i := \tilde{p}_i + |A_i|$  — goods subsidy plus bundle size.

**Conjecture 77** (Target G). *Every dichotomous goods instance admits an allocation whose vector  $q$  has spread at most 1.*

**Proposition 78.** *Conjecture 77 implies Conjecture 2.*

**Proof** By Theorem 16,  $\ell_{\mathcal{A}}(u) = \max_v [\tilde{d}_{\mathcal{A}}(u, v) + |A_u| - |A_v|]$ , where  $\tilde{d}_{\mathcal{A}}(u, v)$  is the heaviest  $u \rightarrow v$  path in the goods envy graph. Appending to any  $u \rightarrow v$  path a heaviest path leaving  $v$  gives  $\tilde{d}_{\mathcal{A}}(u, v) + \tilde{p}_v \leq \tilde{p}_u$ , so  $\tilde{d}_{\mathcal{A}}(u, v) \leq \tilde{p}_u - \tilde{p}_v$  and therefore  $\ell_{\mathcal{A}}(u) \leq q_u - \min_v q_v$ . A spread of at most 1 therefore gives  $\ell_{\mathcal{A}} \leq 1$ , and Observation 12 upgrades this to  $p^* \in \{0, 1\}^n$  with total at most  $n - 1$ .  $\square$

Target G quantifies over *all* allocations, and  $\tilde{p}$  is itself a longest-path weight, so Target G is a condition on paths rather than arcs — exactly what killed Approach 5 (Remark 69). It also survives extensive testing on its own: 0 failures on the full exhaustive  $n = m = 3$  family and on the adversarial families that refuted the invariant of Approach 5, detailed in `targetG.py` and `targetG_adv.py`. But Target G alone does not say how to *find* the allocation. The rest of this section does.

## 11.3 Item-by-item insertion is the wrong template here too

The obvious next move is to run the algorithm of [BKNS22] directly on the size-shifted goods instance  $\{\tilde{v}_i\}$ , since it is a legitimate dichotomous instance and the algorithm’s correctness proofs (Lemmas 3, 7, 8, 9, 10, 11 of that paper) are generic over *which* valid choice EXTEND and FINDSINK make at each step: any valid  $(\rho, \kappa)$  pair, any starting agent in  $M(p)$ . That freedom can be spent trying to steer bundle sizes toward balance while  $\tilde{p} \in \{0, 1\}^n$  is preserved automatically.

It does not work. Greedily steering the choice that locally minimises  $q$ -spread fails on 1,635 of the 9,880 exhaustive  $n = m = 3$  instances with a fixed item order, and on 440 of them even after trying every item order. Worse than a weak heuristic: on the smallest such instance, a full backtracking search over *every* legal execution of the algorithm — every item order, every EXTEND choice, every FINDSINK starting point, 474 states in total — never reaches spread  $\leq 1$ , although spread 0 is achievable (the perfectly balanced allocation  $\{0\}, \{1\}, \{2\}$ ) and is confirmed reachable by direct, algorithm-free enumeration. This is a genuine reachability obstruction, verified by `guidedR3_full.py` and `verify_reach_gap.py`,

in the same family as the peel dead ends of Section 8.5: item-by-item insertion fixes which items end up sharing a bundle the moment they are grouped, and no later permutation can split a bundle back apart, so an early bad grouping can be permanent.

The lesson is to stop building the allocation one item at a time. The rest of this section instead *constructs* a candidate directly and *repairs* it if needed — an operation insertion cannot perform, since repair means moving an item out of a bundle it is already in.

## 11.4 Construct and repair

**Definition 79** (Target G-bal). *A partition of the items into  $n$  groups is cardinality-balanced if the group sizes differ by at most 1. Target G-bal is the statement that every dichotomous goods instance admits a cardinality-balanced partition whose optimal (maximum-welfare) assignment of groups to agents gives  $q$ -spread at most 1.*

**Observation 80.** *Target G-bal implies Target G, since a cardinality-balanced partition is in particular a partition.*

Cardinality-balanced partitions are exactly what the balanced-matching machinery of [BDN<sup>+</sup>20, LMS26] is built to produce, so this restriction moves onto ground the literature already occupies. The algorithm below constructs one directly, using no chores, no replica, and no item-by-item insertion order.

### Algorithm (construct-and-repair).

1. *Construct.* Run one pass of R11’s iterated maximum-weight perfect matching (IMWPM): pad the items with inert dummies to a multiple of  $n$ , and in each of  $\lceil m/n \rceil$  rounds assign every agent exactly one remaining item (real or dummy), by a matching maximising the total *marginal* value  $\tilde{v}_i(A_i \cup \{g\}) - \tilde{v}_i(A_i)$  over the current partial bundles — the non-additive generalisation of the item-value weights of [BDN<sup>+</sup>20], needed because  $\tilde{v}_i$  need not be additive. Discard the dummies.
2. *Repair.* While the current partition’s optimal assignment has  $q$ -spread greater than 1: among all single-item swaps between two groups that keep the sizes cardinality-balanced, take the one that most reduces the ( $q$ -spread,  $-$ welfare) pair, re-optimize the assignment, and repeat.
3. *Restart on failure.* If repair reaches a local optimum with spread still above 1, restart step 1 from a freshly shuffled round-robin partition instead of IMWPM’s output, and repeat step 2. Try up to a small constant number of restarts.

Every step preserves cardinality balance, so termination at spread  $\leq 1$  directly witnesses Target G-bal — and hence, by Proposition 78 and Observation 80, produces an envy-free chore allocation with  $p^* \in \{0, 1\}^n$  once translated back through Theorem 16.

## 11.5 Evidence: zero failures across 389,215 instances

| Test set                                                                                                                                                              | Instances      | Failures |
|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------|----------|
| Exhaustive, $n = m = 3$ , all instances                                                                                                                               | 9,880          | 0        |
| Exhaustive, structured non-additive triples, $n = 3$ , $m \in \{3, 4, 5\}$                                                                                            | 378,176        | 0        |
| Every named hard instance in this project (Theorem 31’s witness, Theorem 59’s no-go instance, Proposition 67’s discrepancy counterexample, Proposition 34’s instance) | 4              | 0        |
| Randomised, $n \in \{3, \dots, 8\}$ , $m \in \{3, \dots, 10\}$ , endpoint-constant biased                                                                             | 1,155          | 0        |
| <b>Total</b>                                                                                                                                                          | <b>389,215</b> | <b>0</b> |

Two components of this deserve to be separated out, since the construction step and the repair step have different track records on their own.

*IMWPM alone*, with no repair, already succeeds on all 9,880 exhaustive  $n = m = 3$  instances and on all four hard instances, but fails on roughly 18% of larger randomised instances (132 of 720 tested) — it is greedy and short-sighted, so it can lock in a bad grouping in an early round exactly as insertion does, just less often.

*Repair from a bare round-robin start*, with no IMWPM warm start, reaches spread  $\leq 1$  on all 9,880 exhaustive instances and on 357,746 of the 357,760 structured triples at  $m = 5$ ; the remaining 14 are genuine local optima of the swap neighbourhood — confirmed by independent exhaustive search over all cardinality-balanced partitions, which finds spread 0 achievable on every one of them (`classify_failures.py`) — and every one of the 14 is recovered by restarting from a freshly shuffled partition, using at most 4 restarts (`restart_check.py`).

Combining both, as in Section 11.4, removed every failure found by either component alone, across every test run against it.

## 11.6 Complexity, and what is proved unconditionally

Three properties of the algorithm are theorems, not conjectures: its outputs are certified correct, it terminates in polynomial time, and its search space is genuinely the set of *unordered* balanced partitions. We state them with proofs, in the value-oracle model with per-query cost  $\tau$ .

**Lemma 81** (Assignment invariance). *Fix an unordered family of  $n$  bundles. The assignments of bundles to agents under which the allocation is envy-freeable are exactly the welfare-maximising ones, and every welfare-maximising assignment induces the same minimum subsidy on each bundle [KMS<sup>+</sup>24, Lemma 2]. Hence the multiset  $\{\tilde{p}(B) + |B|\}$ , and with it the  $q$ -spread, depends only on the unordered partition.*

**Proof** The first claim is condition (ii) of Theorem 10 applied to the goods instance: an assignment is envy-freeable iff no reassignment of its own bundles raises welfare, which

is exactly welfare maximality over assignments. The second is Lemma 2 of [KMS<sup>+</sup>24], which proves that the minimum subsidy vectors of any two maximum-weight assignments coincide bundle-by-bundle. Adding  $|B|$ , a bundle attribute, preserves this.  $\square$

Two consequences. First, the  $n!$  assignment loop in the implementation is mathematically equivalent to a single maximum-weight matching computation, so the complexity below is stated with the Hungarian algorithm,  $O(n^3)$ . Second, the algorithm’s true search space is unordered balanced partitions — the object Conjecture 89 should be read as quantifying over.

**Theorem 82** (Soundness, certified). *Suppose the algorithm outputs a partition and assignment with  $q$ -spread at most 1. Then the same allocation, read in the original chore instance — the size-shift acts on valuations only, so the allocation needs no translation — together with  $p := \ell_{\mathcal{A}}$  is an envy-free solution with  $p \in \{0, 1\}^n$  and total subsidy at most  $n - 1$ . Moreover the output is verifiable in  $O(n^2\tau + n^3)$  time independently of how it was found.*

**Proof** Proposition 78 gives  $\ell_{\mathcal{A}} \leq 1$  on the chore side; Observation 12 upgrades to  $p^* \in \{0, 1\}^n$ ; the endpoint of a heaviest path has  $\ell = 0$ , giving total at most  $n - 1$ . Verification recomputes the  $n^2$  bundle values ( $n^2$  oracle calls), one maximum-weight matching and one longest-path computation ( $O(n^3)$  each).  $\square$

**Theorem 83** (Termination and complexity). *On any dichotomous instance:*

- (i) *evaluating one partition costs  $O(n^2\tau + n^3)$ ;*
- (ii) *the construction pass (IMWPM) costs  $O(m^2\tau + nm^2)$ ;*
- (iii) *each repair iteration examines  $O(nm)$  moves and costs  $O(n^3m\tau + n^4m)$ ;*
- (iv) *the repair loop performs at most  $O(m^2)$  iterations, so with a constant number of restarts the whole algorithm runs in time  $O(n^3m^3\tau + n^4m^3)$ .*

*In particular the algorithm always terminates, in polynomial time, regardless of whether it succeeds.*

**Proof** (i) is the count in Theorem 82, using Lemma 81 to replace the assignment loop by one Hungarian computation. (ii): there are  $\lceil m/n \rceil$  rounds; each builds an  $n \times O(m)$  marginal table ( $O(nm)$  oracle calls, two per entry) and solves one rectangular assignment problem in  $O(n^2m)$ ; summing gives the bound. (iii): a move sends one of  $m$  items to one of  $n - 1$  other groups, and each candidate is evaluated by (i). (iv) is the substantive claim. Because marginals lie in  $\{0, 1\}$ , every bundle value satisfies  $\tilde{v}_i(B) \leq |B|$ ; hence total welfare is an integer in  $[0, m]$ , and every subsidy  $\tilde{p}_i$ , being a path weight, is an integer in  $[0, m]$ , so the spread is an integer in  $[0, m + 1]$ . The repair loop moves only when the pair (spread,  $-$ welfare) strictly decreases lexicographically, and that pair takes at most  $(m + 2)(m + 1)$  values, so at most  $O(m^2)$  moves occur before the loop halts — either at spread  $\leq 1$  or at a local optimum, where a restart or a failure report follows.  $\square$

The gap between these theorems and a full correctness proof is exactly Conjecture 89: soundness says the algorithm never lies, termination says it always answers, and what is missing is that the answer is always “success.” For two agents, that missing piece can be supplied.

## 11.7 A complete proof for two agents

**Theorem 84** (Completeness at  $n = 2$ ). *Every two-agent dichotomous goods instance admits a cardinality-balanced split with  $q$ -spread at most 1. Consequently Target G-bal, Target G, and Conjecture 2 all hold at  $n = 2$  — the last with the additional guarantee, not provided by the complementation route of Remark 9, that the two bundles differ in size by at most one.*

Fix dichotomous  $\tilde{v}_1, \tilde{v}_2$  on  $M$ ,  $|M| = m$ . For an ordered balanced split  $(A, B)$  with  $|A| = \lceil m/2 \rceil \geq |B|$ , write

$$\alpha := \tilde{v}_1(A) - \tilde{v}_1(B), \quad \beta := \tilde{v}_2(A) - \tilde{v}_2(B), \quad u := \tilde{v}_1 + \tilde{v}_2, \quad h := u(A) - u(B) = \alpha + \beta.$$

Call the split *good* if some envy-freeable orientation (assignment of  $A, B$  to the two agents) has  $q$ -spread at most 1.

**Lemma 85** (Swap step). *Exchanging one item of  $A$  with one item of  $B$  changes each of  $\tilde{v}_i(A)$  and  $\tilde{v}_i(B)$  by at most 1, hence each of  $\alpha, \beta$  by at most 2 and  $h$  by at most 4.*

**Proof** Removing an item changes a dichotomous value by 0 or  $-1$  and adding one by 0 or  $+1$ , so the composite changes each side's value by at most 1 in absolute value.  $\square$

**Lemma 86** (Window). *If  $m$  is even and  $|h| \leq 3$ , or  $m$  is odd and  $-1 \leq h \leq 3$ , then the split is good.*

**Proof** *Even  $m$ .* Sizes are equal, so  $q$ -spread equals  $|\tilde{p}_1 - \tilde{p}_2|$ , and it suffices to find an envy-freeable orientation in which both envies are at most 1. Giving  $A$  to agent 1 yields envies  $(-\alpha, \beta)$ , so that orientation is satisfactory iff  $\alpha \geq -1$  and  $\beta \leq 1$ ; the flipped orientation iff  $\alpha \leq 1$  and  $\beta \geq -1$ . Both fail only when  $\alpha, \beta \leq -2$  or  $\alpha, \beta \geq 2$ , and in either case  $|h| \geq 4$ . It remains to check envy-freeability. The two-cycle of the first orientation has weight  $\beta - \alpha$ ; if it is positive while that orientation is satisfactory ( $\alpha \geq -1, \beta \leq 1, \beta > \alpha$ ), then  $\alpha < \beta \leq 1$  and  $\beta > \alpha \geq -1$  show the flipped orientation is satisfactory as well, and its cycle  $\alpha - \beta$  is negative. So a satisfactory envy-freeable orientation exists, its subsidies lie in  $\{0, 1\}$ , and the spread is at most 1.

*Odd  $m$ .* Now  $|A| = |B| + 1$ , and with subsidies in  $\{0, 1\}$  the spread is at most 1 iff the holder of the large bundle needs no subsidy — both agents needing 1 is impossible, since then the two-cycle  $\geq 2$  would be positive. Giving  $A$  to agent 1 is satisfactory iff  $\alpha \geq 0$  and  $\beta \leq 1$ ; flipped, iff  $\beta \geq 0$  and  $\alpha \leq 1$ . Both fail only when  $\alpha, \beta \leq -1$  or  $\alpha, \beta \geq 2$ , i.e. when  $h \leq -2$  or  $h \geq 4$ . Envy-freeability is repaired exactly as in the even case: if the first satisfactory orientation has  $\beta > \alpha$ , then  $\alpha < \beta \leq 1$  gives  $\alpha \leq 0$  and  $\beta > \alpha \geq 0$  gives  $\beta \geq 0$ , so the flipped orientation is satisfactory with a negative cycle.  $\square$

**Proof of Theorem 84** *Even  $m$ .* Take any equal split, ordered so that  $h \geq 0$ . If  $h \leq 3$ , Lemma 86 finishes. Otherwise  $h \geq 4$ ; walk to the complementary split  $(B, A)$  by exchanging one paired item at a time ( $m/2$  swaps). The final value is  $-h \leq -4$ , and by Lemma 85 each step changes  $h$  by at most 4, so at the first step where the value drops below 4 it lies in  $[0, 3]$  — inside the window.

*Odd  $m$ .* The big sides are the  $(k+1)$ -subsets,  $m = 2k + 1$ . Suppose no split is good; then by Lemma 86 every big side  $A$  has  $h(A) \leq -2$  or  $h(A) \geq 4$ . Any two big sides are connected by single exchanges (repeatedly swap an element of the difference), and a step from  $h \geq 4$  to



$h \leq -2$  would change  $h$  by at least  $6 > 4$ , so *all* big sides lie on one side. If all have  $h \geq 4$ : let  $A^*$  minimise  $u$  over big sides, and pick any  $x \in A^*$ . Then  $B := (M \setminus A^*) \cup \{x\}$  is a big side with  $u(B) \leq u(M \setminus A^*) + 2 \leq u(A^*) - 2$ , contradicting minimality — the first inequality because  $u$  has marginals in  $\{0, 1, 2\}$ , the second because  $h(A^*) \geq 4$ . If all have  $h \leq -2$ : let  $A^\circ$  maximise  $u$  over big sides and pick any  $x \in A^\circ$ . Then  $u((M \setminus A^\circ) \cup \{x\}) \geq u(M \setminus A^\circ) \geq u(A^\circ) + 2$ , by monotonicity and  $h(A^\circ) \leq -2$ , contradicting maximality.  $\square$

**Corollary 87** (An exact  $O(m\tau)$  algorithm for two agents). *For  $n = 2$  the proof is constructive and needs no restarts: start from any balanced split; if it is not good, then for even  $m$  walk toward the complement, and for odd  $m$  first apply the jump  $A \mapsto (M \setminus A) \cup \{x\}$  — which maps  $h \geq 4$  to  $h \leq 0$  and  $h \leq -2$  to  $h \geq 2$  — and then, if the jump target is still outside the window, walk between the two splits. Some split within  $m/2 + 2$  evaluations lies in the window and is good. Each evaluation uses  $O(1)$  oracle calls, so the whole procedure runs in  $O(m\tau)$  time.*

**Proof** The even case is the walk in the proof above. For the odd jump, write  $B = (M \setminus A) \cup \{x\}$  with  $x \in A$ , so  $M \setminus B = A \setminus \{x\}$  and  $h(B) = u(B) - u(A \setminus \{x\})$ . Recall that  $u = \tilde{v}_1 + \tilde{v}_2$  has marginals in  $\{0, 1, 2\}$ , so adding or removing a single item changes  $u$  by at most 2. If  $h(A) \geq 4$  then  $u(B) \leq u(M \setminus A) + 2 \leq u(A) - 2$  while  $u(A \setminus \{x\}) \geq u(A) - 2$ , giving  $h(B) \leq 0$ . If  $h(A) \leq -2$  then  $u(B) \geq u(M \setminus A) \geq u(A) + 2$  while  $u(A \setminus \{x\}) \leq u(A)$ , giving  $h(B) \geq 2$ . Both bounds are attained, so neither can be sharpened. In either case the jump target is in the window  $[-1, 3]$  or strictly on the far side of it, and in the latter case the swap walk between  $A$  and  $B$  crosses the window as in the theorem: descending from  $\geq 4$ , the first value below 4 is at least 0; ascending from  $\leq -2$ , the first value above  $-2$  is at most 2.  $\square$

The theorem and the walk are machine-verified by `n2proof_check.py`: exhaustively over all instances at  $m \leq 4$  (490,545 pairs at  $m = 4$  alone) and on 30,000 random instances at  $m \in \{5, 6, 7\}$  — 0 failures of either, with the walk never needing more than 2 evaluations in practice against its  $m/2 + 2$  bound.

**Remark 88** (What the proof teaches the algorithm). Two findings from the proof effort feed back into the implementation. First, the moves that the  $n = 2$  proof walks are *pairwise exchanges*, which preserve group sizes exactly — but the implemented repair neighbourhood contains only single-item *moves*, which change two group sizes by one. When  $n$  divides  $m$ , every group has size  $m/n$  and every single move is rejected by the balance filter, so the repair loop is inert there and the implementation succeeds purely through IMWPM plus restarts; the zero failures at  $n = m$  and at  $(n, m) \in \{(3, 6), (6, 6), (7, 7), (8, 8)\}$  were all achieved in that degenerate mode. Adding exchanges to the neighbourhood is therefore both principled — they are the proof’s own moves — and necessary for the repair step to act at all when  $n \mid m$ . Second, the  $h$ -window technique suggests the right potential for a general proof: a one-dimensional aggregate whose swap-Lipschitz constant is smaller than the width of the region it must cross. At  $n = 2$  the aggregate  $u = \tilde{v}_1 + \tilde{v}_2$  works because the goodness condition is two-sided in a single difference; for  $n \geq 3$  goodness involves  $\binom{n}{2}$  pairwise differences, and no single aggregate is yet known to control them.

## 11.8 What remains

After Sections 11.6 and 11.7, the standing of the approach is: soundness, certification, termination and polynomial complexity are theorems for every  $n$ ; completeness is a theorem for  $n = 2$ ; and for  $n \geq 3$  completeness is the one remaining open statement, sharp and almost entirely combinatorial, with no chores, no replica and no scheduling left in it:

**Conjecture 89** (Construct-and-repair succeeds,  $n \geq 3$ ). *For every dichotomous goods instance, the algorithm of Section 11.4 — with pairwise exchanges added to the repair neighbourhood, per Remark 88 — reaches a cardinality-balanced partition with  $q$ -spread at most 1, using a number of restarts bounded by a constant independent of  $n$  and  $m$ .*

The 14 local optima found at  $n = 3$  show that the move neighbourhood does contain bad local optima, so a proof cannot simply be “every non-optimal balanced partition has an improving move.” What was observed instead is that bad local optima are escapable from other starting partitions, always within a handful of restarts. The  $n = 2$  proof suggests what to look for: a one-dimensional aggregate — there,  $u = \tilde{v}_1 + \tilde{v}_2$  with the window  $|h| \leq 3$  — whose per-move variation is smaller than the forbidden gap it must cross, together with an extremal argument at the ends. Finding the analogous aggregate for  $n \geq 3$ , where goodness involves  $\binom{n}{2}$  pairwise comparisons rather than one, is the whole remaining problem; solving it would upgrade Conjecture 89, and with it Conjecture 2, from extensively tested to proved.

## 11.9 The canonical aggregate, tested and refuted

The aggregate the previous paragraph asks for has an obvious candidate. At  $n = 2$  the quantity is  $h = u(A) - u(B)$  with  $u = \tilde{v}_1 + \tilde{v}_2$ , which is the *spread of  $u$  across the bundles*; for general  $n$  that reads

$$\text{uspread}(A) := \max_j u(A_j) - \min_j u(A_j), \quad u := \sum_i c_i,$$

and at  $n = 2$  it is  $|h|$  exactly, so this is a genuine generalisation rather than an analogy. Splitting Lemma 86 and the walk of Theorem 84 along the same seam gives two statements:

WINDOW( $K$ ) : every balanced partition with  $\text{uspread} \leq K$  is good;

EXISTS( $K$ ) : every instance has a balanced partition with  $\text{uspread} \leq K$ .

Together they give Conjecture 2, and at  $n = 2$  both hold with  $K = 3$  — that is precisely the content of the two-agent proof.

The measurement (`update_22/window_general.py`) first reproduces the proved case: at  $n = 2$  the smallest  $\text{uspread}$  carried by a *bad* balanced split is 4, i.e.  $|h| \leq 3$  suffices, matching Lemma 86 exactly. That agreement is what licenses reading the rest of the table as evidence about the same quantity rather than a different one.

For  $n \geq 3$  the two requirements separate and the route closes:

|            | smallest $\text{uspread}$ of a bad<br>balanced partition | largest $\text{uspread}$ forced<br>on some instance |
|------------|----------------------------------------------------------|-----------------------------------------------------|
| $n = 2$    | 4 (so WINDOW(3))                                         | 3 (so EXISTS(3))                                    |
| $n \geq 3$ | 2 (so $K \leq 1$ )                                       | 3 (so $K \geq 3$ )                                  |

At  $n = 3$ ,  $m = 7$  a bad balanced partition occurs with  $\text{uspread} = 2$ , forcing  $K \leq 1$ ; and 45 of the instances tested admit no balanced partition with  $\text{uspread}$  below 3, forcing  $K \geq 3$ . No  $K$  serves both. *The window shrinks below the achievable spread*, and this is exactly why  $n = 2$  is tractable and  $n \geq 3$  is not: there, both numbers are 3 and the interval is non-empty.

So the canonical aggregate is refuted. Any surviving aggregate must be sensitive to how welfare is distributed among agents, not only to its total —  $u$  discards precisely that information, and at  $n = 2$  it could afford to, because with two agents a single difference determines the split.

### 11.10 Sharpening what needs proving

Conjecture 89 allows a constant number of restarts, and the 14 local optima at  $n = 3$  show the repair neighbourhood really does have bad local optima, so the allowance is not vacuous. But a local optimum matters only if the algorithm can *reach* it, and the warm start is not an arbitrary partition.

**Observation 90** (Restarts are never consumed). *Over 1,580 random instances with  $n \leq 7$  and  $m \leq 9$ , the IMWPM warm start followed by repair alone succeeded every time; the restart branch was never entered (`update_22/restart_necessity.py`).*

This matters for what a proof must establish, because the two readings of Conjecture 89 are of quite different difficulty. Were restarts genuinely needed, completeness would be a *reachability* claim over starting partitions — which presupposes that a good partition exists and then adds that a shuffled round-robin start finds one, making it strictly harder than Conjecture 2 and useless as a route to it. Observation 90 indicates the weight sits elsewhere:

**Conjecture 91** (Warm-start descent). *For every dichotomous goods instance, repair from the IMWPM warm start reaches  $q$ -spread at most 1 without restarts.*

Conjecture 91 implies Conjecture 89 with zero restarts, hence Conjecture 2, and it is the better proof target for a specific reason: unlike the descent lemma of Section 12, whose witnessing move proved not to be local (Section 12.6), it does not quantify over all partitions. It concerns a *single, structured* starting point — the output of an iterated maximum-weight matching — and the properties of that output are available to the proof. Whether the bad local optima are reachable from it is the question; the data say they are not.

### 11.11 What the warm start guarantees, and the two claims that remain

Asking what the IMWPM output looks like, rather than only whether repair eventually succeeds, turns out to answer the question almost completely.

**Observation 92** (The warm start never exceeds spread two). *Across every family tested — the full exhaustive  $n = m = 3$  family, all 357,760 triples of the structured non-additive pool at  $m = 5$ , sampled triples at  $m = 6$ , random instances up to  $n = 7$ ,  $m = 10$ , disjoint-interest instances built for maximal disagreement, and the named hard instances of Section 11.5 — the  $q$ -spread of the raw IMWPM partition, before any repair, never exceeded 2 (`update_23/warmstart.py`, `imwpm_bound.py`).*

Two refinements make this sharper than it first appears. First, the bound is rarely attained: on the  $m = 5$  pool, 352,468 of 357,760 instances start at spread at most 1, so IMWPM *alone* already produces a witness for Target G-bal and repair is never invoked. The discrepancy counterexample, the insertion witness and the W4 no-go instance — the three named obstructions of this section — all start at spread at most 1. Second, when spread 2 does occur, repair is immediate:

**Observation 93** (One move suffices). *Among 3,000 instances whose IMWPM start has  $q$ -spread exactly 2, repair reached spread at most 1 in exactly one improving move in every case, with no stalls.*

Together these decompose Conjecture 91 into two claims, neither of which is a descent along an unbounded chain:

**Conjecture 94** (Warm-start bound). *For every dichotomous goods instance, the IMWPM partition has  $q$ -spread at most 2.*

**Conjecture 95** (Last rung). *Every IMWPM partition with  $q$ -spread exactly 2 admits a balance-preserving single-item move to  $q$ -spread at most 1.*

**Proposition 96.** *Conjectures 94 and 95 together imply Conjecture 91, hence Conjecture 89 with zero restarts, hence Conjecture 2 together with its polynomial-time clause.*

**Proof** By Conjecture 94 the warm start has spread at most 2. If it is at most 1 the partition is already a Target G-bal witness. Otherwise it is exactly 2, and Conjecture 95 supplies a move to spread at most 1; the move preserves cardinality balance by hypothesis. Either way a cardinality-balanced partition with  $q$ -spread at most 1 is produced, which is Target G-bal; Observation 80 and Proposition 78 give Conjecture 2. The running time is that of Theorem 83 with the repair loop executing at most one iteration.  $\square$

The reason this is a better position than anything in Sections 12 and 13 is that Conjecture 94 is not a statement about local search at all. It is a statement about what one round-based matching construction produces, and constructions of exactly this kind are what [BDN<sup>+</sup>20] analyses: for *additive* goods, iterated maximum-weight matching is shown there to need a subsidy of at most one per agent, which combined with cardinality balance is precisely a  $q$ -spread bound of 2. So Conjecture 94 asks for that analysis to survive the passage from additive to dichotomous values, where the matching weights are marginals  $\tilde{v}_i(A_i \cup \{g\}) - \tilde{v}_i(A_i)$  rather than fixed item values. This is the first point in this report at which the remaining gap abuts a theorem that is already proved in the literature rather than a fresh combinatorial obstruction, and it is where effort should go.

Conjecture 95 is the smaller of the two and the more likely to yield to the two-tier machinery of Section 10: spread exactly 2 means the agents split into a top and a bottom tier differing by 2, and what must be shown is that one item can always be moved across.

## 11.12 Reading [BDN<sup>+</sup>20] line by line: what transfers and what does not

Conjecture 94 asks whether [BDN<sup>+</sup>20, Theorem 1.3] — iterated maximum-weight matching needs a subsidy of at most one per agent — survives the passage from additive to dichotomous valuations. The two classes are *incomparable*:  $v(S) = \min(|S|, 1)$  is dichotomous and not

additive, an additive function with an item of value 5 is not dichotomous, and they intersect exactly in the binary additive functions. So nothing is inherited for free, and the question is which parts of the proof survive.

**Where additivity is used.** In [BDN<sup>+</sup>20] it appears in five places. Lemma 3.1 (envy-freeability), Lemma 3.2 (EF1), and equations (1) and (7) all rest on the per-round decomposition  $v_i(A_k) = \sum_t v_i(\mu_k^t)$ ; Claim 4.4, equation (9), invokes additivity by name. These are serious but conceivably routable. The fatal use is the *engine* of their Section 4: the modified profile  $\bar{v}_i$  is *defined by assigning item values*,

$$\bar{v}_i(\mu_k^t) = \max(v_i(\mu_k^t), v_i(\mu_i^{t+1})),$$

and extended additively. A non-additive dichotomous function admits no such extension, so the construction has no direct analogue at all.

**What transfers.** Their Lemma 4.1 — *if  $\mathcal{A}$  is envy-freeable and every arc of its envy graph has weight at least  $-1$ , the subsidy is at most one per agent* — is proved by closing a maximum-weight path into a cycle and using that the cycle has non-positive weight. It is pure envy-graph combinatorics and uses no additivity whatever, so it holds verbatim in our setting. Moreover  $\bar{v}$  enters their argument only through the  $n^2$  numbers  $\bar{v}_i(A_k)$ . That suggests performing the modification on the *matrix* rather than on a valuation: set

$$a_{ik} := \bar{v}_i(A_k), \quad \bar{a}_{ik} := \max(a_{ik}, a_{ii} - 1),$$

so that  $\bar{a}_{ii} = a_{ii}$ , every modified arc is at least  $-1$ , and every modified arc dominates the true one, whence  $\bar{\ell} \geq \ell$  pointwise. If the modified graph has no positive cycle, Lemma 4.1 gives  $\bar{\ell} \leq 1$  and hence  $\ell \leq 1$  — with no additivity used anywhere and no valuation function to construct.

**Why this does not settle it.** The condition “the modified matrix has no positive cycle” implies  $\ell \leq 1$ , so it cannot hold on any instance where  $\ell \geq 2$ ; observing that it fails exactly on those instances is forced by the logic and is not evidence. The measurement is nevertheless decisive, because of what it turned up on the way.

**Proposition 97** (The dichotomous analogue of [BDN<sup>+</sup>20, Thm 1.3] is false). *There are dichotomous goods instances on which the IMWPM allocation requires a subsidy of 2 for some agent. Over 1,034 instances, 2 had  $\max_i \tilde{p}_i = 2$ ; a witness has  $n = 3$ ,  $m = 7$ , bundles  $\{g_1, g_4, g_6\}$ ,  $\{g_0, g_3, g_5\}$ ,  $\{g_2\}$  and  $\tilde{p} = (0, 0, 2)$ .*

So “subsidy at most one per agent” does *not* survive the passage to dichotomous values, and Conjecture 94 cannot be proved by proving that. This also refutes the natural decomposition — prove  $\tilde{p} \leq 1$ , then add cardinality balance to get  $q$ -spread at most 2 — on both legs, because the second leg fails too:

**Proposition 98** (IMWPM does not produce balanced bundles). *When  $m$  is not a multiple of  $n$  the dummies discarded by IMWPM are distributed unevenly, and the resulting bundle sizes need not differ by at most one. Over 1,034 instances, 79 had size spread 2.*

What holds the bound together is instead a *compensation*: in the witness of Proposition 97 the bundle sizes are  $(3, 3, 1)$ , so  $q = \tilde{p} + (|A_1|, |A_2|, |A_3|) = (3, 3, 3)$  and the  $q$ -spread is 0 even though the subsidy is 2. The agent carrying the large subsidy is exactly the agent holding the small bundle. Neither  $\tilde{p}$  nor the bundle size is controlled on its own; only their sum is. Any proof of Conjecture 94 must therefore reason about  $q$  directly, and the two-term split that makes the additive proof work is unavailable.

### 11.13 What $q$ is: the chores-side formula

The compensation observed in Proposition 97 — subsidy 2 against sizes  $(3, 3, 1)$ , giving  $q = (3, 3, 3)$  — is not a coincidence. It follows from an identity that removes the goods picture from this section altogether.

**Lemma 99** (The  $q$ -formula). *Let  $\mathcal{A}$  be envy-freeable and let  $d(i, j)$  denote the heaviest  $i \rightarrow j$  path in the chores envy graph, with  $d(i, i) := 0$ . Then*

$$q_i = \max_{j \in N} [d(i, j) + |A_j|].$$

**Proof** By Lemma 76,  $\tilde{v}_i(S) = |S| - c_i(S)$ , so the goods arc weight is

$$\tilde{w}(i, j) = \tilde{v}_i(A_j) - \tilde{v}_i(A_i) = [c_i(A_i) - c_i(A_j)] + |A_j| - |A_i| = w_{\mathcal{A}}(i, j) + |A_j| - |A_i|.$$

Along a path  $i = u_0 \rightarrow u_1 \rightarrow \dots \rightarrow u_r = j$  the size terms telescope, so  $\tilde{d}(i, j) = d(i, j) + |A_j| - |A_i|$ ; both heaviest-path functions are well defined because the two envy graphs have the same cycle weights (Theorem 16). Taking the maximum over  $j$ , and noting that the empty path contributes the term  $j = i$ ,  $\tilde{p}_i = \max_j [d(i, j) + |A_j|] - |A_i|$ . Adding  $|A_i|$  gives the claim.  $\square$

So  $q_i$  is the largest bundle size reachable from agent  $i$ , inflated by the envy accumulated on the way. The compensation is then forced rather than fortuitous: an agent holding a small bundle can still carry a large  $q$  by envying its way to a big bundle, and an agent holding a big bundle has  $q_i \geq |A_i|$  regardless of its subsidy. Neither term is controlled alone because  $q$  was never a sum of two independent quantities.

**Corollary 100** ( $q$  at an IMWPM allocation). *IMWPM runs  $T = \lceil m/n \rceil$  rounds giving every agent exactly one object per round, so  $|A_j| = T - \delta_j$  where  $\delta_j$  is the number of dummies agent  $j$  received. Hence*

$$q_i = T + \max_j [d(i, j) - \delta_j],$$

and the  $q$ -spread does not depend on  $T$ . In particular, when  $n \mid m$  every  $\delta_j = 0$  and  $q_i = T + \ell_{\mathcal{A}}(i)$ , so on those instances the  $q$ -spread equals  $\max_i \ell_{\mathcal{A}}(i)$  exactly.

Both identities were checked against direct computation on 745 instances with no mismatch, and the last claim held on all 276 instances with  $n \mid m$  (update\_26/q\_formula.py). Corollary 100 is worth stating because it removes the apparent asymmetry between this section and the rest of the report: when  $n$  divides  $m$ , Target G at an IMWPM allocation is Conjecture 2 at that allocation, and Conjecture 94 says exactly  $\max_i \ell_{\mathcal{A}}(i) \leq 2$ .

One more translation is worth recording, since it states the algorithm without any reference to goods at all.



**Lemma 101** (IMWPM in chores terms). *Let  $\gamma_i^t(g) := c_i(A_i^{t-1} \cup \{g\}) - c_i(A_i^{t-1}) \in \{0, 1\}$  be the chore marginal. The goods marginal used as a matching weight satisfies  $\mu_i^t(g) = 1 - \gamma_i^t(g)$ , so maximising  $\sum_i \mu_i^t$  is the same as minimising  $\sum_i \gamma_i^t$ . IMWPM is therefore: in each round, hand one remaining chore to every agent so as to minimise the total marginal cost incurred that round.*

**Proof**  $\mu_i^t(g) = \tilde{v}_i(A_i^{t-1} \cup \{g\}) - \tilde{v}_i(A_i^{t-1}) = 1 - \gamma_i^t(g)$  directly from  $\tilde{v}_i(S) = |S| - c_i(S)$ . Summing over the  $n$  agents turns  $\max \sum_i \mu_i^t$  into  $\min \sum_i \gamma_i^t$ .  $\square$

Lemma 101 also yields the per-round exchange property in a usable form: for any two agents  $i, j$  matched to  $g_i, g_j$  in round  $t$ , optimality against the transposition gives  $\gamma_i^t(g_j) + \gamma_j^t(g_i) \geq \gamma_i^t(g_i) + \gamma_j^t(g_j)$ . Since all four terms lie in  $\{0, 1\}$ , this says: *if agent  $i$  would rather have had  $j$ 's chore, then  $j$  strictly prefers its own*. That is a clean one-round invariant; what Section 11.17 shows is that accumulating it across rounds is exactly what fails.

### 11.14 From paths to arcs

Lemma 99 localises Conjecture 94. If  $j^*$  attains  $q_i$  then  $q_{i'} \geq d(i', j^*) + |A_{j^*}|$ , so  $q_i - q_{i'} \leq d(i, j^*) - d(i', j^*)$  and therefore

$$q\text{-spread} \leq 2 \iff d(i, j) - d(i', j) \leq 2 \text{ for all } i, i', j.$$

Because  $d(i, j) \geq w_{\mathcal{A}}(i, i') + d(i', j)$  whenever  $i \neq i'$  — with no positive cycle a heaviest walk may be taken simple, so concatenation is legitimate — this forces every arc to have weight at most 2. The measurement returns something stronger.

**Observation 102** (Arcs at an IMWPM allocation). *Over 746 instances with  $n \leq 6$ ,  $m \leq 9$ , every chores envy arc at the IMWPM allocation had weight at most 1; the observed distribution of  $w_{\mathcal{A}}(i, j)$  was  $-3:3, -2:648, -1:3551, 0:2356, 1:300$  (`update_27/arc_bound.py`).*

An arc bound of 1 says  $c_i(A_i) \leq c_i(A_j) + 1$ , which — since every marginal is at most 1 — is exactly the EF1 property for chores. Two things follow. First, the lower bound fails badly: arcs reach  $-3$ , so the hypothesis of [BDN<sup>+</sup>20, Lemma 4.1] is unavailable, which is precisely the situation that paper describes before introducing its modified valuation, and which Section 11.12 transplanted to the matrix level. Second, an arc bound alone cannot give a path bound: with arcs at most 1, a path of length  $k$  still admits weight  $k$ . Indeed  $\max_j$ -spread of  $d(\cdot, j)$  reached 3, and  $\max_i \ell_{\mathcal{A}}(i)$  reached 2 (Proposition 97).

What makes Observation 102 worth pursuing is that it is a *pairwise* statement, and the round structure supplies a pairwise tool.

**Lemma 103** (Round dominance). *Let  $g$  be a chore available at the start of round  $t$  that the round- $t$  assignment does not use. Then for every agent  $i$ ,  $\gamma_i^t(g_i^t) \leq \gamma_i^t(g)$ .*

**Proof** By Lemma 101 the round- $t$  assignment minimises  $\sum_k \gamma_k^t$  over all ways of handing one available chore to each agent. Replacing agent  $i$ 's chore  $g_i^t$  by  $g$  is such a way, since  $g$  is available and unused, and it changes only the  $i$ -th term. Minimality gives  $\gamma_i^t(g) \geq \gamma_i^t(g_i^t)$ .  $\square$

**Corollary 104.** *If  $\gamma_i^t(g_i^t) = 1$  then every chore available and unused at round  $t$  has marginal 1 for agent  $i$  over  $A_i^{t-1}$ .*

Both were verified exactly: 0 violations in 6,719 checks over 520 instances. Together with the transposition bound already noted — if agent  $i$  would rather have had  $j$ 's chore then  $j$  strictly prefers its own — Lemma 103 gives a complete description of what one round guarantees, and Corollary 104 is the sharp form: a round in which  $i$  pays is a round in which *nothing available was free for  $i$* .

This is the ingredient that the additive proofs obtain from item values and that Section 11.17 showed cannot be recovered by summing rounds. Whether it can be accumulated *along an envy path* — where Lemma 99 says the quantity of interest actually lives — is the open question, and it is now a question about consecutive agents on a path rather than about all  $T$  rounds at once.

### 11.15 The own-cost decomposition, and a constant bound

Turning Observation 102 into a path bound needs a potential: some  $\phi$  with  $w_{\mathcal{A}}(i, j) \leq \phi_j - \phi_i$  on every arc makes path weights telescope, giving  $\max_i \ell_{\mathcal{A}}(i) \leq \text{spread}(\phi)$ . Taking  $\phi = -\ell_{\mathcal{A}}$  works trivially and is circular, so  $\phi$  must come from the algorithm. Nine candidates read off the round structure were tried and none holds universally; the best,  $\phi = -c_i(A_i)$ , held on 583 of 657 instances (`update_28/potential_search.py`). But that candidate is worth keeping, because the algebra behind it is exact.

**Lemma 105** (Own-cost decomposition). *Write  $\kappa_i := c_i(A_i)$  and  $E_{ij} := c_i(A_j) - c_j(A_j)$ , the excess of  $j$ 's bundle to  $i$  over its cost to its own holder. Then for every arc,*

$$w_{\mathcal{A}}(i, j) = (\kappa_i - \kappa_j) - E_{ij},$$

*and consequently, for every directed path  $P$  from  $a$  to  $b$ ,*

$$w_{\mathcal{A}}(P) = \kappa_a - \kappa_b - \sum_{(i,j) \in P} E_{ij}.$$

**Proof**  $w_{\mathcal{A}}(i, j) = c_i(A_i) - c_i(A_j) = \kappa_i - (\kappa_j + E_{ij})$  by the definition of  $E_{ij}$ . Summing along  $P$ , the  $\kappa$  terms telescope.  $\square$

The identity is exact and holds for every allocation, not only IMWPM's; it was verified with 0 violations over all simple paths of 657 instances. It splits the path weight into a telescoping part, governed by the spread of the agents' own costs, and a defect part. The potential condition  $w_{\mathcal{A}}(i, j) \leq \phi_j - \phi_i$  for  $\phi = -\kappa$  is exactly  $E_{ij} \geq 0$ , i.e. "every bundle is cheapest for the agent holding it". That fails — but only just.

**Observation 106** (Both parts are bounded at an IMWPM allocation — FALSE at larger  $m$ , see Remark 108). *Over 657 instances with  $n \leq 6$ ,  $m \leq 9$  (`update_28/potential_search.py`) and the accompanying path sweep):*

- (i)  $\text{spread}(\kappa) \leq 2$  always, with distribution 0:200, 1:454, 2:3;
- (ii)  $E_{ij} \geq -1$  always — all 80 negative excesses, out of 8,826, equalled exactly  $-1$ ;
- (iii)  $\sum_{(i,j) \in P} E_{ij} \geq -1$  for every simple path  $P$ , with distribution  $-1:74$ ,  $0:569$ ,  $1:14$ .

Part (iii) is the substantive one: individual excesses go negative and paths are long, yet the negative excesses *do not chain*. Combining it with (i) through Lemma 105:

**Corollary 107** (A constant bound, conditionally — WITHDRAWN, see Remark 108). *If (i) and (iii) of Observation 106 hold in general, then at every IMWPM allocation  $\max_i \ell_{\mathcal{A}}(i) \leq 3$ .*

**Proof** By Lemma 105,  $w_{\mathcal{A}}(P) = \kappa_a - \kappa_b - \sum_P E_{ij} \leq \text{spread}(\kappa) + 1 \leq 3$ .  $\square$

**Remark 108** (Both hypotheses of Corollary 107 are false). Observation 106 rested on  $n \leq 6$ ,  $m \leq 9$ , and its part (i) was attained on only 3 of 657 instances — thin evidence, since  $\kappa_i$  counts *rounds* in which agent  $i$  paid and that sweep barely left  $T = 3$ . Re-run at  $m$  up to 12, and on families built to make the agents disagree (disjoint interest blocks, nested interest sets, one agent finding everything costly), both parts fail:

$$\begin{array}{ll} \text{spread}(\kappa) \text{ reached } 4 & (n = 3, m = 12, \text{ nested}), \\ \sum_P E_{ij} \text{ reached } -4 & (n = 3, m = 12, \text{ one-heavy}), \end{array}$$

with distributions  $\text{spread}(\kappa)$  given by 0:58, 1:44, 2:18, 3:31, 4:5 and path defect by  $-4:4$ ,  $-3:21$ ,  $-2:14$ ,  $-1:12$ , 0:105 (`update_29/stress_kappa.py`). So Corollary 107 yields at best 8 at these sizes, and both of its inputs grow with  $m$ . It is withdrawn.

What the same run shows is more interesting than what it refutes: across all 156 of those larger instances,  $\max_i \ell_{\mathcal{A}}(i)$  never exceeded 1 — distribution 0:129, 1:27 — which is *stronger* than Conjecture 94 requires and stronger than the small- $m$  sweeps, where  $\ell_{\mathcal{A}} = 2$  did occur. Since Lemma 105 is an exact identity, the two terms  $\kappa_a - \kappa_b$  and  $\sum_P E_{ij}$  must therefore be *cancelling*: both reach 4 in magnitude while their difference stays at 1.

**Remark 109** (The recurring obstruction, third instance). This is now the third decomposition of the same quantity in which both parts are unbounded and only the combination is controlled. Proposition 97 found it for  $q_i = \tilde{p}_i + |A_i|$ , where the subsidy reached 2 while the bundle size compensated exactly. Proposition 113 found it for the round decomposition, where the per-round defect accumulated to  $-4$  while the true path weight stayed at 2. And Remark 108 finds it for the own-cost decomposition. In each case the identity is exact and the bound on the whole is far better than any bound on the parts. The lesson is not that a particular split was unlucky: it is that  $\ell_{\mathcal{A}}$  at an IMWPM allocation appears not to admit *any* two-term decomposition with separately bounded parts, and a proof must treat  $\kappa_a - \kappa_b - \sum_P E_{ij}$  as a single object.

## 11.16 The closing arc: the one path-as-a-unit tool, and why it fails

Remark 109 says a proof may not split  $\ell_{\mathcal{A}}$  into separately bounded parts. Exactly one classical tool respects that: close the heaviest path into a cycle. If  $P$  is a heaviest path from  $a$  ending at  $b$ , then  $P + (b, a)$  is a cycle of non-positive weight, so

$$\ell_{\mathcal{A}}(a) \leq -w_{\mathcal{A}}(b, a) = c_b(A_a) - c_b(A_b). \quad (\text{CL})$$

This is the mechanism behind Theorem 29 and behind [BDN<sup>+</sup>20, Lemma 4.1], and it is strictly weaker than the uniform balance refuted in Section 10: uniform balance asks  $w_{\mathcal{A}}(v, u) \geq -1$

for *all* pairs, whereas (CL) needs it only for the single pair  $(b, a)$  with  $b$  ending  $a$ 's heaviest path.

Note first that “closing arc  $\geq -1$ ” is *equivalent* to  $\ell_{\mathcal{A}}(a) \leq 1$ , since  $\ell_{\mathcal{A}}(a) \geq 2$  forces  $-w_{\mathcal{A}}(b, a) \geq 2$  by (CL). So it cannot be tested as an independent hypothesis; its interest is as a proof target, being a claim about one pair rather than a path. The question that decides the route is therefore how much (CL) loses.

**Proposition 110** ((CL) is too lossy at an IMWPM allocation). *Over 452 instances with  $n \leq 6$ ,  $m \leq 12$ , on uniform, disjoint-interest and nested-interest families, the 210 pairs  $(a, b)$  with  $b$  ending a heaviest path out of  $a$  satisfied:*

- closing arcs  $w_{\mathcal{A}}(b, a)$  distributed  $-3:18, -2:103, -1:89$  — reaching  $-3$ ;
- slack  $-w_{\mathcal{A}}(b, a) - \ell_{\mathcal{A}}(a)$  distributed  $0:89, 1:103, 2:18$ , so (CL) is strictly lossy on 58% of the pairs and loses up to 2.

Consequently (CL) certifies only  $\ell_{\mathcal{A}}(a) \leq 3$  where the truth is 1, and proving “closing arc  $\geq -1$ ” is a strictly harder statement than proving  $\ell_{\mathcal{A}} \leq 1$  (`update_30/closing_arc.py`).

So the route is worse than what it replaces, and the reason is instructive: endpoint pairs are the *worst* pairs, not the best. Across the same run the arc distribution over all ordered pairs was  $-4:5, -3:94, -2:619, -1:1846, 0:1709, 1:273$ , so a typical arc is  $-1$  or  $0$  while every endpoint arc is at most  $-1$ . That is forced by what  $b$  is: the agent everyone envies towards, hence the agent for whom other bundles look most expensive.

**Remark 111** (Fourth instance of the cancellation). Proposition 110 is Remark 109 again in a fourth guise. (CL) bounds  $\ell_{\mathcal{A}}$  by a single quantity, and that quantity is systematically more extreme than  $\ell_{\mathcal{A}}$  itself. Together with the subsidy/size split, the round decomposition and the own-cost decomposition, every attempt to bound  $\ell_{\mathcal{A}}$  by anything other than  $\ell_{\mathcal{A}}$  has now come in strictly worse — and the closing arc is the sharpest such attempt available, since it is the only one that treats the path as a unit. Whatever keeps  $\ell_{\mathcal{A}}$  small at an IMWPM allocation is not visible in any single arc, any single round, or any two-term split of the path.

The endpoint agents are nevertheless highly structured, which is worth recording even though it did not rescue the route: of the 210 endpoints, 100% had  $\ell_{\mathcal{A}}(b) = 0$  (by construction), 94.3% had the minimum own cost  $c_b(A_b)$  among all agents, 83.8% paid in no round at all, and 73.8% held a smallest bundle.

### 11.17 Auditing [LMS26]: the missing hypothesis is submodularity

IMWPM is taken from [LMS26], whose own analysis had not been read against our setting. Two things had to be checked: whether that paper covers non-additive valuations, and whether its proof technique transplants.

The first is settled by its Section 2, which fixes  $u_i(S) = \sum_{j \in S} u_i(j)$ : the paper is entirely additive. But its Proposition 3.1 recovers the bound of [BDN<sup>+</sup>20] by an argument far simpler than the modified-valuation construction audited in Section 11.12. For a path  $P = (1, \dots, k)$  with terminal agent  $k$ , put  $F_t := \max(\{0\} \cup \{u_k(g) : g \in J^t\})$ , the best item still available to  $k$ . Maximum-weight optimality of the round- $t$  matching, compared against the

alternative that shifts each path item one step back and hands  $k$  the best remaining item, gives  $w_t(P) \leq F_t - F_{t+1}$ ; summing telescopes to  $w_{\mathcal{A}}(P) \leq F_1 \leq 1$ .

Additivity enters this proof in exactly two places, both of the form “sum the rounds”:

$$(A) \quad u_i(A_i) = \sum_t u_i(\mu_i^t), \quad (B) \quad u_i(A_j) = \sum_t u_i(\mu_j^t).$$

**Lemma 112** ((A) survives with marginal weights). *Our IMWPM matches on marginals, so writing  $A_i^t$  for agent  $i$ ’s bundle after round  $t$ ,*

$$\tilde{v}_i(A_i) = \sum_t \left[ \tilde{v}_i(A_i^t) - \tilde{v}_i(A_i^{t-1}) \right] = \sum_t \mu_i^t(\text{own item}),$$

*an exact telescoping identity requiring no additivity whatever.*

So (A) costs nothing, and (B) is the entire gap. What the argument actually needs is only one direction of it,

$$(B') \quad \tilde{v}_i(A_j) \leq \sum_t \left[ \tilde{v}_i(A_i^{t-1} \cup \{\mu_j^t\}) - \tilde{v}_i(A_i^{t-1}) \right],$$

and (B') is precisely *submodularity*: for a submodular  $v$  one has  $v(S \cup T) - v(S) \leq \sum_{g \in T} [v(S \cup \{g\}) - v(S)]$ , which is (B') read along the rounds. Dichotomous functions need not be submodular, and the failure is not marginal.

**Proposition 113** (The defect, and why the transplant fails). *Write  $D_{ij} := \sum_t [\tilde{v}_i(A_i^{t-1} \cup \{\mu_j^t\}) - \tilde{v}_i(A_i^{t-1})] - \tilde{v}_i(A_j)$ , so that (B') is the claim  $D_{ij} \geq 0$ , and*

$$w_{\mathcal{A}}(P) \leq 1 - \sum_{i \in P} D_{i,i+1}.$$

*Over 800 instances, (A) failed 0 times, while (B') failed on 2,103 of 6,660 ordered pairs — 31.6% — with individual defects reaching  $-2$ . Accumulated along a path the defect reached  $-4$ , so the transplanted bound is  $w_{\mathcal{A}}(P) \leq 5$  at  $n \leq 5$ , and in general degrades as  $1 + 2(n - 1)$ .*

The observed maximum path weight is 2. So the technique of [LMS26], carried across as far as it will go, gives a bound of  $1 + 2(n - 1)$  where the truth appears to be 2: no better than the  $n - 1$  per agent already available from [KMS<sup>+</sup>24], and the route yields nothing. The defect accumulates along the path, which is exactly the failure mode that “sum the rounds” arguments cannot survive.

**Remark 114** (Why this corner of the  $2 \times 2$  is the open one). Sections 11.12 and this one identify the same missing hypothesis from two directions. [BDN<sup>+</sup>20] needs additivity to *build* a modified valuation from item values; [LMS26] needs it to *decompose* a bundle’s value across rounds, and what would suffice there is submodularity rather than additivity outright. Dichotomous cost functions lie outside both: they are not additive, and  $c(S) = \min(|S|, 1)$  shows they need not be submodular. The dichotomous *goods* corner is settled in [BKNS22] precisely because binary submodular valuations are matroid rank functions, where (B') is free. Our corner has neither hypothesis available, and Proposition 113 measures how far from free (B') actually is. This is a quantitative account of why the fourth corner has stayed open, rather than a restatement that it is hard.

**Remark 115** (Correction to Section 11.4). Proposition 98 corrects the claim made when the algorithm was stated, that “every step preserves cardinality balance, so termination at spread  $\leq 1$  directly witnesses Target G-bal”. The repair loop does reject moves that leave the partition unbalanced, but it never runs when the warm start already has  $q$ -spread at most 1, and that warm start may itself be unbalanced. Of 1,190 successful runs, 62 terminated at a partition whose bundle sizes differ by 2. The algorithm therefore certifies Target G, not Target G-bal. Nothing is lost for Conjecture 2, which follows from Target G alone by Proposition 78; but the evidence recorded in Section 11.5 should be read as evidence for Target G, and Definition 79 is not what the algorithm delivers.

## 12 Approach 7: descent on a lexicographic potential

### 12.1 The reformulation that suggests it

Section 10 characterises good allocations by the two-tier condition:  $\mathcal{A}$  is good precisely when there is a bipartition  $S \subseteq N$  with

$$c_i(A_i) - [i \in S] \leq c_i(A_j) - [j \in S] \quad \text{for all } i \neq j.$$

Writing  $\delta_j := [j \in S]$  and  $\hat{c}_i(A_j) := c_i(A_j) - \delta_j$ , this says exactly that *every agent holds a bundle minimising its discounted cost*, where the discount is 1 on the paid slots and 0 elsewhere and depends only on the slot, not on the agent. Unpacked by tier:

- within  $S$ , and within  $N \setminus S$ : exact envy-freeness;
- an unpaid agent *strictly* prefers its own bundle to every paid agent’s bundle;
- a paid agent envies an unpaid agent by at most 1.

Two consequences are worth stating before any search is set up. First,  $S = N$  and  $S = \emptyset$  impose *identical* constraints, since the two indicators cancel; this is the structural reason the paid set is never all of  $N$ , and it is the combinatorial shadow of Proposition 13. Second, the existential over subsidy vectors has become an existential over bipartitions, of which there are  $2^n$  — the first formulation in which it is finite.

### 12.2 The search space and the potential

By Theorem 10(ii) an allocation is envy-freeable iff it maximises welfare among reassignments of its own bundles. So we search over *partitions* and let the assignment be determined:

**Definition 116** (Canonical form). *For a partition  $\mathcal{B} = (B_1, \dots, B_n)$  of  $M$ , the canonical allocation  $\mathcal{A}(\mathcal{B})$  assigns bundles to agents by a minimum-cost perfect matching. It is envy-freeable, so  $\ell_{\mathcal{A}(\mathcal{B})}$  is finite.*

**Definition 117** (The potential). *For a partition  $\mathcal{B}$  with canonical allocation  $\mathcal{A}$ , set*

$$\Psi(\mathcal{B}) := \left( \max_{i \in N} \ell_{\mathcal{A}}(i), \sum_{i \in N} |B_i|^2 \right) \in \mathbb{Z}_{\geq 0} \times \mathbb{Z}_{\geq 0},$$

*ordered lexicographically. A transfer moves one chore from one bundle to another.*



The conjecture this approach rests on is a statement about single transfers.

**Conjecture 118** (Descent lemma). *Let  $\mathcal{B}$  be a partition whose canonical allocation has  $\max_i \ell_{\mathcal{A}}(i) \geq 2$ . Then some transfer  $\mathcal{B}'$  satisfies  $\Psi(\mathcal{B}') < \Psi(\mathcal{B})$ .*

### 12.3 The reduction, which is unconditional

**Theorem 119** (Descent suffices). *Conjecture 118 implies Conjecture 2 in full: every instance with negative dichotomous valuations admits an envy-free solution  $(\mathcal{A}, p)$  with  $p \in \{0, 1\}^n$  and total subsidy at most  $n - 1$ , and one is computable in time  $O(m^4 n^4)$  given value-oracle access.*

**Proof** Start from any partition and pass to canonical form. While  $\max_i \ell_{\mathcal{A}}(i) \geq 2$ , apply Conjecture 118 to obtain a transfer strictly decreasing  $\Psi$ .

*Termination.*  $\Psi$  takes values in  $\mathbb{Z}_{\geq 0} \times \mathbb{Z}_{\geq 0}$  under the lexicographic order, which is a well-order, so no infinite strictly decreasing sequence exists. Quantitatively, every arc weight is at most  $m$  in absolute value and  $\ell_{\mathcal{A}}$  is a weight of a simple path, so the first coordinate lies in  $\{0, 1, \dots, m\}$ ; the second lies in  $\{m, \dots, m^2\}$ , since  $\sum_i |B_i|^2$  is minimised by an equal split and maximised by putting everything in one bundle. Hence  $\Psi$  has at most  $(m + 1)(m^2 + 1) = O(m^3)$  values and the loop runs  $O(m^3)$  times.

*Correctness.* The loop can only exit with  $\max_i \ell_{\mathcal{A}}(i) \leq 1$ . By Observation 12 the minimum subsidy vector is integral and non-negative, so  $p^* \in \{0, 1\}^n$ ; by Theorem 11 the pair  $(\mathcal{A}, p^*)$  is an envy-free solution; and by Proposition 13 the total is at most  $n - 1$ .

*Complexity.* One iteration examines  $O(mn)$  transfers. Each requires a minimum-cost perfect matching on  $n$  agents,  $O(n^3)$  by the Hungarian method with  $O(n^2)$  oracle calls, followed by an all-pairs longest-path computation on an  $n$ -vertex graph,  $O(n^3)$ . So an iteration costs  $O(mn \cdot n^3) = O(mn^4)$  and the whole procedure  $O(m^4 n^4)$ .  $\square$

Theorem 119 is the reason this approach is worth the space: unlike every earlier route, the gap between what is proved and Conjecture 2 is a *single* statement about *single* transfers, and closing it delivers the polynomial-time clause for free rather than as a separate construction.

### 12.4 Why this potential and not the obvious one

The first potential tried was  $\Psi_1 = (\max_i \ell_{\mathcal{A}}(i), \#\{i : \ell_{\mathcal{A}}(i) = \max\}, \sum_i c_i(B_i))$ , minimising total cost as a tie-break. It fails, and the witness explains the failure rather than merely exhibiting it. At  $n = 3, m = 4$  with all four chores in one bundle and the other two empty,

$$\mathcal{B} = (\{g_1, g_2, g_3, g_4\}, \emptyset, \emptyset), \quad \ell_{\mathcal{A}} = (0, 0, 2), \quad \sum_i c_i(B_i) = 2,$$

every transfer leaves  $\max_i \ell_{\mathcal{A}}(i) = 2$  and *raises* the total cost. The reason is that dichotomous costs cap: the loaded agent already values its bundle at 2, and removing one chore does not reduce that, while the recipient's cost rises. So the third coordinate pins the search at a utilitarian optimum — precisely the class shown in Section 11 not to be good in general. The tie-break must reward *spreading chores out*, not minimising their total cost. Both  $\sum_i |B_i|^2$  and  $\max_i |B_i|$  do;  $\sum_i \ell_{\mathcal{A}}(i)$  in place of the maximum does not.

| potential                                        | tie-break        | stuck partitions |
|--------------------------------------------------|------------------|------------------|
| $(\max \ell, \# \text{at max}, \sum_i c_i(B_i))$ | total cost       | 449              |
| $(\max \ell, \max_i  B_i )$                      | largest bundle   | 160              |
| $(\sum_i \ell_i, \sum_i  B_i ^2)$                | total envy first | 464              |
| $(\max \ell, \# \text{at max}, \max_i  B_i )$    | —                | 0                |
| $(\max \ell, \# \text{at max}, \sum_i  B_i ^2)$  | —                | 0                |
| $\Psi = (\max \ell, \sum_i  B_i ^2)$             | —                | 0                |

$\Psi$  is the shortest surviving potential, which is why Definition 117 takes it.

## 12.5 Where the remaining burden sits

The second coordinate of  $\Psi$  moves in a completely explicit way, and this splits Conjecture 118 into two unequal halves.

**Lemma 120** (Balance dichotomy). *Transferring a chore from a bundle of size  $s$  to a bundle of size  $t$  changes  $\sum_i |B_i|^2$  by exactly  $2(t - s + 1)$ . Hence the transfer strictly decreases the second coordinate of  $\Psi$  if and only if  $t \leq s - 2$ . Consequently, if  $\mathcal{B}$  is balanced — all bundle sizes within 1 of one another — then no transfer decreases the second coordinate, and Conjecture 118 demands a transfer that strictly decreases  $\max_i \ell_{\mathcal{A}}(i)$  itself.*

**Proof** The sizes change from  $(s, t)$  to  $(s - 1, t + 1)$ , so the sum changes by  $(s - 1)^2 + (t + 1)^2 - s^2 - t^2 = 2(t - s + 1)$ , which is negative exactly when  $t < s - 1$ . If all sizes lie in  $\{q, q + 1\}$  then  $t \geq q \geq s - 1$  for every ordered pair, so no transfer qualifies.  $\square$

So Conjecture 118 is really two claims. On unbalanced partitions one must find a size-equalising transfer that does not *increase* the maximum subsidy; on balanced partitions one must strictly decrease the maximum subsidy by a single transfer. The second is the hard half and the natural next target: it is a statement about balanced partitions only, where every bundle has size  $\lfloor m/n \rfloor$  or  $\lceil m/n \rceil$ , which is a far more rigid object than an arbitrary allocation. Note that the transfer is followed by re-matching, so a single transfer may permute the assignment substantially — this is where the strength of the move comes from, and any proof must use it.

## 12.6 The obstruction: the witnessing transfer is not local

Conjecture 118 asserts that *some* transfer works. A proof wants a *rule* — a named transfer, definable from the envy structure, that always works. The obvious candidates all fail, and the way they fail is informative.

Two mechanisms can reduce an arc  $w_{\mathcal{A}}(i, j) = c_i(A_i) - c_i(A_j)$ : remove a chore from  $A_i$  to lower  $c_i(A_i)$ , or add one to  $A_j$  to raise  $c_i(A_j)$ . Under dichotomous costs *both can fail outright*, because costs cap and marginals need not be monotone:

- $c(S) = \min(|S|, 1)$  has no chore whose removal lowers the cost of a two-element bundle;
- $c(S) = \max(|S| - 1, 0)$  has no chore whose addition raises the cost of the empty bundle.

Note that the second example also rules out the submodular shortcut: a chore may have marginal 1 over a superset of  $A_j$  and marginal 0 over  $A_j$  itself, so “some chore raises  $c_i(A_j)$ ” cannot be inferred from  $c_i$  being large somewhere. This is the same failure of the exchange property that has obstructed every earlier approach in this report.

Five rules were tested against all 7,416 crux partitions — balanced, with  $\max_i \ell_{\mathcal{A}}(i) \geq 2$  — under two criteria: the  $\Psi$  criterion (strictly reduce  $\max_i \ell_{\mathcal{A}}(i)$ ) and the weaker  $P_3$  criterion (reduce  $(\max_i \ell_{\mathcal{A}}(i), \# \text{at max})$  lexicographically).

| rule                                   | transfer                                                | $\Psi$ criterion | $P_3$ criterion |
|----------------------------------------|---------------------------------------------------------|------------------|-----------------|
| R1                                     | argmax- $\ell$ bundle $\rightarrow$ max-path endpoint   | 2,856            | 3,168           |
| R2                                     | argmax- $\ell$ bundle $\rightarrow$ its cheapest bundle | 3,827            | 4,021           |
| R3                                     | argmax- $\ell$ bundle $\rightarrow$ anywhere            | 5,112            | 5,152           |
| R4                                     | anywhere $\rightarrow$ max-path endpoint                | 6,039            | 6,467           |
| R5                                     | largest bundle $\rightarrow$ smallest bundle            | 3,972            | 4,086           |
| <i>every transfer</i> (Conjecture 118) |                                                         | 7,416            | 7,416           |

The decisive row is R3. In 2,264 of 7,416 crux partitions, no transfer *out of the envious agent’s own bundle* works, even under the weaker criterion — the witnessing move involves two bundles neither of which belongs to an agent attaining the maximum. So the move that repairs the envy is not local to the envy. Any proof of Conjecture 118 must therefore be non-constructive, or must name a rule of a kind not yet imagined; the five natural ones are exhausted.

**Remark 121** (The descent lemma is strictly stronger than what is needed). Conjecture 118 asserts more than Conjecture 2: it says both that good allocations exist and that greedy descent on  $\Psi$  finds one, i.e. that  $\Psi$  has no local minimum with  $\max_i \ell_{\mathcal{A}}(i) \geq 2$ . Conjecture 2 is the statement that the *global* minimum of  $\Psi$  has first coordinate at most 1. The extra strength is what buys the polynomial-time clause in Theorem 119 at no cost, but it also means this route is harder than the target, and the obstruction above may be a symptom of exactly that. A proof of Conjecture 2 alone need not go through any local search.

## 12.7 Evidence

Conjecture 118 was tested by enumerating *every* partition of every test instance, computing  $\Psi$  at each and at all  $O(mn)$  of its transfers, and recording any partition with  $\max_i \ell_{\mathcal{A}}(i) \geq 2$  admitting no strict decrease.

| family                                                      | instances | stuck partitions |
|-------------------------------------------------------------|-----------|------------------|
| random, $n \leq 4, m \leq 6$                                | 1,732     | 0                |
| EF-free only, $n \leq 5, m \leq 6$                          | 253       | 0                |
| R10 Set-Splitting, certified hard                           | 4         | 0                |
| random, $n = 3, m \leq 9$                                   | 798       | 0                |
| random, $n = 4, m \leq 7$                                   | 280       | 0                |
| random, $n = 5, m \leq 6$                                   | 75        | 0                |
| <i>crux partitions only</i> (balanced, $\max \ell \geq 2$ ) | 7,416     | 0                |

The last row is the one that matters, and it was measured separately because the others do not imply it: “zero stuck partitions over all partitions” would be consistent with the crux population being empty. It is not. It is empty when  $m = n$  — a balanced partition then gives every agent one chore, and the min-cost matching caps  $\ell_{\mathcal{A}}$  at 1, so there is nothing to prove — and it grows with  $m$ , reaching 990 crux partitions at  $n = 3, m = 9$  and 2,112 at  $n = 4, m = 6$ . Every one of them admits a  $\Psi$ -decreasing transfer.

The large- $m$  rows matter specifically: a size-based potential should fail when bundles grow, since  $\sum_i |B_i|^2$  can only be decreased by equalising sizes while a large bundle may be cheap under a capped cost function. No such failure appears up to  $m = 9$ . Separately, no instance in any of these families lacked a good allocation, so Conjecture 2 itself remains unrefuted throughout.

## 13 Approach 8: the global route

### 13.1 Why global

Approach 7 reduced Conjecture 2 to a descent lemma, but Remark 121 records that the descent lemma is strictly stronger than the conjecture: it forbids bad *local* minima of  $\Psi$ , whereas Conjecture 2 only concerns the *global* minimum. The global route drops the extra strength. It asks:

is there an objective whose global optimum is always a good allocation?

Such a statement would give existence with no reference to local search, and the polynomial-time clause becomes a separate question about computing the optimum. All objectives below are evaluated over partitions in canonical form (Definition 116), so every candidate is envy-freeable and  $\ell_{\mathcal{A}}$  is finite.

### 13.2 What is refuted

For an objective  $f$ , two questions differ: are *all*  $f$ -optima good (a usable theorem, since then any optimum may be returned), or is merely *some*  $f$ -optimum good (which needs a tie-break before it is a theorem)? Over 1,760 instances at  $n \leq 5, m \leq 7$ :

| objective                 |             | all optima good | some optimum good |
|---------------------------|-------------|-----------------|-------------------|
| $\sum_i c_i(B_i)$         | utilitarian | 389 fail        | 50 fail           |
| sorted cost vector        | leximax     | 54 fail         | 21 fail           |
| $\max_i c_i(B_i)$         | egalitarian | 382 fail        | 0                 |
| $\sum_i  B_i ^2$          | balance     | 68 fail         | 0                 |
| $\max_i c_i - \min_i c_i$ | spread      | 201 fail        | 0                 |

The utilitarian row restates the refutation of Section 11. The leximax row is new and worth recording, because leximax is the objective the goods literature reaches for first: on 21 instances *every* leximax optimum is bad, so no tie-break can rescue it. Egalitarian, balance and spread all survive in the weaker sense.

A tie-break chain would upgrade a “some” to a theorem, and twelve were tried. None succeeds:

| chain                                        | failures | chain                                                           | failures |
|----------------------------------------------|----------|-----------------------------------------------------------------|----------|
| bal.                                         | 75       | bal. $\rightarrow$ lex.                                         | 4        |
| bal. $\rightarrow$ util.                     | 21       | bal. $\rightarrow$ lex. $\rightarrow$ util.                     | 4        |
| bal. $\rightarrow$ egal.                     | 30       | bal. $\rightarrow$ egal. $\rightarrow$ lex. $\rightarrow$ util. | 4        |
| bal. $\rightarrow$ egal. $\rightarrow$ util. | 8        | egal. $\rightarrow$ bal.                                        | 48       |
| bal. $\rightarrow$ util. $\rightarrow$ egal. | 8        | egal. $\rightarrow$ bal. $\rightarrow$ util.                    | 9        |
| bal. $\rightarrow$ spread                    | 15       | bal. $\rightarrow$ #at max $\rightarrow$ util.                  | 29       |

Adding components helps only down to 4 failures out of 1,082 and then stops; the last three chains are progressively finer and all fail on the same residue. So the good allocations are not picked out by any natural cost-based ordering. Note also that appending “# agents at maximum cost” actively *breaks* the objective, turning 0 “some” failures into 21.

### 13.3 What survives: the balance lemma

The balance objective is special because its optima are describable outright:  $\sum_i |B_i|^2$  is minimised exactly by the *balanced* partitions, those with all bundle sizes within 1 of one another, i.e. every bundle of size  $\lfloor m/n \rfloor$  or  $\lceil m/n \rceil$ . So its “some optimum is good” row reads as a structural restriction:

**Conjecture 122** (Balance lemma). *Every instance with negative dichotomous valuations admits a good allocation whose partition is balanced.*

**Proposition 123.** *Conjecture 122 implies Conjecture 2, except for its polynomial-time clause.*

**Proof** Immediate: a good allocation is one with  $\max_i \ell_{\mathcal{A}}(i) \leq 1$ , which by Observation 12 and Proposition 13 yields  $p^* \in \{0, 1\}^n$  with total at most  $n - 1$ .  $\square$

**Remark 124** (Relation to Target G-bal, and a correction). Conjecture 122 was arrived at here by scanning global objectives, but the restriction to cardinality-balanced partitions is not new to this section: it is Definition 79, introduced in Section 11 on the goods side of the size-shift involution. The two are not the same statement, and the direction matters. Proposition 78 bounds  $\ell_{\mathcal{A}}(u)$  by  $q_u - \min_v q_v$  in one direction only, so Target G-bal implies Conjecture 122 but not conversely — and the gap is real, not merely unproved. Over the balanced partitions enumerated in update\_22, 666 are good while carrying  $q$ -spread at least 2, whereas none has  $q$ -spread at most 1 without being good. So Conjecture 122 is a strict weakening of Target G-bal, and to that extent a better target; but the idea of restricting to balanced partitions was already in hand, and this section should not be read as introducing it.

Conjecture 122 is stronger than Conjecture 2, and it is worth being explicit about why that strengthening is of a more useful kind than the one in Approach 7. The descent lemma strengthens the conjecture by constraining the *optimisation landscape* — it must have no bad local minimum — which is a statement about every partition. The balance lemma

strengthens it by constraining the *witness*, and a constrained witness is extra structure a proof may use: on a balanced partition every bundle has one of two sizes,  $m$  is nearly determined by  $n$  and the common size, and the pigeonhole arguments that drive Example 1 and the small-bundle theorem of Section 5 apply directly.

| family                       | instances | no good balanced allocation |
|------------------------------|-----------|-----------------------------|
| random, $n \leq 5, m \leq 7$ | 1,760     | 0                           |
| random, $n \leq 5, m \leq 8$ | 1,082     | 0                           |
| EF-free only                 | 30        | 0                           |
| balanced partitions examined |           | 466,560                     |

The two approaches also fit together rather than competing. The second coordinate of  $\Psi$  in Definition 117 is exactly  $\sum_i |B_i|^2$ , so descent under  $\Psi$  drives the partition towards balance; Conjecture 122 asserts that the balanced region is where the good allocations are. That the empirically successful potential and the empirically successful global objective agree on the same quantity is weak evidence that balance, not cost, is the operative structure.

### 13.4 The induction the balance lemma invites, and why it does not close

Restricting the witness suggests building it one chore at a time. Three statements of increasing strength were tested; the first holds, and the two that would prove anything fail.

**Proposition 125** (Incremental structure, empirical). *Write (INC) for: there is a good balanced allocation  $\mathcal{A}$ , with its assignment fixed, and an ordering  $g_1, \dots, g_m$  of the chores such that  $\mathcal{A}$  restricted to  $\{g_1, \dots, g_k\}$  is good for every  $k$ . Over 1,125 instances, (INC) never failed — not even when the assignment was forbidden to change between steps.*

Restrictions of dichotomous cost functions to a subset of chores are again dichotomous, so each prefix is a legitimate smaller instance and (INC) has exactly the shape of an induction on  $m$ . It is not one. (INC) *presupposes* a good balanced allocation and then asserts it can be reached; it says nothing about existence, and the induction does not close, because the good allocation supplied by the hypothesis on  $m - 1$  chores need not be the restriction of any good allocation on  $m$ .

The statement that would prove existence has to be self-contained:

(NS) if  $\mathcal{A}$  is a good allocation of a proper subset  $T$  of the chores, some chore outside  $T$  can be added to some bundle leaving it good.

(NS) implies Conjecture 2 by induction from the empty allocation, which is good since all costs are 0, and gives an  $O(m^2 n^4)$  algorithm with no backtracking. It is false. Over 502,528 good partial allocations, 846 admit no good one-chore extension, and 150 resist even when re-matching is permitted. The first witness is

$$A = (\emptyset, \{g_1, g_2, g_3\}, \emptyset), \quad g_4 \text{ unassigned},$$

three chores piled on one agent with two agents empty — a state a balanced construction would never reach. That suggests the repair: keep the partial allocation balanced, inserting



only into minimum-size bundles, which preserves balance and would prove the balance lemma and Conjecture 2 together.

**Proposition 126** (The balanced induction fails, and fails worse). *Write (NSB) for (NS) restricted to balanced partial allocations with insertion into a minimum-size bundle. Over 236,421 good balanced partial allocations, 11,723 are stuck — 5.0%, against 0.17% for the unrestricted (NS). Allowing insertion into any balance-preserving bundle rather than a minimum-size one changes nothing: the same 11,723 states are stuck.*

So balance does not rescue the induction; it makes it worse, by removing exactly the freedom — piling chores onto one agent for a step — that lets greedy insertion escape. Note that this does not contradict Proposition 125: a good insertion order exists, but it cannot be found by extending an arbitrary good balanced partial allocation. The order must be chosen with foresight.

**Remark 127** (The pattern across Approaches 7 and 8). Three independent constructive routes have now failed in the same way. The descent lemma is true in every test but its witnessing transfer is not local to the envy (Section 12.6); (INC) is true in every test but the insertion order is not locally determined; and no lexicographic objective picks out the good allocations. In each case the object exists and no rule computable from local data finds it. That is weak evidence that Conjecture 2, if true, is not provable by exhibiting a construction, and that the search should turn to non-constructive existence arguments — counting, averaging, or a fixed-point or topological argument — with the algorithmic clause pursued separately through Section 11, whose algorithm is polynomial by construction and has never failed.

### 13.5 The non-local route: is there room, and can anything reach it?

Remark 111 closes the local and compositional routes: four independent attempts to bound  $\ell_{\mathcal{A}}$  by a part of itself all came in strictly worse than the truth. What remains is the non-local class — counting, averaging, probabilistic-method and parity arguments, which reason about all allocations at once. Two questions decide whether that class is usable: whether the good allocations are numerous enough to be reached by such an argument, and whether any natural distribution or identity actually reaches them.

**Observation 128** (There is room). *Exhaustively over all  $n^m$  allocations of 1,169 instances with  $n \leq 5$ ,  $m \leq 6$ , the number of good allocations never fell below 11, and on the 80 EF-free instances — the residual class, since Conjecture 2 is trivial where an exactly envy-free allocation exists — never below 27 (update\_31/counting\_room.py).*

So the good set is not a needle: an argument that merely shows it non-empty has something to aim at, unlike the situation where only one witness exists and nothing short of exhibiting it can work. But the two standard ways of reaching it both fail.

**Proposition 129** (Parity fails). *The number of good allocations has no fixed parity: over 737 instances it was even 469 times and odd 268 times. No Sperner- or Tucker-style identity forcing an odd count can hold.*

**Proposition 130** (No natural distribution concentrates). *The fraction of allocations that are good is not bounded below. Under the uniform distribution its minimum fell from 0.136 at  $n = 3$ ,  $m = 4$  to 0.025 at  $n = m = 5$ . Restricting to cardinality-balanced partitions does not repair this — the minimum fraction there was 0.042 at  $n = m = 4$ , worse than uniform — and restricting to minimum-cost allocations gives a minimum fraction of 0, reconfirming Section 11’s refutation of that family.*

The probabilistic method needs a distribution under which the good set has probability bounded away from zero, and Proposition 130 rules out the three natural candidates; a counting proof needs an identity forcing the count non-zero, and Proposition 129 rules out the parity identity. What is *not* ruled out is a cleverer weighting: the fractions above are for specific distributions, and the space of weightings is unbounded. But no weighting suggested by the structure of the problem survives, and Observation 128 shows the target is generous enough that if a natural one existed it would likely have been among these.

**Remark 131** (Status of the technique classes). Taken with Remark 111, every named technique class has now been tried and refuted on this problem: constructive (Sections 12.6, 13.4), decompositional (Remark 109, four instances), transplanted from the additive literature (Sections 11.12, 11.17), and now non-local. This is not a claim that Conjecture 2 is false — it survives every test at every size — nor that it is unprovable. It is a claim about where a proof will not come from, and it is the most useful thing this report currently establishes.

### 13.6 The algorithmic clause

Conjecture 2 also asks for a polynomial-time algorithm, and Proposition 123 deliberately excludes it. The position is as follows.

- Minimising  $\sum_i |B_i|^2$  is trivial — every balanced partition is optimal — but that is exactly why the balance lemma gives no algorithm on its own: it does not say *which* balanced partition is good, and there are exponentially many.
- The objectives that would name one are not obviously tractable. Minimising  $\max_i c_i(B_i)$  is a makespan-type problem and NP-hard already for identical additive costs, so “compute the egalitarian optimum” is not a polynomial algorithm even where it is correct.
- Two polynomial candidates remain and are unaffected by this section: the algorithm of Section 11, whose completeness is Conjecture 89 and which has never failed in 389,215 instances; and descent on  $\Psi$ , which Theorem 119 shows runs in  $O(m^4 n^4)$  and is correct given Conjecture 118.

So the division of labour is now explicit: the global route is the cleanest line on *existence*, and the algorithmic clause rests on Approach 6 or Approach 7. A proof of Conjecture 122 would settle the mathematical content of Conjecture 2 and leave a well-posed, separately attackable computational question — which is a better position than the current one, where existence and computation are entangled in a single conjecture that has resisted both.

## 14 Computational Evidence

All searches operate directly on the graph-theoretic form of Conjecture 2: enumerate allocations, build the envy graph, detect positive-weight cycles and compute  $\ell_{\mathcal{A}}$  by a Bellman–Ford iteration on maximum-weight walks, and record  $\min_{\mathcal{A}} \max_i \ell_{\mathcal{A}}(i)$ .

| Experiment                                                                                     | Script                  | Result                                                                                                 |
|------------------------------------------------------------------------------------------------|-------------------------|--------------------------------------------------------------------------------------------------------|
| Exhaustive: $n = m = 3$ , all 9,880 instances up to agent symmetry                             | <code>search.py</code>  | worst-case minimum per-agent subsidy = 1                                                               |
| Randomised: $n \in \{3, 4, 5\}$ , $m \in \{4, 5, 6\}$ , $\approx 3.5 \times 10^4$ instances    | <code>fast.py</code>    | no counterexample                                                                                      |
| Binary additive construction of Theorem 19, $2 \times 10^5$ instances, $n \leq 5$ , $m \leq 7$ | <code>binadd.py</code>  | maximum per-agent subsidy ever observed = 1                                                            |
| The obstruction of Theorem 31                                                                  | <code>deadend.py</code> | all three insertions require subsidy 2; full instance solvable with 0                                  |
| Does <i>some</i> utilitarian-optimal allocation attain $\leq 1$ ?                              | <code>msw.py</code>     | yes, in all but one of $\approx 2.6 \times 10^4$ instances                                             |
| Does <i>every</i> utilitarian-optimal allocation attain $\leq 1$ ?                             | <code>msw.py</code>     | no — fails in $\approx 2.5\%$ of instances at $n = 3, m = 4$ , rising to $\approx 5\%$ at $m = 5$      |
| The single exception                                                                           | <code>mswcex.py</code>  | refutes Conjecture 18 (Proposition 34)                                                                 |
| Peel schedule and trap-state refusal on the witness of Theorem 31                              | <code>peel.py</code>    | six legal peels reach the zero-subsidy allocation; the trap state has $\ell(1) = 2$                    |
| Full peel state space of that witness, $7^3 = 343$ states, both move types                     | <code>peel3.py</code>   | 175 legal, 166 good, <b>9 dead ends</b> , root good, all six reachable terminals are perfect matchings |

and, for the approaches attempted after the replica transform,

| Experiment                                                                                                                                                                                      | Script                  | Result                                                                                                             |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------|--------------------------------------------------------------------------------------------------------------------|
| No dichotomous reweighting separates: all $990^3$ triples on the instance of Theorem 59, plus the separable $P$ -family on three instances                                                      | dupsep.py               | <b>0</b> separating triples (Theorem 60)                                                                           |
| Type-ordered peel schedules, witness plus 890 random instances, $n \in \{3, 4\}$ , $m \leq 4$                                                                                                   | typeorder.py, stress.py | every instance admitting a legal schedule admits a type-ordered one; arg max recipient rule fails (Proposition 53) |
| Double-refinement selection rule, $\approx 1050$ random instances                                                                                                                               | rules.py                | no failure — but see the next row                                                                                  |
| The same rule on the instance of Proposition 34                                                                                                                                                 | checkD.py               | <b>fails:</b> selected allocation needs 2, another needs 0 (Remark 37)                                             |
| Iterated minimum-marginal perfect matching, non-additive costs                                                                                                                                  | rules.py                | fails on 13/250 at $n = 3, m = 4$ and 10/120 at $n = 3, m = 5$                                                     |
| Uniformly balanced families: exhaustive $n = m = 3$ , then $\approx 1900$ random instances up to $n = 5$                                                                                        | routeA.py               | no failure, and the implication chain verified on every instance                                                   |
| The same, hunted adversarially over structured families                                                                                                                                         | routeA_adv.py           | 10 <b>failures</b> at $n = 3, m = 4$ , plus the certified-hard set-splitting instance (Proposition 67)             |
| Arc weights of the good allocations on that counterexample                                                                                                                                      | routeA_cex.py           | all 19 have an arc $\leq -2$ ; minimum $-3$ (Proposition 68)                                                       |
| <b>Conjecture 2</b> itself, hunted adversarially over non-additive instances: exhaustive structured triples at $m = 3, 4$ plus $\approx 8.4 \times 10^3$ endpoint-constant instances to $n = 5$ | twotier.py              | <b>no counterexample</b> (Remark 132)                                                                              |
| How often a subsidy is needed at all                                                                                                                                                            | twotier.py              | exactly envy-free in 98.6% of 11,884 instances (Remark 133)                                                        |
| Structure of the paid set on the 164 instances that need one                                                                                                                                    | hardcases.py            | $\min  S  \leq 3$ , always inside $\arg \max_i c_i(A_i)$ (Remark 75)                                               |
| The $n = 3$ characterisation, against the true longest-path computation                                                                                                                         | n3check.py              | 0 mismatches over $5.8 \times 10^5$ (instance, allocation) pairs (Proposition 71)                                  |

and, for the construct-and-repair algorithm of Section 11,

| Experiment                                                                       | Script                                                      | Result                                                                                                                                                                               |
|----------------------------------------------------------------------------------|-------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Item-by-item insertion on the transformed goods instance, greedily steered       | <code>guidedR3.py</code> ,<br><code>guidedR3_test.py</code> | 1,635/9,880 (fixed order),<br>440/9,880 (best of 6)                                                                                                                                  |
| Full backtracking over every legal insertion execution on the smallest failure   | <code>guidedR3_full.py</code>                               | spread 2 is the best reachable; spread 0 is unreachable                                                                                                                              |
| Independent, algorithm-free confirmation that spread 0 is achievable there       | <code>verify_reach_gap.py</code>                            | confirmed — a genuine reachability obstruction                                                                                                                                       |
| IMWPM warm start alone, no repair                                                | <code>imwpm_raw.py</code>                                   | exhaustive $n = m = 3$ and all hard instances pass; $\approx 18\%$ of larger random instances fail                                                                                   |
| Swap-based repair from a round-robin start alone                                 | <code>targetGbal_local.py</code>                            | 9,880/9,880 exhaustive; 357,746/357,760 structured triples at $m = 5$                                                                                                                |
| The 14 residual local optima, classified                                         | <code>classify_failures.py</code>                           | all 14 are local-search artefacts — exhaustive balanced search finds spread 0 on every one                                                                                           |
| The same 14, retried with shuffled restarts                                      | <code>restart_check.py</code>                               | 14/14 recovered, at most 4 restarts each                                                                                                                                             |
| <b>The combined algorithm:</b> IMWPM warm start, swap repair, restart on failure | <code>final_algorithm.py</code>                             | <b>0 failures across 389,215 instances</b> — exhaustive $n = m = 3$ , exhaustive structured triples to $m = 5$ , every named hard instance, and randomised sweeps to $n = 8, m = 10$ |
| The $n = 2$ completeness theorem and its $O(m\tau)$ constructive walk            | <code>n2proof_check.py</code>                               | 0 failures over 521,310 instances — exhaustive $m \leq 4$ (490,545 pairs at $m = 4$ ), random $m \in \{5, 6, 7\}$ ; walk never exceeded 2 evaluations (Theorem 84, Corollary 87)     |

The exhaustive row is the only unconditional evidence for Conjecture 2. The dichotomous functions on  $m$  items number 2, 6, 38, 990 for  $m = 1, 2, 3, 4$ . Agents are interchangeable, so an instance is a multiset of three of them, and at  $m = 3$  there are  $\binom{38+3-1}{3} = \binom{40}{3} = 9,880$  of these — feasible to enumerate, and already out of reach at  $m = 4$ .

The last three rows are the ones that changed the shape of the problem. The utilitarian-optimal set almost always contains a witness, but not always, and one exception is enough: Proposition 34 is built from the single instance in which it did not.

## 14.1 Why the randomised evidence is weak

The count of  $\approx 3.5 \times 10^4$  instances without a counterexample overstates the support for Conjecture 2, and the reason is structural rather than a matter of sample size.

The generator builds a dichotomous function by visiting subsets in order of increasing size and, at each subset  $S$ , computing the interval  $[\ell, h]$  of values consistent with what has already been fixed — namely  $\ell = \max_{g \in S} c(S \setminus \{g\})$  and  $h = \min_{g \in S} c(S \setminus \{g\}) + 1$  — then choosing an endpoint by an unbiased coin whenever  $\ell \neq h$ . Independent fair coins put almost all of the mass on functions that wander in the middle of the feasible region. Structured extremal functions, which take the same endpoint at every subset, are exponentially unlikely: on  $m$  items there are  $2^m - 1$  nonempty subsets and a function that requires a prescribed endpoint at each of them is sampled with probability  $2^{-(2^m - 1)}$ , which is  $1/128$  at  $m = 3$  and about  $3 \times 10^{-5}$  at  $m = 4$ , per agent.

This matters because the two instances that carry the results of this paper are of exactly that kind. The obstruction of Theorem 31 uses  $c_1(S) = \max(0, |S| - 1)$  alongside  $c_2 = c_3 = |S|$ , all three of which are endpoint-constant; drawing that triple at  $m = 3$  has probability  $2^{-21} \approx 5 \times 10^{-7}$ , so a run of  $3.5 \times 10^4$  samples would be expected to produce it about once in every sixty runs. Both instances were found by hand or by a targeted search, not by the sampler. The sampler has, in effect, never looked in the region where the interesting behaviour lives.

**Remark 132** (The experiment worth running). If Conjecture 2 is false, the counterexample is in the corner the current generator cannot reach, and no amount of additional uniform sampling will find it. An adversarial generator — biased towards endpoint-constant functions, towards supermodular costs, towards agents whose free chores overlap heavily, and towards instances in which the utilitarian optimum is unique — is the single highest-value experiment outstanding against the *conjecture*.

It has now been built (`routeA_adv.py`, `twotier.py`) and it works: it refuted the invariant of Section 10 in a single run, on an instance that  $\approx 1.2 \times 10^4$  uniformly sampled instances had missed. Pointed at Conjecture 2 itself it finds nothing. Since binary additive instances are settled by [LMS26], a counterexample must be non-additive and the search was restricted accordingly: exhaustive over all structured non-additive triples at  $m = 3$  (120 instances) and  $m = 4$  (5,984), then  $\approx 8.4 \times 10^3$  endpoint-constant instances up to  $n = 5$ ,  $m = 6$ . **No instance without a good allocation was found.** This is the first evidence for the conjecture produced by a method demonstrated to find counterexamples on this problem, and it is a genuine update: the accumulating failures of Sections 6, 7, 9 and 10 are evidence about the *methods*, not about the statement.

**Remark 133** (Where the difficulty is concentrated). The same sweep records how often a subsidy is needed at all. Across 11,884 adversarial non-additive instances an *exactly* envy-free allocation — zero subsidy — exists in 98.6% of them; only 164 require any subsidy, and among those the minimum number of paid agents is 1 or 2 at  $n = 3$  and never exceeds 3 at  $n = 4$ .

Two cautions. The pool is adversarially structured rather than representative, so the percentage describes that pool and is not a claim about the model. And the residual 1.4% is not negligible in the way the number suggests: deciding whether an exactly envy-



free allocation exists is NP-complete already for binary additive chores [BSV21], so those instances are hard for a reason and they are precisely the ones any construction must handle.

## 14.2 The peel-frame search

The highest-value experiment against the *approach* of Section 8 is different, and it is cheap. Reuse the `gen.py/search.py` harness over the full  $n = 3$ ,  $m = 3$  dichotomous family — all 9,880 instances up to agent symmetry — and for each instance run the fixed-point computation of `peel3.py`: is the root good, and how many dead ends are there? An instance with a bad root refutes Conjecture 50 and kills the peel approach.

It would *not*, by Remark 51, refute Conjecture 2: a good terminal state can in principle exist while every schedule reaching it passes through an illegal intermediate state. The two searches therefore test different things and both are worth running — the allocation-space search above tests the conjecture, and the peel-frame search tests whether this route to it survives.

Two follow-ups, in order. First, test the hypothesis of Remark 49 across that family: is every dead end a state already committed to an unbalanced terminal? If so, formulate the balance rule precisely and check whether greedy-under-that-rule reaches a terminal without backtracking. Second, try candidate rules in the order (a) balance-first, peel so that a balanced terminal remains reachable; (b) spectator-free peels first by Lemma 45, ties broken by balance; (c)  $\arg \max_i \ell(i)$  with a balance tie-break — noting that (c) alone already fails on the witness.

## 14.3 Reproducing the results

| Script                      | Purpose                                                                           |
|-----------------------------|-----------------------------------------------------------------------------------|
| <code>gen.py</code>         | enumerate all dichotomous functions on $m$ items                                  |
| <code>search.py</code>      | exhaustive search, arguments $m$ $n$                                              |
| <code>fast.py</code>        | randomised counterexample hunt, $m$ $n$ $T$ $\text{seed}$                         |
| <code>msw.py</code>         | utilitarian-optimal test, $m$ $n$ $T$ $\text{seed}$                               |
| <code>tiebreak.py</code>    | selection rules of Section 7.2                                                    |
| <code>localsearch.py</code> | local minima of single-transfer search                                            |
| <code>deadend.py</code>     | Theorem 31                                                                        |
| <code>binadd.py</code>      | Theorem 19                                                                        |
| <code>mswcex.py</code>      | Proposition 34                                                                    |
| <code>peel.py</code>        | the schedule and trap-state refusal of Section 8.4                                |
| <code>peel3.py</code>       | the state-space table of Proposition 48                                           |
| <code>dupsep.py</code>      | Theorems 57, 59, 60; <code>dupsep.py</code> $p$ family for the $P$ -family sweeps |
| <code>typeorder.py</code>   | type-ordered schedules on the witness and small families                          |
| <code>stress.py</code>      | randomised type-ordering sweep, seeded                                            |
| <code>rules.py</code>       | the refinement rules and the matching lift                                        |
| <code>checkD.py</code>      | Remark 37                                                                         |

and, continuing,

| Script                | Purpose                                                                   |
|-----------------------|---------------------------------------------------------------------------|
| routeA.py             | Proposition 66; exhaustive $n = m = 3$ and randomised sweeps              |
| routeA_adv.py         | adversarial hunt; finds Proposition 67                                    |
| routeA_cex.py         | Propositions 67 and 68                                                    |
| twotier.py            | Remarks 132 and 133                                                       |
| hardcases.py          | Remark 75                                                                 |
| n3check.py            | Proposition 71 and Corollary 72                                           |
| monotone_check.py     | Remark 40: monotonicity alone does not suffice                            |
| targetG.py            | Conjecture 77, exhaustive $n = m = 3$                                     |
| targetG_adv.py        | Target G against the Approach 5 adversary                                 |
| guidedR3.py,          | guided-insertion greedy; Section 11.3                                     |
| guidedR3_test.py      |                                                                           |
| guidedR3_full.py      | the insertion reachability obstruction                                    |
| verify_reach_gap.py   | algorithm-free confirmation of that gap                                   |
| targetGbal.py         | Definition 79, exhaustive over balanced partitions                        |
| targetGbal_local.py   | the swap-based repair, round-robin start                                  |
| imwpm_raw.py          | the IMWPM warm start alone                                                |
| targetGbal_stress.py  | the stress sweep that found the 14 local optima                           |
| classify_failures.py  | classifies the 14 as search artefacts                                     |
| restart_check.py      | confirms all 14 recovered by restarts                                     |
| final_algorithm.py    | the combined algorithm; the 389,215-instance sweep                        |
| n2proof_check.py      | Theorem 84 and Corollary 87, exhaustive $m \leq 4$ plus random $m \leq 7$ |
| balanced_msw.py,      | the min-cost-balanced rule and its tie-breaks                             |
| tiebreak_balanced.py, |                                                                           |
| ruleD_adv.py          |                                                                           |
| structural.py         | Theorem 23, Corollary 24, two-type sweep                                  |

The instance of Proposition 34 was produced by `msw.py 4 3 3000 11` at trial 2460 and is re-derived independently by `mswcex.py`, which re-checks that all three cost functions are negative dichotomous, that the utilitarian optimum is unique, and that a zero-subsidy allocation exists outside it.

## 15 Ideas and Open Threads

A running log, newest last. Nothing in this section is claimed to be true.

### 2026-08-01 — What the two negative results jointly require

Taking Theorem 31 and Proposition 34 together, any correct algorithm must be free to (i) move already-allocated chores between bundles and (ii) give up total cost while doing so. Every construction that currently works does so by exhibiting a per-agent potential  $\phi$  with  $w_{\mathcal{A}}(i, j) \leq \phi_i - \phi_j$ , which telescopes along paths and vanishes around cycles. *Finding*

such a potential for general negative dichotomous costs, or proving that none exists, is the problem. In Theorem 19 the potential is  $\phi_i = u_i$ , the count of universally disliked chores held; in Theorem 21 it is  $\phi_i = c(A_i)$ . What is the right  $\phi$  in general?

Worth noting that the existence of such a  $\phi$  is not merely sufficient but close to necessary: if  $p^* \in \{0, 1\}^n$  then  $\phi_i := p_i^*$  satisfies  $w_{\mathcal{A}}(i, j) \leq \phi_i - \phi_j$  by definition of an envy-free solution. So the search for a potential and the search for a proof are the same search, and the content is in constructing  $\phi$  and the allocation *together*.

## 2026-08-03 — routes surviving the two no-go results

Section 9 closes valuation-level encoding and algorithm-level weight forcing, and Corollary 36 closes every rule that selects inside the utilitarian-optimal set. What remains splits by which of those failures it avoids.

**Route M — iterated matching.** [LMS26, §3] develops the objective-chores analogue of iterated maximum-weight *perfect* matching, padding with zero-valued dummies so every agent takes one object per round, which makes the output balanced. The obvious lift to our setting — each round assign every agent one remaining chore minimising the total marginal  $\sum_i [c_i(A_i \cup \{g_i\}) - c_i(A_i)]$  — does *not* give  $\ell \leq 1$  for non-additive dichotomous costs: it fails on 13/250 instances at  $n = 3, m = 4$  and 10/120 at  $n = 3, m = 5$  (rules.py). That is diagnostic rather than discouraging, and [LMS26] flags the obstacle itself: the telescoping needs additivity, because an additive round- $t$  cost is independent of what came before whereas a dichotomous marginal depends on the bundle already held. The repair to attempt is a round invariant replacing the additive domination — each round’s matching minimum-cost *conditioned on the current bundles*, with marginals non-increasing across rounds, a “peel the free chores first” discipline. **Status:** the naive lift is refuted; the conditioned version is untested.

**Route A — induct on agents, not items.** Every approach so far inducts on items. [LMS26, §4.3.1] instead maintains a set of active agents, adds one at a time, and restores two invariants: compensated utility at least  $\lambda$  with  $\lambda \leq p_i \leq 1$ , and an equality path from every active agent to the set of minimum compensated utility. For chores this is attractive because a newcomer starts with an empty bundle, which is the *cheapest* bundle, so the newcomer is the natural sink and the orientation problem of Section 6 does not arise. The invariant carried is the equality graph rather than a subsidy scalar — which is precisely the extra information Theorem 62 and Proposition 53 say a recipient rule needs. **Status:** untried, and the most promising of these.

**Route N — non-redundancy instead of coverage.** Theorem 62 says the failing placements are exactly the zero-marginal ones, and the literature already forbids those under the name non-wastefulness. A duplicate is redundant for its holder by construction, so ask for a *non-redundant* allocation of the replica set — which General Yankee Swap is built to maintain — and measure how far non-redundancy falls short of coverage. The two are not equivalent: a duplicate can be redundant for everyone when every other agent already holds the type or has zero marginal for it. But the residual gap is narrower than coverage itself and is exactly characterisable. **Status:** cheap to test, untested.

**Route X — matroid exchange.** As in Section 8.8. Highest ceiling, highest cost; the parallel-extension step needs a full proof before anything is built on it.

**Ranking.** Route A, then Route M’s conditioned repair, then Route N, then Route X. Conjecture 52 — the type-ordered peel — stays live alongside them and has the best evidence of anything currently open.

## 2026-08-06 — a cleaner rule than construct-and-repair

**The idea.** Among cardinality-balanced allocations of minimum total cost, take one minimising the number of envious ordered pairs  $\# \{(i, j) : c_i(A_i) > c_i(A_j)\}$ . A single two-level optimisation: no local search, no restarts, no schedule, no warm start.

**Why it might work.** It sidesteps Corollary 36 because the counterexample there has a *unique* optimum of sizes  $(0, 3, 1)$ , which is not balanced — restricting the optimisation to balanced allocations removes it from the feasible set entirely. A cost-minimising balanced allocation is automatically envy-freeable, since any reassignment of its own bundles is another balanced allocation and so cannot be cheaper, which means only  $\ell \leq 1$  remains to be shown. The second criterion is the one aligned with Proposition 15: in a  $\{0, 1\}$  solution envy runs only from the paid tier to the unpaid one, so minimising envy count pushes towards exactly the two-tier shape.

**Status.** Zero failures in the *strong* form — every allocation the rule can select is good — over 41,196 instances: exhaustive over all dichotomous triples at  $n = m = 3$ , exhaustive over structured families at  $m \leq 4$ , and endpoint-constant sweeps to  $n = 5, m = 6$  (update\_8/ruleD\_adv.py). Without the tie-break the cost-minimising balanced set always *contains* a good allocation but is not uniformly good; the three other natural tie-breaks — leximin own-cost, minimum total perceived load, leximin perceived load — each fail at  $n = 3, m = 4$  (tiebreak\_balanced.py).

**What would kill it.** A balanced instance where every minimum-cost, minimum-envy allocation needs subsidy 2. **Not proved**, at any  $n$ : an attempt at  $n = 3$  stalled on the exchange step, because dichotomous costs need not be submodular and a bundle of positive cost need not contain any single item whose removal lowers it — so “ $c_u(A_u)$  is large, therefore some swap improves” is unavailable. That is the same missing exchange property that makes the matroid-rank subclass (Section 8.8) the natural next target: there it *is* available.

## 2026-08-06 — at most two distinct cost functions

**The idea.** Suppose the  $n$  agents carry only *two* distinct cost functions  $c$  and  $c'$ , in groups of sizes  $k$  and  $n - k$ . This strictly generalises the identical-cost case of Theorem 21 and is not covered by anything cited here: the literature settles identical, additive and submodular costs, but “two types of agents” is a different slice — one that does appear in the house-allocation work of [DCW<sup>+</sup>24] as a tractable restriction, which is mild evidence it is the right kind of hypothesis.

**Why it might work.** The two-tier conditions collapse dramatically. For two agents of the same type in the same tier, (a) forces  $c(A_i) = c(A_j)$  exactly; for two of the same type in different tiers, (b) and (c) force  $c(A_i) = c(A_j) + 1$  exactly. So within each type group the bundle costs take at most two values, differing by exactly one, and the tier split inside that group is determined by which value a bundle takes. What remains is a finite system of

cross-group inequalities in two functions rather than  $n$ , which is the kind of thing a direct construction can attack.

**Status.** Untried as a proof. Zero failures over 7,805 two-type instances at  $m \in \{3, 4\}$ ,  $n \in \{3, 4\}$ , both for existence and for the rule of the previous entry (`update_9/structural.py`). Worth one session: it is the first class in a while where the conditions genuinely simplify rather than merely being restricted.

**What would kill it.** A two-type instance with no good allocation. Note this would also refute Conjecture 2 outright, so the search is worth running for its own sake.

## 2026-08-06 — induct on the depth of the cost functions

**The idea.** Stop inducting on items, on agents, or on schedules, and induct on the *cost functions themselves*. Every dichotomous  $c$  decomposes into nested monotone Boolean layers: putting  $\mathcal{F}_t := \{S : c(S) \geq t\}$ , each  $\mathcal{F}_t$  is upward-closed,  $\mathcal{F}_1 \supseteq \cdots \supseteq \mathcal{F}_K$  with  $K = c(M)$ , and

$$c(S) = \#\{t : S \in \mathcal{F}_t\} = \sum_{t=1}^K \mathbb{1}[S \in \mathcal{F}_t].$$

Crucially the class is closed under truncation:  $\min(c, k)$  is dichotomous for every  $k$ , so layers can be peeled one at a time without leaving the model. Define the *depth* of an instance to be  $K = \max_i c_i(M)$ .

**Why it might work.** The base case is already proved. A depth-one instance has every  $c_i$  taking values in  $\{0, 1\}$ , so every bundle costs at most 1 to everybody — which is exactly the hypothesis of Theorem 23. So *Conjecture 2 holds unconditionally at depth 1, for every  $n$  and every  $m$* . No other induction in this project has a proved base case, and none of the failed approaches touches the representation of the cost functions at all: they all manipulate allocations while treating  $c_i$  as a black-box oracle.

**Evidence for the inductive step.** Over all 1,330 depth- $\geq 2$  instances at  $n = m = 3$ , the full instance and its truncation  $(\min(c_i, K - 1))_i$  *always* share at least one common good allocation — 100%, with no instance where one is solvable and the other is not (`update_10/layers.py`). That is precisely the compatibility an induction on depth needs.

**The kill condition, hunted and not found.** A depth- $\geq 2$  instance whose good allocations are disjoint from its truncation's would end the approach. Across 39,692 depth- $\geq 2$  instances — exhaustive over all dichotomous triples at  $n = m = 3$ , exhaustive over structured families at  $m \leq 4$ , and endpoint-constant and deliberately high-depth sweeps to  $n = 5$  — there is **no such instance**, in either peeling direction (`update_10/layer_hunt.py`). This is the generator that refuted Proposition 67 on its first run.

The same sweep says which direction to peel. Measuring the fraction of the truncation's good allocations that survive into the full instance, top-peeling  $c \mapsto \min(c, K - 1)$  carries far more than bottom-peeling  $c \mapsto \max(c - 1, 0)$  — worst case 0.21 against 0.05 — and on high-depth instances top-peeling carries *all* of them, at  $n \in \{3, 4, 5\}$ .

**The strengthened hypothesis, and why it is the right one.** Sharing a witness is necessary but not sufficient: an induction must *name* which good allocation to carry forward. The fix is induction loading — prove something stronger, so that the inductive hypothesis is strong enough to reproduce itself.

**Chain conjecture.** For every dichotomous instance of depth  $K$  there is a *single* allocation that is good for  $(\min(c_i, k))_i$  *simultaneously* for every  $k = 1, \dots, K$ .

Its top level  $k = K$  is Conjecture 2; its bottom level  $k = 1$  is Theorem 23, already proved. It has **zero failures over 30,405 depth- $\geq 2$  instances** across the same exhaustive and adversarial sweeps. Being stronger it is a far better induction hypothesis than Conjecture 2 itself, since carrying a whole chain of certificates upward is exactly what the step needs and what the plain statement cannot supply.

**Progress on the step: the chain collapses to two levels.** Write  $a_{ij} := c_i(A_j)$  for the cost matrix of an allocation. Every level- $k$  arc depends on the matrix alone,

$$w^{(k)}(i, j) = \min(a_{ii}, k) - \min(a_{ij}, k),$$

and is therefore *sign-constant and saturating*: it has the sign of  $w(i, j)$  throughout, its modulus is non-decreasing in  $k$ , and it reaches  $|w(i, j)|$  once  $k \geq \max(a_{ii}, a_{ij})$ . One consequence is immediate and proved.

**Lemma 134** (Level one is exactly envy-freeability). *An allocation is good for  $(\min(c_i, 1))_i$  if and only if it is envy-freeable for  $(\min(c_i, 1))_i$ . No subsidy condition is needed.*

**Proof**  $\min(c_i, 1)$  takes values in  $\{0, 1\}$ , so every bundle costs at most 1 to every agent and Theorem 23 applies to any envy-freeable allocation. The converse holds because goodness includes envy-freeability.  $\square$

**A two-level reduction was proposed here and is now REFUTED.** It rested on two statements about a single allocation's set of good levels:

**(I) Interval.** That set is always an interval.

**(II) Endpoints.** Good at level 1 and level  $K$  implies good at every level between.

Both survived 2,288,509 (instance, allocation) pairs, and both are **false**.

The refutation came from a change of search space that is worth keeping. Since the level- $k$  arcs depend only on the matrix  $a$ , and *every* non-negative integer matrix is realisable — take  $c_i(S) := \sum_j \min(|S \cap A_j|, a_{ij})$ , whose marginals are all 0 or 1, with bundles of size  $\max_i a_{ij}$  — claims (I) and (II) are equivalent to statements about all integer matrices, which can be *enumerated* rather than sampled. Doing so finds 972 counterexamples among the 261,632 matrices with entries at most 3 (`update_11/matrix_interval.py`). The smallest is

$$a = (c_i(A_j)) = \begin{pmatrix} 0 & 0 & 0 \\ 2 & 2 & 1 \\ 3 & 3 & 2 \end{pmatrix},$$

realised at  $n = 3, m = 8$  by  $c_1 \equiv 0$  and the two formulas above. That allocation is good at level 1, has a *positive cycle* at level 2, and is good again at levels 3 through 8 — so its good-level set  $\{1, 3, 4, \dots, 8\}$  is not an interval, and (II) fails at its endpoints (`update_11/verify_cex.py`). Every earlier sweep missed it only because none reached  $m = 8$ .

**What survives.** The chain conjecture is untouched: (I) and (II) concern one allocation's level set, whereas the chain conjecture asserts the *existence* of some allocation good at every



level. On the refuting instance a chain witness does exist — give everything to agent 1, who finds all of it free. So the collapse to two levels is gone, but the target is not.

**The chain conjecture survives past the frontier.** Every earlier chain test stopped at  $m = 6$ , while the counterexample above lives at  $m = 8$ , so the conjecture had never been tried where the level structure is known to misbehave. Re-run at  $m = 7, 8, 9$  for  $n = 3$  and at  $m \leq 8$  for  $n \in \{4, 5\}$  — using both endpoint-constant sampling and the *matrix-realising* family  $c_i(S) = \sum_j \min(|S \cap B_j|, a_{ij})$  that produced the refutation — it has **zero failures** (update\_12/chain\_hunt.py). Neither did any instance fail Conjecture 2. So the conjecture held exactly where its own shortcut broke.

**A proved reduction.** The natural sufficient condition turns out to be the strongest possible one, and it is two lines.

**Lemma 135** (Exact envy-freeness implies chain-goodness). *If  $\mathcal{A}$  is exactly envy-free for  $(c_i)_i$  — that is  $c_i(A_i) \leq c_i(A_j)$  for all  $i, j$  — then  $\mathcal{A}$  is good at every level, with zero subsidy at every level.*

**Proof** Exact envy-freeness says  $a_{ii} \leq a_{ij}$  for all  $i, j$ . Since  $x \mapsto \min(x, k)$  is non-decreasing,  $\min(a_{ii}, k) \leq \min(a_{ij}, k)$ , so every level- $k$  arc  $w^{(k)}(i, j)$  is at most 0, for every  $k$ . Every cycle is then a sum of non-positive terms, so no positive cycle exists at any level; and every directed path likewise has weight at most 0, so  $\ell^{(k)} \equiv 0$ .  $\square$

**Corollary 136.** *The chain conjecture holds for every instance that admits an exactly envy-free allocation. Equivalently, the chain conjecture is equivalent to its own restriction to instances admitting no exactly envy-free allocation.*

This localises the difficulty sharply. By Remark 133 an exactly envy-free allocation exists in about 98.6% of adversarial instances, so Corollary 136 disposes of almost everything and leaves a residual of roughly 1.4%. That residual is not a rounding error: it is precisely the class on which deciding envy-freeness is NP-complete even for binary additive costs [BSV21]. Across the 24 envy-free-free instances encountered in the sweeps above, every one still had a chain witness, so the residual is not obviously where a counterexample lives either.

**Status.** Proved: the decomposition, truncation closure, the base case (Theorem 23), Lemma 134, Lemma 135 and Corollary 136. Refuted: (I), (II) and the two-level reduction. Open: the chain conjecture restricted to instances with no exactly envy-free allocation — a strictly smaller target than before, and one that now coincides with the known hardness core rather than straddling it.

**The residual class, generated deliberately.** Since EF-free instances are about 1.4% of the population they must be constructed rather than sampled. Two sources were used: rejection sampling from the endpoint-constant and matrix-realising families, and R10’s Set-Splitting reduction, which is EF-free *by construction* on a hypergraph with no proper 2-colouring — the triangle  $U = \{v_1, v_2, v_3\}$ ,  $F = \{v_1v_2, v_1v_3, v_2v_3\}$ , giving  $n = 5$ ,  $m = 6$  with binary additive costs. Across 785 EF-free instances so obtained, **every one has a chain witness** (update\_13/residual.py). On the Set-Splitting instance itself there are 1620 of them.

Two hypotheses this suggested, and their fates. On Set-Splitting the chain-good and top-good sets coincide exactly, both 1620, which raised the possibility that on the residual the chain conjecture collapses to Conjecture 2. It does not: over 363 further EF-free instances, 78 have strictly fewer chain-good than top-good allocations. So the chain conjecture remains

strictly stronger than Conjecture 2 even after Corollary 136 removes the easy instances — there is no collapse, and the layer structure retains content exactly where the difficulty lives. Separately, the min-cost-balanced/min-envy rule of the earlier entry selects a chain witness on 420 of 422 residual instances but not all, so it is not the construction either.

**The Set-Splitting family, and a correction.** R10’s reduction is the only generator of *certified* EF-free instances, so it was rebuilt as a family over hypergraphs rather than the single triangle instance of `update_3`. The reconstruction taken from that addendum’s summary is *wrong* and failed its own correctness check — it produced an EF-free instance from a splittable hypergraph. Reading [BSV21, proof of Theorem 2] directly supplies the missing clause: when  $r < r'$  the padding agents  $e_{r+1}, \dots, e_{r'}$  are “imaginary hyperedges adjacent to the *entire* set of vertices”, so they cost 1 on *every* vertex chore, not 0. With that fixed the reduction verifies in both directions — an exactly envy-free allocation exists precisely on splittable hypergraphs — across seven hypergraphs including two YES controls (`update_14/setsplit_family.py`).

On the corrected NO instances the structure is sharp. Chain-good and top-good coincide *exactly* on every one of them (6, 108, 72, 1620, 1620, 3840 witnesses respectively), and the total subsidy is the same for every witness of a given instance, always  $n - 2$ . The dominant form is describable: the  $r'$  dummy chores are spread one per agent over  $r'$  of the  $n = r' + 2$  agents, the two dummy-free agents hold only chores they value at 0, and *the paid set is exactly the set of dummy holders*. That account was exact up to  $n = 5$  but covered only 96% of the 30,276 witnesses at  $n = 6, m = 8$ . The exceptions are fully explained, and the explanation removes them.

**The exceptions are an artefact of padding.** The reduction sets  $r' = \max(q, r)$  and, when  $r < r'$ , fills the gap with imaginary hyperedges adjacent to every vertex. Those padding agents do not behave like real edge agents, which are free on vertex chores outside their own edge. Comparing two instances at the same  $n = 6$  settles it: with  $q = 4, r = 2$  — two padding agents — the form covers 96.3% and total subsidies of 1 and 2 appear alongside  $n - 2 = 4$ ; with  $q = 3, r = 4$  — *no* padding — the form covers 100% of all 28,800 witnesses, the paid set is the dummy-holder set in every one, and the total subsidy is  $n - 2$  with no exceptions. Every zero-padding instance tested behaves this way.

**Characterisation (exact on all zero-padding instances tested).** Let  $(U, F)$  be a NO instance of Set Splitting with  $r \geq q$ , so that  $r' = r$  and every edge agent is real. In R10’s reduced instance an allocation is a chain witness *if and only if* the  $r$  dummy chores are distributed one per agent among  $r$  of the  $n = r + 2$  agents and the two dummy-free agents hold only chores they value at 0. Its paid set is exactly the set of dummy-holders and its total subsidy is exactly  $n - 2 = r$ .

The restriction  $r \geq q$  costs nothing: duplicating hyperedges raises  $r$  without changing splittability, so Set Splitting stays NP-hard on it and the family stays certified-hard. This is now an exact description of every chain witness on an infinite family of hard instances, stated in terms of the hypergraph rather than of the allocation search — which is what a proof needs as its starting point.

**The non-additive residual, which is the class that actually matters.** R10’s family is binary *additive*, so Conjecture 2 already holds on all of it via [LMS26], and the characterisation above is the binary-additive theorem’s shape reappearing. The class that carries the open

problem is the *non-additive* envy-free-free instances, which that reduction cannot produce. Generating 553 of them by rejection sampling with an explicit additivity filter settles three things (`update_15/nonadditive_residual.py`).

*The chain conjecture survives.* All 553 have a chain witness — 79,206 of them in total, none missing. This is the most adversarial test the open statement has had, being simultaneously envy-free-free and non-additive.

*The Set-Splitting characterisation does not survive.* “Paid set = holders of universally costly chores” holds on 22 of 553, and “total subsidy constant across the witnesses of an instance” on 56 of 553. Both were exact on the additive family. So they are artefacts of additivity, and a proof of the characterisation on R10’s family would not have generalised — the check was worth running before the proof attempt rather than after.

*And it corrects an earlier lead.* The containment  $\text{paid} \subseteq \arg \max_i c_i(A_i)$ , recorded as 164/164 in Remark 75, fails on 141 of the 553. That sample had not been filtered for additivity. The remark is amended accordingly; what survives there is Proposition 74, which is proved rather than observed.

**Back at the induction: what set is it aiming at?** Since level- $k$  goodness depends only on  $a$ , and every non-negative integer matrix is realisable, “which matrices are chain-good” is decidable by enumeration. Doing it over all 261,632 matrices with entries at most 3 at  $n = 3$  (`update_16/chain_matrices.py`) kills the two natural invariants and leaves one lemma.

*Refuted: carry a bipartition.* By Proposition 15, goodness at level  $k$  means *some* bipartition  $S_k$  satisfies the two-tier inequalities at that level. If one  $S$  worked at every level the induction could simply carry it. It cannot: 4,911 matrices are chain-good with no single bipartition valid throughout. *Refuted: the tier set is stable.* The canonical set  $S_k = \{i : \ell^{(k)}(i) = 1\}$  genuinely moves with  $k$  — 12,822 violations. The smallest witness is

$$a = \begin{pmatrix} 0 & 0 & 0 \\ 2 & 1 & 2 \\ 2 & 1 & 2 \end{pmatrix}, \quad \ell^{(1)} = (0, 0, 0), \quad \ell^{(2)} = (1, 0, 1),$$

which is chain-good while its paid set grows from  $\emptyset$  to  $\{1, 3\}$ . So no invariant of the form “fix who gets paid” can work.

*What survives, and it is provable.*

**Lemma 137** (Breakpoints). *Let  $T := \{k \in [1, K] : k = a_{ij} \text{ for some } i, j\} \cup \{K\}$ . An allocation is good at every level in  $[1, K]$  if and only if it is good at every level in  $T$ . Hence the chain condition amounts to at most  $n^2 + 1$  tests, independently of the depth  $K$ .*

**Proof** Between two consecutive elements of  $T$  each  $\min(a_{ij}, k)$  is affine in  $k$  — equal to the constant  $a_{ij}$  if  $a_{ij}$  lies below the interval, and to  $k$  if it lies above — so every arc, and hence every cycle weight, is affine there. An affine function that is at most 0 at both endpoints is at most 0 throughout, so no positive cycle can appear strictly inside an interval. Given that,  $\ell^{(k)}$  is the maximum of the finitely many affine functions given by simple paths, together with 0; a maximum of affine functions is convex, and a convex function at most 1 at both endpoints is at most 1 throughout.  $\square$

Verified with 0 violations over the same 261,632 matrices. The consequence is worth stating: the chain conjecture’s difficulty does *not* grow with depth. Layers were introduced

to give the induction a parameter to climb, and Lemma 137 says the climb is bounded by  $n^2 + 1$  steps no matter how tall the stack. An induction on  $K$  is therefore the wrong shape — the right parameter is the number of distinct cost values an allocation exhibits, which is at most  $n^2$  and has nothing to do with  $K$ .

**Method note, worth more than any single result here.** For any question about a *fixed* allocation, work with the cost matrix and enumerate rather than sampling instances: every non-negative integer matrix is realisable, so matrix enumeration is exhaustive over a space that instance enumeration only samples, and it reaches effective sizes instance search cannot. The (I)/(II) counterexample sat two items beyond every sweep that preceded it and fell out of matrix enumeration immediately.

## 2026-08-06 — how big must the paid set be, and what the target really is

**Two corrections first.** (i) The previous entry closed by suggesting the  $n = 3$  analysis was now small enough to do by hand, on the grounds that Corollary 72 makes the 2-path  $i \rightarrow j \rightarrow k$  free whenever  $a_{ki} \leq a_{kk} + 1$ . Counting how many *hard pairs* — ordered pairs with  $a_{ki} > a_{kk} + 1$  — a good matrix can carry gives up to 6 at  $n = 3$  (update\_17), so a case analysis branched on the hard-pair pattern has  $2^6$  cases, not a handful. That framing was oversold. (ii) The first measurement run minimised the total subsidy over *all* allocations, which is the wrong quantity: it returns spikes such as  $p = (2, 0, 0)$ , cheap in total but outside  $\{0, 1\}^n$  and so irrelevant to Conjecture 2. The corrected script update\_17/minimal\_S.py minimises  $|S|$  over *good* allocations.

**What the measurement says.** Over 11,163 instances at  $n \in \{3, 4, 5\}$ ,  $m \leq 6$ :

| minimal $ S $ over good allocations | instances |         |
|-------------------------------------|-----------|---------|
| 0 (exactly envy-free)               | 10,559    | (94.6%) |
| 1                                   | 486       | (4.4%)  |
| 2                                   | 93        | (0.8%)  |
| 3                                   | 24        | (0.2%)  |
| 4                                   |           | 1       |
| no good allocation at all           | 0         |         |

The maximum of the minimal  $|S|$  is exactly  $n - 1$  at every  $n$  tested, always with witness  $p = (1, \dots, 1, 0)$ . So the answer to “is the truth stronger than Conjecture 2?” is *no*: no constant bound, and no bound growing more slowly than  $n - 1$ , can hold. Extracting the smallest extremal instance returns  $n$  identical additive agents with  $n - 1$  unit chores — which is Example 1, already in the report. The run therefore re-derives the known lower bound rather than finding a new one; what it does add is exhaustive confirmation of that example’s “no allocation does better” clause, for  $n \leq 7$  (update\_17/tight\_family.py).

**Two things that are new.** First, Proposition 13:  $p^*$  always has a zero coordinate, so the total-subsidy clause of Conjecture 2 follows from the per-agent one and the target is the single inequality  $\max_i \ell_{\mathcal{A}}(i) \leq 1$ . Second, in 9 of the instances the minimum total over all allocations is strictly below the minimum over good ones — at  $n = 4$ , 2 against 3. So Conjecture 2 is *not* the statement that the utilitarian-optimal subsidy is at most  $n - 1$ ; the

$\{0, 1\}$  shape genuinely costs money, and any argument that reasons about the total alone will prove the wrong theorem.

**The reframing worth carrying forward.** Reading the two-tier characterisation as a statement about matrices removes the subsidy vector entirely:

$$\text{Conjecture 2} \iff \exists \mathcal{A}, \exists S \subseteq N : c_i(A_j) \geq c_i(A_i) - [i \in S] + [j \in S] \quad \forall i \neq j.$$

What is left is a partition and a *bipartition* — and there are only  $2^n$  of the latter. This is the first form of the problem in which the existential over subsidies is finite and small: guess  $S$ , then ask whether a partition meeting a fixed system of local constraints exists. Every earlier formulation had to discover the subsidies and the bundles simultaneously.

**Status.** Proposition 13 proved. Tightness confirmed for  $n \leq 7$ . Conjecture 2 unrefuted over these 11,163 instances. Next: for fixed  $S$ , characterise when the constraint system is satisfiable — note that  $S = N$  and  $S = \emptyset$  impose *the same* constraints, namely exact envy-freeness, which is why the paid set is never everybody.

**Followed up the same day**, in Section 12. Taking the discounted-cost reading seriously turns the problem into local search over partitions, and Theorem 119 shows that a single descent lemma about single chore transfers implies Conjecture 2 outright, polynomial-time clause included. That is the shortest remaining gap any approach in this report has produced, and Lemma 120 narrows the hard half of it to balanced partitions. It is where work should continue.

## 16 Conclusion and Future Work

We have not settled Conjecture 2 in general, but Section 11 now carries a construct-and-repair algorithm together with a proof of most of what one would want from it. Four properties are theorems for every  $n$ : the algorithm’s outputs are *certified*, so it never lies (Theorem 82); it always terminates, in polynomial time  $O(n^3 m^3 \tau + n^4 m^3)$  (Theorem 83); its search space is genuinely the unordered balanced partitions (Lemma 81); and for *two agents* it is *complete* (Theorem 84) — every two-agent dichotomous instance admits a balanced split of  $q$ -spread at most 1, found in  $O(m\tau)$  time with no restarts. That last result even strengthens the known  $n = 2$  chore case, adding a balancedness guarantee the complementation route of Remark 9 does not provide.

For  $n \geq 3$  what is missing is completeness alone, and it is backed by **zero observed failures across 389,215 test instances**, including every hard instance in this project. So the headline is a proved algorithm whose only unproved property is that it always succeeds — a single, sharply stated combinatorial claim (Conjecture 89), with a proved base case and a named proof technique to generalise.

Getting there required ruling out three strategies a reader would reasonably try first: item-by-item insertion, the strategy that works in [BKNS22] (Theorem 31, and again in the pure goods coordinates of Section 11.3); restriction to utilitarian-optimal allocations, together with *every* refinement of that set (Proposition 34, Corollary 36); and encoding the coverage constraint into the numbers, at either the valuation or the algorithm level (Section 9). Alongside those sit a small set of positive results — binary additive costs (Theorem 19, now subsumed by [LMS26]), identical costs (Theorem 21), two agents — and structural tools that say precisely how the chore problem differs from its goods mirror.



**The lesson that ties the whole investigation together.** Every insertion-based strategy failed for the same reason, in three different coordinate systems: chores directly (Section 6), the replica/peel frame (Section 8), and even the pure goods frame (Section 11.3). Building an allocation one item at a time fixes decisions before their consequences are visible, and by the time a bad grouping shows itself, insertion has no way to undo it — only permutation between agents, never a genuine re-split of a bundle’s contents, is available. The algorithm that finally worked abandoned that template entirely: *construct* a whole candidate partition directly, then *repair* it by moves that can freely undo an earlier bad grouping. That is the one idea from this project worth carrying into any future attempt on a similar problem.

The three negative results are more useful together than separately, because they fail in the same way. Insertion fixes an owner before the consequences are visible; a selection rule fixes a criterion before knowing which allocation it will pick out; a reweighting fixes a penalty before knowing who will hold the duplicate. Each is *a rule fixed ex ante against a quantity determined ex post* (Section 9.6). What survives is procedures that decide late — which is exactly the design of the peel frame, where ownership is settled only by the last move touching a chore.

Both negative results are failures of *timing*, and the constraint they jointly impose is that *a correct algorithm must be free to move already-allocated chores between bundles, and free to give up total cost while doing so*. The replica transform of Section 8 is the response: Theorem 39 turns the instance into a goods instance with an identical envy graph, Corollary 41 reduces the conjecture to the single constraint that no agent be relieved of a chore twice, and reading that dynamically defers every ownership decision to the last move touching a chore, so there is nothing to re-open. This is currently the most promising route. It is not known to work — Proposition 48 shows the state space has genuine dead ends, and Remark 51 notes that the reduction is sufficient but not known to be complete — and what is missing is a peel rule that provably never enters a deadlock.

Every construction here that succeeds does so by exhibiting a per-agent potential  $\phi$  with  $w_{\mathcal{A}}(i, j) \leq \phi_i - \phi_j$ , which telescopes along paths. Finding such a potential for general negative dichotomous costs, or proving none exists, remains the problem; the balance signal of Remark 49 is the current best guess at its shape.

## Directions.

1. **Prove Conjecture 89 for  $n \geq 3$**  (Section 11.8). This is now the top-ranked direction by a wide margin, and it has a proved base case: completeness holds at  $n = 2$  for every  $m$  (Theorem 84), constructively and in  $O(m\tau)$  time, by the  $h$ -window technique — a discrete intermediate-value argument on the aggregate  $u = \tilde{v}_1 + \tilde{v}_2$  along swap paths, closed off by extremal endpoint arguments. Soundness, certification, termination and polynomial complexity are theorems for every  $n$  (Theorems 82 and 83). What remains is the  $n \geq 3$  analogue of the window argument: a one-dimensional aggregate whose per-move variation is smaller than the forbidden gap it must cross, controlling  $\binom{n}{2}$  pairwise comparisons at once. The 14 known local optima at  $n = 3$  bound what any such proof must accommodate, and Remark 88 records the move-set refinement (pairwise exchanges) the proof already forced on the implementation.
2. **The type-ordered peel search** (Section 8.7, Section 14.2). Conjecture 52 has the best



evidence of anything currently open and a much smaller state space than Conjecture 50. Run the fixed-point computation over the full  $n = m = 3$  dichotomous family; a bad root anywhere kills the approach of Section 8, a clean sweep is strong evidence for it. Cheapest decisive experiment available. Note what a recipient rule must now satisfy: by Theorem 62 the recipient in the hard case has zero marginal, and by Proposition 53  $\arg \max_i \ell(i)$  is insufficient even under the type-ordering, so the rule has to read the residual workload.

3. **Choose the paid set first** (Proposition 15, Proposition 73). The target is exactly: an envy-freeable allocation together with a bipartition of the agents such that all strict envy runs from one side to the other. This is a condition on a bipartition rather than on a schedule, so it inherits none of the dead ends of Section 8 and none of the ex-ante failures of Section 9. It is the one direction that four separate lines of attack have now pointed at and nobody has tried.
4. **Agent induction with a path-based invariant** (Section 10). The arc-based invariant is refuted, so any revival must carry path weights, which are not assignment-invariant. The equality graph of [LMS26, §4.3.1] is the natural candidate for what to carry. Harder than it looked a session ago, but not closed.
5. **Characterise the dead ends** (Remark 49). Test across that family whether every dead end is a state already committed to an unbalanced terminal. If so, the balance rule is the missing ingredient and can be stated precisely.
6. **Settle the converse** (Remark 51). Is a good terminal state, when one exists, always reachable by a legal schedule? A positive answer upgrades the peel frame from a sufficient reformulation to a complete one, and makes the search of item 1 a direct test of Conjecture 2.
7. **Adversarial instance generation** (Section 14.1). The randomised evidence is drawn from a distribution that essentially never produces the structured extremal functions from which both negative results are built. This is the cheapest way to find out whether Conjecture 2 is true at all.
8. **Exchange paths rather than single transfers** (Section 7.2). Single-chore local search has bad local minima. The move set should match the exchange structure of dichotomous costs, as Yankee Swap and matroid union do in the matroid-rank literature [BEF21]; and by Remark 35 the search must range over all allocations, not the utilitarian-optimal ones.
9. **The matroid-rank subclass** (Section 8.8). Write out the supermodular-to-matroid reduction properly, in particular the parallel-extension step, then attempt basis exchange in the parallel-copy matroid. This is where the surrounding literature's machinery is strongest.
10. **Build the two tiers first** (Proposition 15). Choose the partition into paid and unpaid agents and only then allocate. This is the one natural approach that Theorem 31 does not obviously kill, since it is not incremental in the chores. In the proof of Theorem 19

the paid set is exactly the agents holding a  $\lceil q/n \rceil$ -share of the universally disliked chores; identifying what plays that role in general is the first question.

11. **Climb the milestone ladder.** Binary additive is done (Theorem 19); binary submodular is the next rung, by item 6 above; general dichotomous is the summit. Note that Theorem 16 and Remark 49 agree on what the middle rung needs — cardinality balancing — which is what matroid union would supply.
12. **The mixed model.** Allow each item to be a good or a chore for each agent, with all marginals in  $\{-1, 0, 1\}$ : the dichotomous restriction of the doubly monotone setting of [KMS<sup>+</sup>24]. This is the natural common generalisation of [BKNS22] and of the present paper.

**Where this would sit.** Conjecture 2 is the fourth corner of the square in Section 1: additive goods [BDN<sup>+</sup>20], dichotomous goods [BKNS22] and additive chores [LMS26] all give one dollar per agent and  $n - 1$  in total, and dichotomous chores is what remains. For valuations that are dichotomous but not additive the best published bound is still that of [KMS<sup>+</sup>24],  $n - 1$  per agent and  $n(n - 1)/2$  in total from any EF1 allocation, so the conjecture would improve the total by a factor of  $n$  — the same factor [BKNS22] gains over [BDN<sup>+</sup>20] on the goods side.

Two cautions follow from the literature check. First, Theorem 19 is already subsumed by [LMS26] (Remark 20); what survives of it is the mechanism, not the statement. Second, [LMS26] appeared in July 2026, which is a reminder that the surrounding corners are being filled in quickly and that the search should be repeated before circulation rather than treated as done.

## References

- [BCE<sup>+</sup>16] Felix Brandt, Vincent Conitzer, Ulle Endriss, Jérôme Lang, and Ariel D. Procaccia. *Handbook of Computational Social Choice*. Cambridge University Press, 1st edition, 2016.
- [BDN<sup>+</sup>20] Johannes Brustle, Jack Dippel, Vishnu V. Narayan, Mashbat Suzuki, and Adrian Vetta. One dollar each eliminates envy. In *Proceedings of the 21st ACM Conference on Economics and Computation (EC)*, pages 23–39, 2020.
- [BEF21] Moshe Babaioff, Tomer Ezra, and Uriel Feige. Fair and truthful mechanisms for dichotomous valuations. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 35. AAAI, 2021.
- [BKNS22] Siddharth Barman, Anand Krishna, Y. Narahari, and Soumyarup Sadhukhan. Achieving envy-freeness with limited subsidies under dichotomous valuations, 2022.
- [BMS05] Anna Bogomolnaia, Hervé Moulin, and Richard Stong. Collective choice under dichotomous preferences. *Journal of Economic Theory*, 122(2):165–184, 2005.

- [BSV21] Umang Bhaskar, A. R. Sricharan, and Rohit Vaish. On approximate envy-freeness for indivisible chores and mixed resources. In *Approximation, Randomization, and Combinatorial Optimization (APPROX/RANDOM)*, LIPIcs, 2021.
- [BSY23] Xiaolin Bu, Jiaxin Song, and Ziqi Yu. EFX allocations exist for binary valuations. In *Proceedings of the International Conference on Frontiers of Algorithmic Wisdom (FAW)*, 2023.
- [Bud11] Eric Budish. The combinatorial assignment problem: Approximate competitive equilibrium from equal incomes. *Journal of Political Economy*, 119(6):1061–1103, 2011.
- [CI21] Ioannis Caragiannis and Stavros Ioannidis. Computing envy-freeable allocations with limited subsidies. In *Workshop on Internet and Network Economics (WINE)*, 2021.
- [DCW<sup>+</sup>24] Sijia Dai, Yankai Chen, Xiaowei Wu, Yicheng Xu, and Yong Zhang. Weighted envy-freeness in house allocation, 2024.
- [End17] Ulle Endriss. *Trends in Computational Social Choice*. Lulu.com, 2017.
- [Fol67] Duncan Karl Foley. Resource allocation and the public sector. *Yale Economic Essays*, 7(1):45–98, 1967.
- [GIK<sup>+</sup>21] Hiromichi Goko, Ayumi Igarashi, Yasushi Kawase, Kazuhisa Makino, Hanna Sumita, Akihisa Tamura, Yu Yokoi, and Makoto Yokoo. Fair and truthful mechanism with limited subsidy, 2021.
- [GSW25] Jugal Garg, Eklavya Sharma, and Xiaowei Wu. Proportional and Pareto-optimal allocation of chores with subsidy, 2025.
- [HS19] Daniel Halpern and Nisarg Shah. Fair division with subsidy. In *Proceedings of the 12th International Symposium on Algorithmic Game Theory (SAGT)*, volume 11801 of *Lecture Notes in Computer Science*, pages 374–389. Springer, 2019.
- [KMS<sup>+</sup>24] Yasushi Kawase, Kazuhisa Makino, Hanna Sumita, Akihisa Tamura, and Makoto Yokoo. Towards optimal subsidy bounds for envy-freeable allocations. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2024.
- [LMMS04] Richard J. Lipton, Evangelos Markakis, Elchanan Mossel, and Amin Saberi. On approximately fair allocations of indivisible goods. In *Proceedings of the 5th ACM Conference on Electronic Commerce (EC)*, pages 125–131, 2004.
- [LMS26] Xinhang Lu, Simon Mackenzie, and Mashbat Suzuki. Optimal subsidy bounds for goods and chores: One dollar each suffices, 2026.
- [Mas87] Eric S. Maskin. On the fair allocation of indivisible goods. In G. Feiwel, editor, *Arrow and the Foundations of the Theory of Economic Policy*, pages 341–349. MacMillan, 1987.

- [NSV21] Vishnu V. Narayan, Mashbat Suzuki, and Adrian Vetta. Two birds with one stone: Fairness and welfare via transfers. In *International Symposium on Algorithmic Game Theory (SAGT)*, pages 376–390. Springer, 2021.
- [TWYZ23] Biaoshuai Tao, Xiaowei Wu, Ziqi Yu, and Shengwei Zhou. On the existence of EFX (and pareto-optimal) allocations for binary chores, 2023. Journal version in *Theoretical Computer Science*, 2025.
- [Var74] Hal Varian. Equity, envy, and efficiency. *Journal of Economic Theory*, 9(1):63–91, 1974.